How to Teach Programming (And Other Things)

*what everyone in tech ought to know
about teaching and learning*

Edited by Greg Wilson

Contents

1 Introduction

In Brief

This book shows readers how to build and deliver high-quality learning experiences to people who want to learn how to program (and other things as well). It is based on Software Carpentry[1]'s instructor training course[2], and all material can be freely distributed and re-used under the Creative Commons - Attribution license. Please see http://github.com/gvwilson/thirdbit/ for the source, http://third-bit.com/teaching/ for the online version, and http://third-bit.com/teaching.epub and http://third-bit.com/teaching.mobi for e-book versions, or buy the book at cost.

Thousands of grassroots "learn to code" groups have sprung up in the past few years. They exist so that people don't have to figure out how to program on their own, but ironically, that's exactly what most of their founders are doing when it comes to teaching.

As many have discovered, there's more to teaching than talking. Good teachers break subjects up into digestible pieces, design lessons with verifiable goals in mind, check their students' progress at short intervals, and encourage collaboration and improvisation. Like good programming practices, these don't have to be reinvented by every teacher: they can and should be taught and learned. And while these practices won't magically make someone a *great* teacher, they do make most people *better* teachers.

This book is a brief introduction to modern evidence-based teaching practices and how to use them to teach programming to free-range learners. It covers:

- how people's thinking changes as they go from being novices to competent practitioners and then to being experts;

---

[1] http://software-carpentry.org
[2] https://swcarpentry.github.io/instructor-training/

- how to tell if your learners are keeping up with you, and what to do or say when they're not;

- how to design and improve lessons efficiently and collaboratively;

- how and why live coding (§8) (i.e., writing programs step by step in front of learners) is a better way to teach programming than lectures or self-directed practice; and

- how insights and techniques borrowed from the performing arts can make you a better teacher.

1.1 History

I started teaching people how to program in the late 1980s. At first, I went too fast, used too much jargon, and had little idea of how much my learners actually understood. I got better over time, but still had no idea how effective I was compared to other teachers.

In 2010, I rebooted a project called Software Carpentry[3], whose aim is to teach basic computing skills to researchers. In the years that followed, I discovered resources like Mark Guzdial's blog [Guzdial2017] and the book *How Learning Works* [Ambrose2010]. These led me to [Lemov2014], [Huston2009], [Green2014], and other sources that showed me how to make my teaching better and why I should believe it would work.

We started using these ideas in Software Carpentry[4] in 2012, and the results were everything we'd hoped for. We also started offering a short course to introduce people to these techniques and the ideas behind them. This course was originally delivered online over multiple weeks, but by 2014 we were teaching it in two intensive days (just like our regular software skills workshops). Since then, I have run the course more than forty times for people who want to teach programming to children, recent immigrants, women re-entering the workforce, and a wide variety of other groups. Those experiences are the basis of this book.

1.2 Who You Are

Learner Personas (§6.1) will explain how to define who a class is for. Here, we present personas of two typical participants in a workshop based on this book.

---

[3]http://software-carpentry.org
[4]http://software-carpentry.org

Samira is an undergraduate student in mechanical engineering who first encountered the subject in an after-school club for girls and would now like to pass on her love for it. She has done one programming class and one robotics class, and was a lab assistant for a couple of weekend introductions to engineering for high school students at her university. Feels insecure about standing up and teaching a subject that she isn't an expert in ("I'm not a professor!").

Samira would like to learn techniques for explaining ideas and handling unexpected questions or situations. This workshop will introduce her to some basic classroom practices and give her a chance to try them out in front of a supportive audience.

Moshe is a professional programmer with two young children. Their school doesn't offer a programming class, so he has volunteered to put one together. He has been programming in Visual Basic and C# for almost twenty years, during which time he has frequently given presentations to colleagues and management, but after reading a dozen different "programming for kids" books, he feels more confused than ever about what to do.

Moshe wants to learn how to build lessons that both he and other people can use and maintain. This class will show him how to design and deliver lessons tailored for his students, how to tell how well those lessons are working, and how to keep those lessons up to date.

Moshe is partially deaf, and most of his students have hearing disabilities as well.

1.3 Teaching Practices

We suggest that instructor training workshops use these three teaching practices right from the start:

- Have a code of conduct (§9.1).

- Take notes together (§9.5).

- Pre-assess (§9.6) learners' motivation and prior knowledge.

1.4 Acknowledgments

This book is the product of many contributors, including Erin Becker, Neil Brown, Francis Castro, Warren Code, Karen Cranston, Katie Cunningham, Neal Davis, Mark Degani, Brian Dillingham, Bob Freeman, Mark Guzdial, Rayna Harris, Ian Hawke, Kate Hertweck, Toby Hodges, Christina Koch, Colleen Lewis, Sue McClatchy, Lex

Nederbragt, Jeramia Ory, Elizabeth Patitsas, Aleksandra Pawlik, Emily Porta, Alex Pounds, Danielle Quinn, Erin Robinson, Ariel Rokem, Pat Schloss, Malvika Sharan, Tracy Teal, Richard Tomsett, Matt Turk, Fiona Tweedie, Allegra Via, Anelda van der Walt, Stéfan van der Walt, Belinda Weaver, Hadley Wickham, Jason Williams, and Andromeda Yelton. I am grateful to them, and to everyone who has gone through classes based on this material over the years.

This book is dedicated to my mother, Doris Wilson, who taught hundreds of children to read and to believe in themselves.

1.5 Challenges

Favorite Class (10 minutes)

In the online notes, write down your name, the best class you ever took, and what made it so great.

2   Helping Novices Build Mental Models

Objectives

- *Learners can explain the cognitive differences between novices and competent practitioners in terms of mental models, and explain the implications of these differences for teaching.*
- *Learners can define and differentiate formative and summative assessment.*
- *Learners can construct multiple-choice questions with plausible distractors that have diagnostic power.*

The first task in teaching is to figure out who your learners are and how best to help them. Our approach is based on the Dreyfus model of skill acquisition[1], and more specifically on the work of researchers like Patricia Benner, who studied how nurses progress from being novices to being experts [Benner2000]. Benner identified five stages of cognitive development that most people go through in a fairly consistent way. (We say "most" and "fairly" because human beings are variable, and there will always be outliers. However, that shouldn't prevent us from making strong statements about what's true for the majority.)

For our purposes, we simplify the five stages to three:

1. A *novice* is someone who doesn't know what they don't know, i.e., they don't yet know what the key ideas in the domain are or how they relate. They reason by analogy and guesswork, borrowing bits and pieces of their mental models of other domains which seem superficially similar.

---

[1]https://en.wikipedia.org/wiki/Dreyfus_model_of_skill_acquisition

2. A *competent practitioner* is someone who has a mental model that's good enough for everyday purposes: they can do normal tasks with normal effort under normal circumstances. This model does not have to be completely accurate in order to be useful: for example, the average driver's mental model of how a car works probably doesn't include most of the complexities that a mechanical engineer would be concerned with.

3. An *expert* is someone who can easily handle situations that are out of the ordinary, diagnose the causes of problems, and so on. We will discuss expertise in more detail in Memory (§4).

One sign that someone is a novice is that the things they say are not even wrong[2], e.g., they think there's a difference between programs they type in character by character and identical ones that they have copied and pasted. As we will discuss later (§7), it is very important not to shame novices for this.

One example of a mental model is the ball-and-spring model of molecules that most of us encountered in high school chemistry. Atoms aren't actually balls, and their bonds aren't actually springs, but the model does a good job of helping people reason about chemical compounds and their reactions. Another model of an atom has a small central ball (the nucleus) surrounded by orbiting electrons. Again, this model is wrong, but useful for many purposes.

Novices, competent practitioners, and experts need to be taught differently. In particular, presenting novices with a pile of facts early on is counter-productive, because they don't yet have a model to fit those facts into. In fact, presenting too many facts too soon can actually reinforce the incorrect mental model they've cobbled together. As Derek Muller wrote about this [Muller2011] in the context of video instruction for science students:

> Students have existing ideas about scientific phenomena before viewing a video. If the video presents scientific concepts in a clear, well illustrated way, students believe they are learning but they do not engage with the media on a deep enough level to realize that what was is presented differs from their prior knowledge.

> There is hope, however. Presenting students' common misconceptions in a video alongside the scientific concepts has been shown to increase learning by increasing the amount of mental effort students expend while watching it.

---

[2]https://en.wikipedia.org/wiki/Not_even_wrong

The goal with novices is therefore *to help them construct a working mental model* so that they have somewhere to put facts. As an example of what this means in practice, Software Carpentry's lesson on the Unix shell[3] introduces fifteen commands in three hours. Twelve minutes per command may seem glacially slow, but the lesson's real purpose isn't to teach those fifteen commands: it's to teach learners about paths, history, tab completion, wildcards, pipes and filters, command-line arguments, redirection, and all the other big ideas that the shell depends on. Once they understand those concepts, people can quickly learn a repertoire of commands. What's more, later lessons on how to build functions in a programming language can refer back to pipes and filters, which helps solidify both ideas.

Different Kinds of Lessons

The cognitive differences between novices and competent practitioners underpin the differences between two kinds of teaching materials. A tutorial's purpose is to help newcomers to a field build a mental model; a manual's role, on the other hand, is to help competent practitioners fill in the gaps in their knowledge. Tutorials frustrate competent practitioners because they move too slowly and say things that are obvious (though of course they are anything but to newcomers). Equally, manuals frustrate novices because they use jargon and don't *explain things. One of the reasons Unix and C became popular is that Kernighan et al's trilogy [Kernighan1982], [Kernighan1984], [Kernighan1988] somehow managed to be good tutorials* and *good manuals at the same time. Ray and Ray's book on Unix [Ray2014] and Fehily's introduction to SQL [Fehily2008] are among the very few other books in computing that have accomplished this.*

One of the challenges in building a mental model is to clear away things that *don't* belong. As Mark Twain said, "It ain't what you don't know that gets you into trouble. It's what you know for sure that just ain't so."

Broadly speaking, learners' misconceptions fall into three categories:

1. Simple *factual errors*, such as believing that Vancouver is the capital of British Columbia (it's Victoria). These are simple to correct, but getting the facts right is not enough on its own.

---

[3]http://swcarpentry.github.io/shell-novice/

2. *Broken models*, such as believing that motion and acceleration must be in the same direction. We can address these by having them reason through examples to see contradictions.

3. *Fundamental beliefs*, such as "the world is only a few thousand years old" or "some kinds of people are just naturally better at programming than others" [Patitsas2016]. These are often deeply connected to the learner's social identity, and so are resistant to evidence and cannot be reasoned away in class.

2.1 Formative Assessment

Teaching is most effective when instructors have a way to identify and clear up learners' misconceptions *while they are teaching*. The technical term for this is *formative assessment*, which is assessment that takes place during the lesson in order to form or shape it. Learners don't pass or fail formative assessments; instead, its main purpose is to tell both the instructor and the learner how the learner is doing, and what to focus on next. For example, a music teacher might ask a student to play a scale very slowly in order to see whether she is breathing correctly, and if she is not, what she should change.

The counterpoint to formative assessment is *summative assessment*, which is used at the end of the lesson to tell whether the desired learning took place and whether the learner is ready to move on. One example is a driving exam, which reassures the rest of society that someone can safely be allowed on the road.

When the cook tastes the soup, that's formative. when the guests taste the soup, that's summative.
- Michael Scriven, as quoted by Debra Dirksen.

In order to be useful during teaching, a formative assessment has to be quick to administer and give an unambiguous result. The most widely used kind of formative assessment is probably the multiple choice question (MCQ). When designed well, these can do much more than just tell whether someone knows something or not. For example, suppose we are teaching children multi-digit addition. A well-designed MCQ would be:

Q: what is 37 + 15 ?

(a) 52 (b) 42 (c) 412 (d) 43

The correct answer is 52, but each of the other answers provides valuable insight:

- If the child answers 42, she is throwing away the carry completely.

- If she answers 412, she knows that she can't just discard the carried 1, but doesn't understand that it's actually a ten and needs to be added into the next column. In other words, she is treating each column of numbers as unconnected to its neighbors.

- If she answers 43 then she knows she has to carry the 1, but is carrying it back into the same column it came from.

Each of these incorrect answers is a *plausible distractor* with *diagnostic power*. "Plausible" means that it looks like it could be right, while "diagnostic power" means that each of the distractors helps the instructor figure out what to explain to that particular learner next.

A good MCQ tests for conceptual misunderstanding rather than simple factual knowledge. If you are having a hard time coming up with diagnostic distractors, then either you need to think more about your learners' mental models, or your question simply isn't a good starting point for an MCQ.

Connecting Formative and Summative Assessment

One rule to use when designing lessons is that formative assessments should prepare people for summative assessments: no one should ever encounter a question on an exam for which the teaching did not prepare them. This doesn't mean that novel problems should not appear, but that if they do, learners should have had practice with and feedback on tackling novel problems beforehand.

When you are trying to come up with distractors, think about questions that learners asked or problems they had the last time you taught this subject. If you haven't taught it before, think about your own misconceptions or ask colleagues about their experiences. You can also ask open-ended questions in one class to collect misconceptions about material to be covered in a later class.

Instructors should use MCQs or some other kind of formative assessment every 10-15 minutes in order to make sure that the class is actually learning. That way, if a significant number of people have fallen behind, only a short portion of the lesson will have to be repeated. Additionally, most learners can only focus intensely for roughly this long, so using formative assessments this frequently also helps them re-focus.

Formative assessments can also be used preemptively: if you start a class with an MCQ and everyone can answer it correctly, then you can safely skip the part of the lecture in which you were going to explain something that your learners already know. Doing this also helps show learners that the instructor cares about how much they are learning, and respects their time enough not to waste it.

But what should you do if most of the class votes for one of the wrong answers? What if the votes are evenly spread between options? The answer is, "It depends." If the majority of the class votes for a single wrong answer, you should go back and work on correcting that particular misconception. If answers are pretty evenly split between options, learners are probably guessing randomly and it's a good idea to go back to a point where everyone was on the same page.

If most of the class votes for the right answer, but a few vote for wrong ones, you have to decide whether you should spend time getting the minority caught up, or whether it's more important to keep the majority engaged. This is just one example of one of the most important rules of teaching: no matter how hard you work, or what teaching practices you use, you won't always be able to give everyone the help they need.

Concept Inventories

Given enough data, MCQs can be made surprisingly precise. The best-known example is the Force Concept Inventory[4], which gauges understanding of basic Newtonian mechanics. By interviewing a large number of respondents, correlating their misconceptions with patterns of right and wrong answers to questions, and then improving the questions, its creators constructed a diagnostic tool to pinpoint specific misconceptions. However, it's very costly to do this, and students' ability to search for answers on the internet is an ever-increasing threat to the validity of tools like this.

Designing an MCQ with plausible distractors is useful even if it is never used in class because it forces the instructor to think about the learners' mental models and how they might be broken–in short, to put themselves into the learners' heads and see the topic from their point of view.

---

[4]https://en.wikipedia.org/wiki/Force_Concept_Inventory

2.2 Teaching Practices

If you haven't done so already, you should start using these three teaching practices in your instructor training workshop:

- Use sticky notes as status flags (§9.7) so that you can quickly see who needs help, who has questions, and who's ready to move on.

- Use sticky notes to distribute attention (§9.8) so that everyone gets a fair share of the instructor's time.

- Use sticky notes as minute cards (§9.11) to encourage learners to reflect on what they've just learned and to give instructors actionable feedback while they are still in a position to act on it.

2.3 Challenges

Your Mental Models (5 minutes)

What is one mental model you use to frame and understand your work? Write a few sentences describing it in the shared notes, and give feedback on other learners' contributions.

Symptoms of Being a Novice (5 minutes)

What are the symptoms of being a novice? I.e., what does someone do or say that leads you to classify them as a novice in some domain?

Modelling Novice Mental Models (20 minutes)

Create a multiple choice question related to a topic you intend to teach and explain the diagnostic power of each its distractors (i.e., what misconception each distractor is meant to identify).

When you are done, give your MCQ to a partner, and have a look at theirs. Is the question ambiguous? Are the misconceptions plausible? Do the distractors actually test for them? Are any likely misconceptions *not* tested for?

Other Kinds of Formative Assessment (20 minutes)

A good formative assessment requires people to think through a problem. For example, consider this question from [Epstein2002]. Imagine that you have placed a cake of ice in a bathtub and then filled the tub to the rim with water. When the ice melts, does the water level go up (so that the tub overflows), go down, or stay the same?

The correct answer is that the level stays the same: the ice displaces its own weight in water, so it exactly fills the "hole" it has made when it melts. Figuring this out why helps people build a model of the relationship between weight, volume, and density.

Describe another kind of formative assessment you have seen or used and explain how it helps both the instructor and the learner figure out where they are and what they need to do next.

A Different Progression (15 minutes)

Another progression often used to describe the path from novice to expert is the four stages of competence[5]:

- Unconscious incompetence: the person doesn't know what they don't know.
- Conscious incompetence: the person realizes that they don't know something.
- Conscious competence: the person has learned how to do something, but can only do it while concentrating, and may still need to break things down into steps.
- Unconscious competence: the skill has become second nature, and the person can do it reflexively.

Describe one or more subjects related to programming for which you are at each of these levels.

Where You Are (5 minutes)

Identify one field in which you are a novice (but hoping to improve), one in which you feel you are competent, and one where you feel you are an expert, and share with the rest of the group.

---

[5]https://en.wikipedia.org/wiki/Four_stages_of_competence

3   Teaching as a Performance Art

Objectives

- *Learners can define* jugyokenkyu *and lateral knowledge transfer and explain their relationship to each other.*
- *Learners can describe and enact at least three techniques for giving and receiving feedback on teaching performance.*
- *Learners can explain at least two ways in which using a rubric makes feedback more effective.*

Many people assume that teachers are born, not made. From politicians to researchers and teachers themselves, reformers have designed systems to find and promote those who can teach and eliminate those who can't. But as Elizabeth Green explains [Green2014], that assumption is wrong, which is why educational reforms based on it have repeatedly failed.

The book is written as a history of the people who have put that puzzle together in the US. Its core begins with a discussion of what James Stigler discovered during a visit to Japan in the early 1990s:

Some American teachers called their pattern "I, We, You": After checking homework, teachers announced the day's topic, demonstrating a new procedure (I)... Then they led the class in trying out a sample problem together (We)... Finally, they let students work through similar problems on their own, usually by silently making their way through a worksheet (You)...

The Japanese teachers, meanwhile, turned "I, We, You" inside out. You might call their version "You, Y'all, We." They began not with an introduction, but a single problem that students spent ten or twenty minutes working through alone (You)... While the students worked, the teacher wove through the students' desks, studying what they came up with and taking notes to remember who had which idea. Sometimes the teacher then deployed the

students to discuss the problem in small groups (Y'all). Next, the teacher brought them back to the whole group, asking students to present their different ideas for how to solve the problem on the chalkboard. . . Finally, the teacher led a discussion, guiding students to a shared conclusion (We).

It's tempting but wrong to think that this particular teaching technique is some kind of secret sauce. The actual key is a practice called *jugyokenkyu*, which means "lesson study":

Jugyokenkyu is a bucket of practices that Japanese teachers use to hone their craft, from observing each other at work to discussing the lesson afterward to studying curriculum materials with colleagues. The practice is so pervasive in Japanese schools that it is. . . effectively invisible.

In order to graduate, [Japanese] education majors not only had to watch their assigned master teacher work, they had to effectively replace him, installing themselves in his classroom first as observers and then, by the third week, as a wobbly. . . approximation of the teacher himself. It worked like a kind of teaching relay. Each trainee took a subject, planning five days' worth of lessons. . . [and then] each took a day. To pass the baton, you had to teach a day's lesson in every single subject: the one you planned and the four you did not. . . and you had to do it right under your master teacher's nose. Afterward, everyone–the teacher, the college students, and sometimes even another outside observer–would sit around a formal table to talk about what they saw.

Putting work under a microscope in order to improve it is commonplace in sports and music. A professional musician, for example, will dissect half a dozen different recordings of "Body and Soul" or "Smells Like Teen Spirit" before performing it. They would also expect to get feedback from fellow musicians during practice and after performances. Many other disciplines work this way too: the Japanese drew inspiration from Deming's ideas on continuous improvement in manufacturing[1], while the adoption of code review over the last 15 years has done more to improve everyday programming than any number of books or websites.

But this kind of feedback isn't part of teaching culture in the US, the UK, Canada, or Australia. There, what happens in the classroom stays in the classroom: teachers don't watch each other's lessons on a regular basis, so they can't borrow each other's good ideas. The

---

[1]https://en.wikipedia.org/wiki/W._Edwards_Deming

result is that *every teacher has to invent teaching on their own*. They may get lesson plans and assignments from colleagues, the school board, a textbook publisher, or the Internet, but each teacher has to figure out on their own how to combine that with the theory they've learned in education school to deliver an actual lesson in an actual classroom for actual students.

Demonstration lessons, in which one teacher is in front of a room full of students while other teachers observe, seem like a way to solve this. However, Fincher and her colleagues studied how teaching practices are actually transferred using both a detailed case study [Fincher2007] and analysis of change stories [Fincher2012]. The abstract of the latter paper sums up their findings:

> *Innovative tools and teaching practices often fail to be adopted by educators in the field, despite evidence of their effectiveness. Naïve models of educational change assume this lack of adoption arises from failure to properly disseminate promising work, but evidence suggests that dissemination via publication is simply not effective... We asked educators to describe changes they had made to their teaching practice... Of the 99 change stories analyzed, only three demonstrate an active search for new practices or materials on the part of teachers, and published materials were consulted in just eight... Most of the changes occurred locally, without input from outside sources, or involved only personal interaction with other educators.*

Barker et al found something similar [Barker2015]:

> *Adoption is not a "rational action," however, but an iterative series of decisions made in a social context, relying on normative traditions, social cueing, and emotional or intuitive processes... Faculty are not likely to use educational research findings as the basis for adoption decisions... Positive student feedback is taken as strong evidence by faculty that they should continue a practice.*

This phenomenon is sometimes called *lateral knowledge transfer*: someone sets out to teach X, but while watching them, their audience actually learns Y as well (or instead). For example, an instructor might set out to show people how to do a particular statistical analysis in R, but what her learners might take away is some new keyboard shortcuts in RStudio. Live coding makes this much more likely because it allows learners to see the "how" as well as the "what", and *jugyokenkyu* works because it creates more opportunities for this to happen.

3.1 Feedback

As the cartoon below suggests, sometimes it can be hard to receive feedback, especially negative feedback. The process is easier and more productive when the people involved share ground rules and expectations. This is especially important when they have different backgrounds or cultural expectations about what's appropriate to say and what isn't.
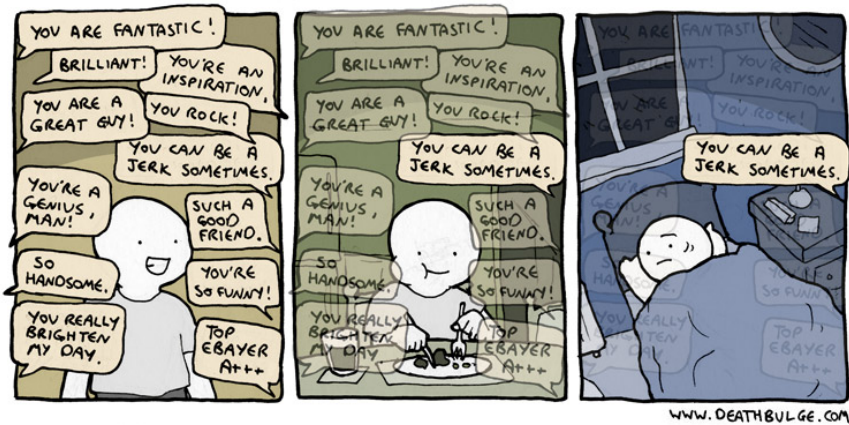
Figure 3.1: Feedback Feelings (copyright (c) Deathbulge 2013)

You can get better feedback on your work from other people using techniques like these:

1. *Initiate feedback.* It's better to ask for feedback than to receive it unwillingly.

2. *Choose your own questions*, i.e., ask for specific feedback. It's a lot harder for someone to answer, "What do you think?" than to answer either, "What is one thing I could have done as an instructor to make this lesson more effective?" or "If you could pick one thing from the lesson to go over again, what would it be?"

   Directing feedback like this is also more helpful to you. It's always better to try to fix one thing at once than to change everything and hope it's for the better. Directing feedback at something you have chosen to work on helps you stay focused, which in turn increases the odds that you'll see progress.

3. *Use a feedback translator.* Have a fellow instructor (or other trusted person in the room) read over all the feedback and give an executive summary. It can be easier to hear "It sounds like most people are following, so you could speed up" than to read several notes all saying, "this is too slow" or "this is boring".

4. Most importantly, *be kind to yourself.* Many of us are very critical of ourselves, so it's always helpful to jot down what we thought of ourselves *before* getting feedback from others. That allows us to compare what we think of our performance with what others think, which in turn allows us to scale the former more accurately. For example, it's very common for people to think that they're saying "um" and "err" all the time, when their audience doesn't notice it. Getting that feedback once allows instructors to adjust their assessment of themselves the next time they feel that way.

You can give feedback to others more effectively as well:

1. *Balance positive and negative feedback.* One method is a "compliment sandwich" made up of one positive, one negative, and a second positive observation.

2. *Organize your feedback using a rubric.* Most people are more comfortable giving and receiving feedback when they feel that they understand the social rules governing what they are allowed to say and how they are allowed to say it. A facilitator can then transcribe items into a shared document (or onto a whiteboard) during discussion.

Two by Two

The rubric we find most useful for feedback on teaching is a 2x2 grid whose vertical axis is labelled "positive" and "negative", and whose horizontal axis is labelled "content" (what was said) and "presentation" (how it was said). Observers write each of their comments in one of the grid's four squares as they are watching the demonstration.

Whatever methods are used, the most important thing to remember is feedback on teaching is meant to be formative: its goal is to help people figure out what they are doing well and what they still need to work on.

Studio Classes

Architecture schools often include studio classes, in which students solve small design problems and get feedback from their peers right then and there. These classes are most effective when the instructor critiques both the designs and the peer critiques, so that participants are learning not only how to make buildings, but how to give and get feedback [Schön1984]. Master classes in music serve a similar purpose, and a few people have experimented with using live coding at conferences or online in similar ways.

Tells

Everyone has nervous habits. For example, many of us become "Mickey Mouse" versions of ourselves when we're nervous, i.e., we talk more rapidly than usual, in a higher-pitched voice, and wave our arms around more than we usually would.

Gamblers call nervous habits like this "tells". While these are often not as noticeable as you would think, it's good to know whether you pace, fiddle with your hair, look at your shoes, or rattle the change in your pocket when you don't know the answer to a question.

You can't get rid of tells completely, and trying to do so can make you obsess about them. A better strategy is to try to displace them, e.g., to train yourself to scrunch your toes inside your shoes instead of cracking your knuckles.

If you are interested in knowing more about giving and getting feedback, you may want to read [Gormally2014] and discuss ways you could make peer-to-peer feedback a routine part of your teaching. You may also enjoy [Gawande2011], which looks at the value of having a coach.

3.2   How to Practice Teaching

One of the key elements of instructor training is recording trainees and having them, and their peers, critique those recordings. We were introduced to this practice by UBC's Warren Code, who got it from the Instructional Skills Workshop [ISW2017], and it has evolved to the following:

1. Split into groups of three.

2. Each person rotates through the roles of instructor, audience, and videographer. As the instructor, they have two minutes to

explain one key idea from their research (or other work) as if they were talking to a class of interested high school students. The person pretending to be the audience is there to be attentive, while the videographer records the session using a cellphone or similar device.

3. After everyone in the group of three has finished teaching, watch the videos as a group. Everyone gives feedback on all three videos, i.e., people give feedback on themselves as well as on others.

4. After everyone has given feedback on all of the videos, return to the main group and put all of the feedback into the notes. Again, try to divide positive from negative and content from presentation. Try also to identify each person's tells: what do they do that betrays nervousness, and how noticeable is it?

It's important to record all three videos and then watch all three: if the cycle is teach-review-teach-review, the last person to teach runs out of time. Doing all the reviewing after all the teaching also helps put a bit of distance between the teaching and the reviewing, which makes the exercise slightly less excruciating.

In order for this exercise to work well:

- Let people know at the start of the class that they will be asked to teach something so that they have time to choose a topic. (In our experience, telling them this in advance of the class can be counter-productive, since some people will fret over how much they should prepare.)

- Groups must be physically separated to reduce audio cross-talk between their recordings. In practice, this means 2-3 groups in a normal-sized classroom, with the rest using nearby breakout spaces, coffee lounges, offices, or (on one occasion) a janitor's storage closet.

- People must give feedback on themselves, as well as giving feedback on each other, so that they can calibrate their impressions of their own teaching according to the impressions of other people. (We find that most people are harder on themselves than others are, and it's important for them to realize this.)

- Try to make at least one mistake during the demonstration of live coding so that trainees can see you talk through diagnosis and recovery, and draw attention afterward to the fact that you did this.

19

The announcement of this exercise is often greeted with groans and apprehension, since few people enjoy seeing or hearing themselves. However, it is consistently rated as one of the most valuable parts of the class, and also serves as an ice breaker: we want pairs of instructors at actual workshops to give one another feedback, and that's much easier to do once they've had some practice and have a rubric to follow.

Setting Up Your Teaching Environment

If the room setup allows it, try to set up your environment (§9.19) to mimic what you would use in an actual classroom: have a glass of water handy, stand instead of sitting, and so on.

3.3 Challenges

Give Feedback (20 minutes)

1. Watch [Wilson2016] as a group and give feedback on it. Organize feedback along two axes: positive vs. negative and content vs. presentation.

2. Have each person in the class add one point to a 2x2 grid on a whiteboard (or in the shared notes) without duplicating any points that are already up there.

What did other people see that you missed? What did they think that you strongly agree or disagree with?

Practice Giving Feedback (45 minutes)

Use the process described above to practice teaching in groups of three. When your group is done, the instructor will add one point of feedback from each participant to a 2x2 grid on the whiteboard or in the shared notes, without accepting duplicates. Participants should not say whether the point they offer was made by them, about them, or neither: the goal at this stage is primarily for people to become comfortable with giving and receiving feedback, and to establish a consensus about what sorts of things to look for.

4   Expertise and Memory

Objectives

- *Learners can define expertise and explain its operation using a graph metaphor for cognition.*
- *Learners can explain the difference between repetition and deliberate practice.*
- *Learners can define and construct concept maps, and explain the benefits of externalizing cognition.*
- *Learners can differentiate long-term and short-term memory, describe the capacity limits of the latter, and explain the the impact of these limits on teaching.*

The previous chapter looked at what distinguishes novices from competent practitioners. Here, we will look at expertise: what it is, how people acquire it, and how it can be harmful as well as helpful. We will then see how concept maps can be used to figure out how to turn knowledge into lessons.

To start, what do we mean when we say someone is an expert? The usual response is that they can solve problems much faster than people who are "merely competent", or that they can recognize and deal with the cases where the normal rules don't apply. They also somehow make this look effortless: in most cases, they just know what the right answer is.

What makes someone an expert? The answer isn't just that they know more facts: competent practitioners can memorize a lot of trivia without any noticeable improvement to their performance. Instead, imagine for a moment that we store knowledge as a graph in which facts are nodes and relationships are arcs. (This is emphatically *not* how our brains work, but it's a useful metaphor.) The key difference between experts and people who are "merely competent" is that experts have many more connections, i.e., their mental models are much more densely connected.

This metaphor helps explain many observed aspects of expert behavior:

- Experts can jump directly from a problem to its solution because there actually is a direct link between the two in their mind. Where a competent practitioner would have to reason "A, B, C, D, E", the expert can go from A to E in a single step. We call this *intuition*, and it isn't always a good thing: when asked to explain their reasoning, experts often can't, because they didn't actually reason their way to the solution–they just recognized it.

- Experts are frequently so familiar with their subject that they can no longer imagine what it's like to *not* see the world that way. As a result, they are often less good at teaching the subject than people with less expertise who still remember what it's like to have to learn the things. This phenomenon is called *expert blind spot*, and while it can be overcome with training, it's part of why there is no correlation between how good someone is at doing research in an area and how good they are at teaching it [Marsh2002].

- Densely-connected knowledge graphs are also the basis for experts' *fluid representations*, i.e., their ability to switch back and forth between different views of a problem [Petre2016]. For example, when trying to solve a problem in mathematics, we might switch between tackling it geometrically and representing it as a set of equations to be solved.

- Finally, this metaphor also explains why experts are better at diagnosis than competent practitioners: more linkages between facts makes it easier to reason backward from symptoms to causes. (And this in turn is why asking programmers to debug during job interviews gives a more accurate impression of their ability than asking them to program.)

The J Word

Experts often betray their blind spot by using the word "just" in explanations, as in, "Oh, it's easy, you just fire up a new virtual machine and then you just install these four patches to Ubuntu and then you just re-write your entire program in a pure functional language." As we discuss later in Motivation (§7), the J word (also sometimes called the passive dismissive adjective) should be banned from classrooms, primarily because using it gives learners the very clear signal that the instructor thinks their problem is trivial and that they therefore must be stupid.

The graph model of knowledge explains why helping learners make connections is as important as introducing them to facts. To use another analogy, the more people you know in a group, the more likely you are to remain part of that group. Similarly, the more connections a fact has to other facts, the more likely the fact is to be remembered.

4.1   Repetition vs. Deliberate Practice

The idea that ten thousand hours of practice will make someone an expert in some field is widely quoted, but reality is more complex. Doing exactly the same thing over and over again is much more likely to solidify bad habits than perfect performance. What actually works is *deliberate practice* (also sometimes called *reflective practice*), which is doing similar but subtly different things, paying attention to what works and what doesn't, and then changing behavior in response to that feedback to get cumulatively better.

A common progression is for people to go through three stages:

1. They *learn how to do something given feedback from others*. For example, they might write an essay about what they did on their summer holiday, and get feedback from a teacher telling them how to improve it.

2. They *learn how to give feedback*. For example, they might write an essay about character development in *The Catcher in the Rye*, and get feedback on their critique from a teacher.

3. They *apply what they've learned about feedback to themselves*. At some point, they start critiquing their own work in real time (or nearly so) using the critical skills they have now built up. Doing this is so much faster than waiting for feedback from others that proficiency suddenly starts to take off.

A meta-study conducted in 2014 [Macnamara2014] found that "...deliberate practice explained 26% of the variance in performance for games, 21% for music, 18% for sports, 4% for education, and less than 1% for professions." One explanation for this variation is that deliberate practice works best when the rules for evaluating success are very stable, but is less effective when there are more factors at play (i.e., when it's harder to connect cause to effect).

4.2   Concept Maps

Our tool of choice to represent a knowledge graph (expert or otherwise) is a *concept map*. A concept map is simply a picture of someone's mental model of a domain: facts are bubbles, and connections are labelled arcs. It is important that they are labelled: saying "X and Y are related" is only helpful if we explain what the relationship *is*. (And yes, one person's fact may be another person's connection, but one of the benefits of concept mapping is that it makes those differences explicit.) The figure below reproduces a concept map taken from the IHMC CMap site[1] showing why the Earth has seasons.
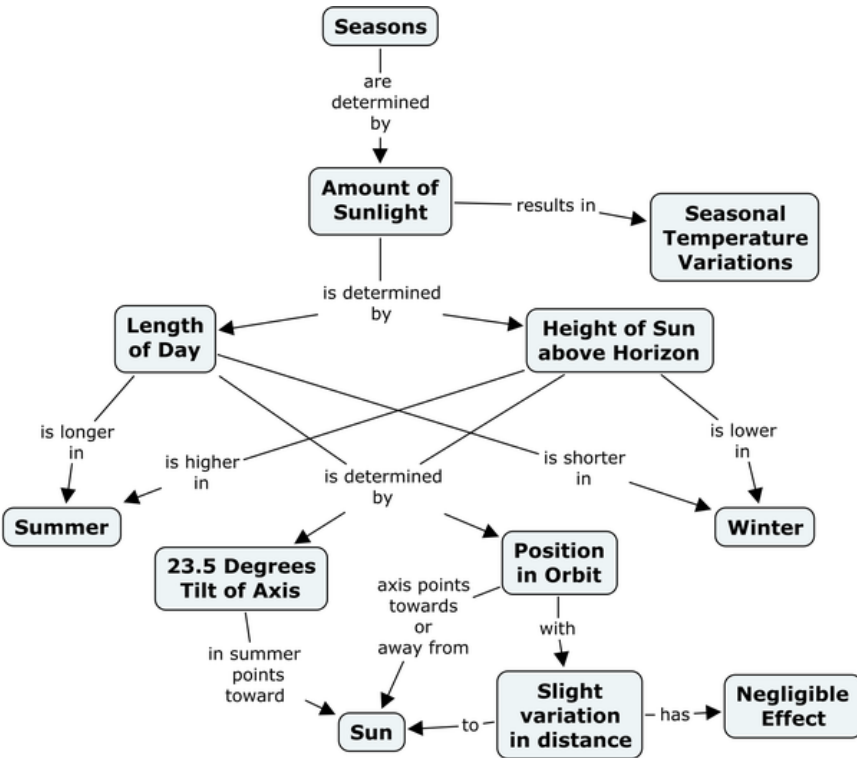
Figure 4.1: Concept Map for Seasons

To show what concept maps look like, consider this simple `for` loop in Python:

```
for letter in "abc":
      print('*' + letter)
```

whose output is:

```
*a
*b
*c
```

The three key "things" in this loop are shown in the first part of this concept map:
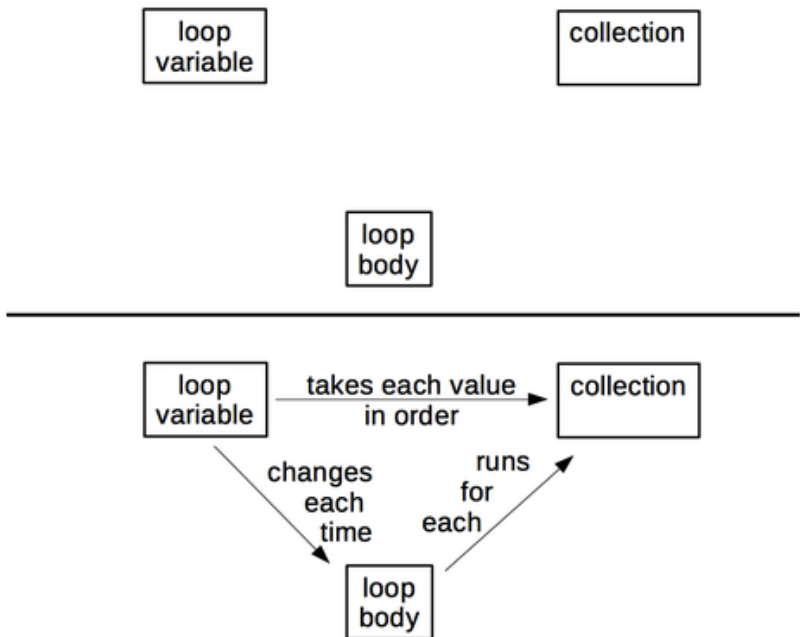
Figure 4.2: Concept Map

This is only half the story, though, and arguably the less important half. An expanded concept map shows the *relationships* between those things, which are as important as the concepts themselves.

Concept maps can be used in many ways:

1. Concept maps aid design of a lesson by helping authors figure out what they're trying to teach. Crucially, a concept map separates content from order: in our experience, people rarely wind up teaching things in the order in which they first drew them.

2. They also aid communication between lesson designers. Instructors with very different ideas of what they're trying to teach are likely to pull their learners in different directions. Drawing and sharing concept maps isn't guaranteed to prevent this, but it certainly helps.

3. Concept maps also aid communication with learners. While it's possible to give learners a pre-drawn map at the start of a lesson for them to annotate, it's better to draw it piece by piece while teaching to reinforce the ties between what's in the map and what the instructor said. (We will return to this idea when we discuss Mayer's work on multimedia learning in Cognitive Load (§5)).

4. Concept maps are also a useful for assessment: having learners draw concept maps of what they think they just heard shows the instructor what was missed and what was mis-understood. However, reviewing learners' concept maps is too time-consuming for use in class, but very useful in weekly lectures *once learners are familiar with the technique.* The qualification is necessary because any new way of doing things initially slows people down–if a student is trying to make sense of basic programming, asking them to figure out how to draw their thoughts at the same time is an unfair load. Finally, some instructors are skeptical of whether novices can effectively map their understanding, since introspection and explanation of understanding are generally more advanced skills than understanding itself.

Externalizing Cognition

Concept maps are just one way to represent our understanding of a subject. For example, Andrew Abela's decision tree [Abela2009] presents a mental model of how to choose the right kind of chart for different kinds of questions and data. Maps, flowcharts, and blueprints can also be useful in some contexts. What each does is externalize cognition, *i.e., make thought processes and mental models visible so that they can be compared, contrasted, and combined.*

Meetings, Meetings, Meetings

The next time you have a team meeting, give everyone a sheet of paper and have them spend a few minutes drawing a concept map of the project you're all working on–separately. On the count of three, have everyone reveal their concept maps simultaneously. The discussion that follows everyone's realization of how different their mental models of the project's aims and organization are is always interesting. . .

4.3   Seven Plus or Minus Two

The graph model of knowledge is wrong but useful, but another simple model has a sound physical basis. As a rough approximation, human memory can be divided into two distinct layers. The first is called *long-term* or *persistent memory*. It is where we store things like our password, our home address, and what the clown did at our eighth birthday party that scared us so much. It is essentially unbounded: barring injury or disease, we will die before it fills up. However, it is also slow to access–too slow to help us handle hungry lions and disgruntled family members.

Evolution has therefore given us a second system called *short-term* or *working memory*. It is much faster, but also much smaller: in 1956, Miller estimated that the average adult's working memory could hold 7±2 items for a few seconds before things started to drop out. This is why phone numbers are typically 7 or 8 digits long[2]: back when phones had dials instead of keypads, that was the longest string of numbers most adults could remember accurately for as long as it took the dial to go around and around. It's also why sports teams tend to have about half a dozen members, or be broken down into smaller groups (such as the forwards and backs in rugby).

Chunking

Our minds can store larger numbers of facts in short-term memory by creating chunks. *For example, most of us will remember a word we read as a single item, rather than as a sequence of letters. Similarly, the pattern made by five spots on cards or dice is remembered as a whole rather than as five separate pieces of information.*

Chunks allow us to manage larger problems, but can also mislead us if we mis-identify something, i.e., see it as something it isn't.

---

[2]https://www.quora.com/Why-did-Bell-Labs-create-phone-numbers-of-7-digits-10-digits-Is-there-a-reason-that-dashes-and-brackets-are-used

We will discuss this in more detail later (§5.1). Recent estimates have suggested that the size of short-term memory might be as low as 4±1 [Didau2016] which means that effective chunking (discussed below) is even more important than first thought.

7±2 is probably the most important number in programming. When someone is trying to write the next line of a program, or understand what's already there, she needs to keep a bunch of arbitrary facts straight in her head: what does this variable represent, what value does it currently hold, etc. If the number of facts grows too large, her mental model of the program comes crashing down (something we have all experienced).

7±2 is also the most important number in teaching. An instructor cannot push information directly into a learner's long-term memory. Instead, whatever she presents is first represented in the learner's short-term memory, and is only transferred to long-term memory after it has been held there and rehearsed. If we present too much information too quickly, the new will displace the old before it has a chance to consolidate in long-term memory.

This is one of the reasons to create a concept map for a lesson before teaching it: an instructor needs to identify how many pieces of separate information will need to be "stored" in memory as part of the lesson. In practice, it's common to draw a concept map, realize that there's far too much in it to teach in a single pass, and then carve out tightly-connected sub-sections to divide the overall lesson into teachable episodes.

Building Concept Maps Together

Concept maps can be used as a classroom discussion exercise. Put learners in small groups (2-4 people each), give each group some sticky notes on which a few key concepts are written, and have them build a concept map on a whiteboard by placing those sticky notes, connecting them with labelled arcs, and adding any other concepts they think they need.

4.4 Six Strategies

The Learning Scientists[3] have catalogued six strategies for effective learning, which are summarized in these downloadable posters[4]:

1. *Spaced practice*: creating a study schedule that spreads study activities over time. Block off 30 minutes to study each day

---

[3]http://www.learningscientists.org/
[4]http://www.learningscientists.org/downloadable-materials

rather than trying to cram everything in the night before an exam.

2. *Retrieval practice*: bringing learned information to mind from long-term memory. Take practice tests, or practice summarizing the details of a topic from memory.

3. *Interleaving*: creating a study schedule that mixes different kinds of problems, or different content subjects within one study session. Solve practice problems A-B-C-A-B-C-A-B-C rather than A-A-A-B-B-B-C-C-C.

4. *Elaboration*: asking and explaining why things work. Follow up each practice answer with an explanation of why that answer is correct.

5. *Concrete examples*: illustrating abstract concepts with specific real-life examples. Follow each statement of a general principle with one or more examples, or each particular problem with a statement of the general principles it embodies.

6. *Dual coding*: combining words with visuals. Draw or label diagrams while explaining things or working through problems.

Teaching these strategies to students, and mentioning them by name when you use them in class, can help them learn how to learn faster and better.

4.5 Challenges

Concept Mapping (30 minutes)

Create a hand drawn concept map for something you would teach in five minutes. (If possible, do it for the same subject that created a multiple choice question for earlier.) Trade with a partner, and critique each other's maps. Do they present concepts or surface detail? Which of the relationships in your partner's map do you consider concepts and vice versa?

Learning Strategies (20 minutes)

1. Which of the learning strategies discussed above do you regularly use? Which ones do you not?
2. Write down three general concepts that you want your learners to master, and then give two specific examples of each. (This uses practice #5).

3. For each of those concepts, work backward from one of your examples to explain how the concept explains it. (This uses practice #4).

The Cost of Multi-Tasking (10 minutes)

This post[5] from the Learning Scientists describes a simple experiment you can do without preparation or equipment (other than a timer) to demonstrate the mental cost of multi-tasking. Try the activity in pairs, and then report your results to the group.

---

[5]http://www.learningscientists.org/blog/2017/7/28-1

5 Cognitive Load

Objectives

- *Learners can define cognitive load and explain how consideration of it can be used to shape instruction.*
- *Learners can explain what faded examples are and construct faded examples for use in programming workshops.*
- *Learners can explain what Parsons Problems are and construct Parsons Problems for use in programming workshops.*
- *Learners can describe ways they differ from their own students and what effect those differences have on instruction.*

In 2006, Kirschner, Sweller, and Clark published a paper titled "Why Minimal Guidance During Instruction Does Not Work: An Analysis of the Failure of Constructivist, Discovery, Problem-Based, Experiential, and Inquiry-Based Teaching" [Kirschner2006]. Its abstract says:

> *Although unguided or minimally guided instructional approaches are very popular and intuitively appealing...these approaches ignore both the structures that constitute human cognitive architecture and evidence from empirical studies over the past half-century that consistently indicate that minimally guided instruction is less effective and less efficient than instructional approaches that place a strong emphasis on guidance of the student learning process. The advantage of guidance begins to recede only when learners have sufficiently high prior knowledge to provide "internal" guidance.*

The paper set off a minor academic firestorm, because beneath the jargon the authors were claiming that *inquiry-based learning* doesn't actually work very well. Inquiry-based learning is the practice of allowing learners to ask their own questions, set their own goals, and find their own path through a subject, just as they would

when solving problems in real life. It is intuitively appealing, but Kirschner argued that it overloads learners, since it requires them to simultaneously master both a domain's factual content and its problem-solving strategies.

More specifically, *cognitive load theory* posits that people have to deal with three things when they're learning:

1. *Intrinsic* load is what people have to keep in mind in order to carry out a learning task. In a programming class, this might be understanding what a variable is, or understanding how assignment in a programming language is different from creating a reference to a cell in a spreadsheet.

2. *Germane* load is the (desirable) mental effort required to create linkages between new information and old, which is one of the things that distinguishes learning from memorization. An example might be learning how to loop through a collection in Python.

3. *Extraneous* load is everything else that distracts or gets in the way, such as knowing that tabs look like multiple characters but only count as one when indenting Python code.

According to this theory, searching for a solution strategy is an extra burden on top of applying that strategy. We can therefore accelerate learning by giving learners worked examples that show them a problem and a detailed step-by-step solution, followed by a series of *faded examples*. The first of these presents a nearly-complete use of the same problem-solving strategy just demonstrated, but with a small number of blanks for the learner to fill in. The next problem is also of the same type, but has more blanks, and so on until the learner is asked to solve the entire problem. (The material that *isn't* blank is often referred to as *scaffolding*, since it serves the same purpose as the scaffolding set up temporarily at a building site.)

For example, someone teaching Python might start by explaining this:

```
# total_length(["red", "green", "blue"]) => 12
def total_length(words):
    total = 0
    for word in words:
        total += len(word)
    return total
```

then ask learners to fill in the blanks in:

```
# word_lengths(["red", "green", "blue"]) => [3, 5, 4]
def word_lengths(words):
    lengths = ____
    for word in words:
        lengths ____
    return lengths
```

The next problem might be:

```
# join_all(["red", "green", "blue"]) => "redgreenblue"
def join_all(words):
    result = ____
    for ____ in ____:
        ____
    return result
```

and learners would finally be asked to tackle:

```
# acronymize(["red", "green", "blue"]) => "RGB"
def acronymize(words):
    ____
```

Faded examples work because they introduce the problem-solving strategy piece by piece. At each step, learners have one new problem to tackle. As discussed later (§9.4), this is less intimidating than a blank screen or a blank sheet of paper. It also encourages learners to think about the similarities and differences between various approaches, which helps create the linkages in the mental model that instructors want them to form.

The key to constructing a good faded example is to think about the problem-solving strategy or solution pattern that it is meant to teach. For example, the series of problems are all examples of the *accumulator pattern*, in which the results of processing items from a collection are repeatedly added to a single variable in some way to create the final result.

Cognitive load theory has been criticized as being unfalsifiable[1]: since there's no way to tell in advance of an experiment whether something is germane or not, any result can be justified after the fact by labelling things that hurt performance as "extraneous" and things that don't "germane". However, there is no doubt that faded examples are effective.

---

[1] https://edtechdev.wordpress.com/2009/11/16/cognitive-load-theory-failure/

Split Attention

Research by Mayer and colleagues on the split-attention effect[2] is closely related to cognitive load theory [Mayer2003]. Linguistic and visual input are processed by different parts of the human brain, and linguistic and visual memories are stored separately as well. This means that correlating linguistic, auditory, and visual streams of information takes cognitive effort: when someone reads something while hearing it spoken aloud, their brain can't help but check that it's getting the same information on both channels.

Learning is therefore more effective when redundant information is not presented simultaneously in two different channels. For example, people find it harder to learn from a video that has both narration and on-screen captions than from one that has either the narration or the captions but not both.

The key word in the previous paragraph is "redundant". It turns out that it's more effective to draw a diagram piece by piece while teaching rather than to present the whole thing at once. If parts of the diagram appear at the same time as things are being said, the two will be correlated in the learner's memory. Pointing at part of the diagram later is then more likely to trigger recall of what was being said when that part was being drawn.

Another way to use cognitive load theory to construct exercises is called a *Parsons Problem*. If you are teaching someone to speak a new language, you could ask them a question, and then give them the words they need to answer the question, but in jumbled order. Their task is to put the words in the right order to answer the question grammatically, which frees them from having to think simultaneously about what to say *and* how to say it.

Similarly, when teaching people to program, you can give them the lines of code they need to solve a problem, and ask them to put them in the right order. This allows them to concentrate on control flow and data dependencies, i.e., on what has to happen before what, without being distracted by variable naming or trying to remember what functions to call. Multiple studies have shown that Parsons Problems take less time for learners to do, but produce equivalent educational outcomes [Ericons2017].

---

[2]https://en.wikipedia.org/wiki/Split_attention_effect

5.1 Pattern Recognition

An earlier section (§4.3) described how people chunk related or correlated information together so that they can fit more into short-term memory. One key finding in cognition research is that experts have more and larger chunks than non-experts, i.e., experts "see" larger patterns, and have more patterns to match things against. This allows them to reason at a higher level, and to search for information more quickly and more accurately.

It is therefore tempting to try to teach patterns directly–in fact, supporting this is one of the reasons programmers have been so enthusiastic about design patterns[3]. In practice, though, pattern catalogs are too large to flick through and too dry to memorize directly. Giving names to a small number of patterns, though, does seem to help with teaching, primarily by giving the learners a richer vocabulary to think and communicate with [Kuittinen2004].

5.2 You Are Not Your Learners

People learn best when they care about the topic and believe they can master it. Neither fact is particularly surprising, but their practical implications have a lot of impact on what we teach, and the order in which we teach it.

First, as noted in Motivation (§7), most people don't actually want to program: they want to build a website or check on zoning regulations, and programming is just a tax they have to pay along the way. They don't care how hash tables work, or even that hash tables exist; they just want to know how to process data faster. We therefore have to make sure that everything we teach is useful right away, and conversely that we don't teach anything just because it's "fundamental".

Second, believing that something will be hard to learn is a self-fulfilling prophecy. This is why it's important not to say that something is easy: if someone who has been told that tries it, and it doesn't work, they are more likely to become discouraged.

It's also why installing and configuring software is a much bigger problem for us than experienced programmers like to acknowledge. It isn't just the time we lose at the start of boot camps as we try to get a Unix shell working on Windows, or set up a version control client on some idiosyncratic Linux distribution.

It isn't even the unfairness of asking students to debug things that depend on precisely the knowledge they have come to learn,

---

[3]https://en.wikipedia.org/wiki/Software_design_pattern

but which they don't yet have. The real problem is that every such failure reinforces the belief that computing is hard, and that they'd have a better chance of making next Thursday's deadline at work if they kept doing things the way they always have. For these reasons, we have adopted a "teach most immediately useful first" approach described in Motivation (§7).

5.3   Challenges

Create a Faded Example (30 minutes)

It's very common for programs to count how many things fall into different categories: for example, how many times different colors appear in an image, or how many times different words appear in a paragraph of text.

1. Create a short example (no more than 10 lines of code) that shows people how to do this, and then create a second example that solves a similar problem in a similar way, but has a couple of blanks for learners to fill in. How did you decide what to fade out? What would the next example in the series be?

2. Define the audience for your examples. For example, are these beginners who only know some basics programming concepts? Or are these learners with some experience in programming but not in Python?

3. Show your example to a partner, but do *not* tell them what level it is intended for. Once they have filled in the blanks, ask them what level they think it is for.

(If there are people among the trainees who don't program at all, try to place them in different groups, and have them play the part of learners for those groups.)

Create a Parsons Problem (20 minutes)

Write five or six lines of code that does something useful, jumble them, and ask your partner to put them in order. If you are using an indentation-based language like Python, do not indent any of the lines; if you are using a curly-brace language like Java, do not include any of the curly braces.

6 Designing Lessons

Objectives

- *Learners can describe the steps in reverse instructional design and explain why it generally produces better lessons than the usual "forward" lesson development process.*
- *Learners can define "teaching to the test" and explain why reverse instructional design is* not *the same thing.*
- *Learners can construct and critique five-part learner personas.*
- *Learners can construct good learning objectives and critique learning objectives with reference to Bloom's Taxonomy and/or Fink's Taxonomy.*

Most people design lessons as follows:

1. Someone tells you that you have to teach something you haven't thought about in ten years.

2. You start writing slides to explain what you know about the subject.

3. After two or three weeks, you make up an assignment based more or less on what you've taught so far.

4. You repeat step 3 several times.

5. You stay awake into the wee hours of the morning to create a final exam.

There's a better way, but to explain it, we first need to explain how *test-driven development* (TDD) is used in software development. Programmers who are using TDD don't write software and then (possibly) write tests. Instead, they write the tests first, then write just enough new software to make those tests pass, and then clean up a bit.

TDD works because writing tests forces programmers to specify exactly what they're trying to accomplish and what "done" looks like. It's easy to be vague when using a human language like English or Korean; it's much harder to be vague in Python or R.

TDD also reduces the risk of endless polishing, and also the risk of confirmation bias: someone who hasn't written a program is much more likely to be objective when testing it than its original author, and someone who hasn't written a program *yet* is more likely to test it objectively than someone who has just put in several hours of hard work and really, really wants to be done.

A similar "backward" method works very well for lesson design. This method is something called *reverse instructional design* and was developed independent in [Wiggins2005], [Biggs2011], and [Fink2013] (a summary of which is freely available online [Fink2003].) In brief, lessons should be designed as follows:

1. Brainstorm to get a rough idea of what you want to cover, how you're going to do it, what problems or misconceptions you expect to encounter, what's *not* going to be included, and so on. You may also want to draw some concept maps at this stage.

2. Create or recycle learner personas (discussed in the next section) to figure out who you are trying to teach and what will appeal to them.

3. Create assessments that will give the learners a chance to practice the things they're trying to learn and tell you and them whether they're making progress and where they need to focus their work.

4. Put the assessments in order based on their complexity and dependencies to construct a course outline.

5. Write just enough to get learners from one formative assessment to the next. An actual classroom lesson will typically then consist of three or four such episodes, each building toward a short check that learners are keeping up.

This method helps to keep teaching focused on its objectives. It also ensures that learners don't face anything on the final exam that the course hasn't prepared them for.

Building Lessons by Subtracting Complexity

One way to build a programming lesson is to write the program you want learners to finish with, then remove the most complex

part that you want them to write and make it the last exercise. You can then remove the next most complex part you want them to write and make it the penultimate exercise, and so on. Anything that's left–i.e., anything you don't want them to write as an exercise–becomes the starter file(s) that you give them. This typically includes things like importing libraries or helper functions to access data.

How and Why to Fake It

One of the most influential papers in the history of software engineering was Parnas and Clements' "A Rational Design Process: How and Why to Fake It". In it, the authors pointed out that in real life we move back and forth between gathering requirements, interface design, programming, and testing, but when we write up our work it's important to describe it as if we did these steps one after another so that other people can retrace our steps. The same is true of lesson design: while we may change our mind about what we want to teach based on something that occurs to us while we're writing an MCQ, we want the notes we leave behind to present things in the order described above.

Teaching to the Test

Reverse instructional design is not *the same thing as "teaching to the test". When using RID, teachers set goals to aid in lesson design, and may never actually give the final exam that they wrote. In many school systems, on the other hand, an external authority defines assessment criteria for all learners, regardless of their individual situations, and the outcomes of those summative assessments directly affect the teachers' pay and promotion. Green's* Building a Better Teacher *[Green2014] argues that this focus on measurement is appealing to those with the power to set the tests, but is unlikely to improve outcomes unless it is coupled with support for teachers to make improvements based on test outcomes. This is often missing, because as [Scott1999] pointed out, large organizations usually value uniformity over productivity.*

6.1   Learner Personas

A key step in the process above is figuring out who your audience is. One way to do this is to write two or three *learner personas*. This technique is borrowed from user interface design, where short profiles of typical users are created to help designers think about

their audience's needs, and to give them a shorthand for talking about specific cases.

Learner personas have five parts: the person's general background, what they already know, what *they* think they want to do, how the course will help them, and any special needs they might have. A learner persona for a weekend workshop aimed at new college students might be:

1. Jorge has just moved from Costa Rica to Canada to study agricultural engineering. He has joined the college soccer team, and is looking forward to learning how to play ice hockey.

2. Other than using Excel, Word, and the Internet, Jorge's most significant previous experience with computers is helping his sister build a WordPress site for the family business back home in Costa Rica.

3. Jorge needs to measure properties of soil from nearby farms using a handheld device that sends logs in a text format to his computer. Right now, Jorge has to open each file in Excel, crop the first and last points, and calculate an average.

4. This workshop will show Jorge how to write a little Python program to read the data, select the right values from each file, and calculate the required statistics.

5. Jorge can read English proficiently, but still struggles sometimes to keep up with spoken conversation (especially if it involves a lot of new jargon).

A single learner persona is sometimes enough, but two or three that cover the whole range of potential learners is better. One of the ways they help is by serving as a shorthand for design issues: when speaking with each other, lesson authors can say, "Would Jorge understand why we're doing this?" or, "What installation problems would Jorge face?"

Our Learners Revisited

The personas of Samira and Moshe in the introduction (§1) have the five points listed above, rearranged to flow more readably.

Deciding What to Teach

There are two ways to decide what to teach: pick material and then find an audience, or decide on an audience and then figure

out what they want to learn. Either way, Guzdial's "Five Principles for Programming Languages for Learners" offers essential guidance:

1. *Connect to what learners know.*
2. *Keep cognitive load (§5) low.*
3. *Be honest (i.e., use authentic tasks).*
4. *Be generative and productive.*
5. *Test your ideas rather than trusting your instincts.*

6.2 Learning Objectives

Summative and formative assessments help instructors figure out what they're going to teach, but in order to communicate that to learners and other instructors, a course description should also have *learning objectives* (sometimes also called a *learning goal*). A learning objective is a single sentence describing what a learner will be able to do once they have sat through the lesson in order to demonstrate what they have learned.

Learning objectives are meant to ensure that everyone has the same understanding of what a lesson is supposed to accomplish. For example, a statement like "understand Git" could mean any of the following, each of this would be backed by a very different lesson:

- Learners can describe three scenarios in which version control systems like Git are better than file-sharing tools like Dropbox, and two in which they are worse.

- Learners can commit a changed file to a Git repository using a desktop GUI tool.

- Learners can explain what a detached HEAD is and recover from it using command-line operations.

Objectives vs. Outcomes

A learning objective is what a lesson strives to achieve. A learning outcome is what it actually achieves, i.e., what learners actually take away. The role of summative assessment is therefore to compare outcomes with objectives.

More specifically, a good learning objective has a *measurable or verifiable verb* that states what the learner will do, and specifies the *criteria for acceptable performance*. Writing these kinds of learning objectives may initially seem restrictive or limiting, but will make

both you, your fellow instructors, and your learners happier in the long run. You will end up with clear guidelines for both your teaching and assessment, and your learners will appreciate the clear expectations.

One way to understand what makes for a good learning objective is to see how a poor one can be improved:

- "Learner will be given opportunities to learn good programming practices." *Describes the lesson's content, not the attributes* of successful students.*

- "Learner will have a better appreciation for good programming practices." *Doesn't start with an active verb or define the* level of learning, and the subject of learning has no context and *is not specific.*

- "Learner will understand how to program in R." Starts with an *active verb, but doesn't define the level of learning, and the* subject of learning is still too vague for assessment.*

- "Learner will write one-page read-filter-summarize-print data analysis scripts for tabular data using R and R Studio." *Starts with an active verb, defines the level of learning, and provides context to ensure that outcomes can be assessed.*

Bloom's taxonomy can be used to organize learning objectives. First published in 1956, it attempts to define levels of understanding in a way that is hierarchical, measurable, stable, and cross-cultural. The list below defines its levels and shows some of the verbs typically used in learning objectives written for each level.

- Knowledge: recalling learned information (name, define, recall).

- Comprehension: explaining the meaning of information (restate, locate, explain, recognize).

- Application: applying what one knows to novel, concrete situations (apply, demonstrate, use).

- Analysis: breaking down a whole into its component parts and explaining how each part contributes to the whole (differentiate, criticize, compare).

- Synthesis: assembling components to form a new and integrated whole (design, construct, organize).

- Evaluation: using evidence to make judgments about the relative merits of ideas and materials (choose, rate, select).

Another way to think about learning objectives comes from [Fink2013], which defines learning in terms of the change it is intended to produce in the learner. Fink's Taxonomy has six categories:

- Foundational Knowledge: understanding and remembering information and ideas (remember, understand, identify).
- Application: skills, critical thinking, managing projects (use, solve, calculate, create).
- Integration: connecting ideas, learning experiences, and real life (connect, relate, compare).
- Human Dimension: learning about oneself and others (come to see themselves as, understand others in terms of, decide to become).
- Caring: developing new feelings, interests, and values (get excited about, be ready to, value).
- Learning How to Learn: becoming a better student (identify source of information for, frame useful questions about).

A set of learning objectives based on this taxonomy for an introductory course on HTML and CSS might be:

By the end of this course, students will:

- *Understand the difference between markup and presentation, the nested nature of HTML, what CSS properties are, and how CSS selectors work.*

- *Know how to write and style a web page using common tags and CSS properties.*

- *Be able to compare and contrast authoring with HTML and CSS to authoring with desktop publishing tools.*

- *Understand how the visually impaired interact and people in low-bandwidth environments interact with web pages and take their needs into account when designing new pages.*

- *Understand the role that JavaScript plays in styling web pages and want to learn more about how to use it.*

- *Be familiar with W3Schools[1] and other free tutorials for HTML and CSS, and know what search terms to use to find answers on Stack Overflow[2].*

---

[1]https://www.w3schools.com/
[2]https://stackoverflow.com/

6.3 Maintainability

Good courses take a lot of effort to build, but building them is only the first challenge. Once they have been written, someone needs to maintain them, and doing that is a lot easier if the lessons have been built in a maintainable way.

But what exactly does "maintainable" mean? The short answer is that a course is maintainable if it's cheaper to update it than to replace it. This equation depends on many factors, only some of which are under our control:

1. *How well documented the course's design is.* If the person doing maintenance doesn't know (or doesn't remember) what the course is supposed to accomplish or why topics are introduced in a particular order, it will take her more time to update it. One of the reasons to use the template described earlier is to capture decisions about why each course is the way it is.

2. *How the course's content is structured.* Version control is the secret sauce that allows software development to scale, but today's version control systems (still) can't handle widely-used file formats like Word and PowerPoint. Lessons should therefore either be written in plain-text formats like HTML, Markdown, or LaTeX, or stored online in systems like Google Docs that allow many people to edit the same files. (The next section discusses this in more detail.)

3. *How easy it is for collaborators to collaborate technically.* Lesson authors usually share material by passing it from hand to hand (or equivalently, by emailing files to each other or putting them in a shared drive. Collaborative writing tools like Google Docs[3] and wikis are a big improvement, as they allow many people to update the same document and comment on other people's updates. The version control systems used by programmers, such as GitHub[4], are another big advance, since they let any number of people work independently and then merge their changes back together in a controlled, reviewable way. Unfortunately, version control systems have a long, steep learning curve, which makes shared online authoring systems like Google Docs[5] and wikis the best technical choice for most groups.

The True Cost of Video

---

[3]http://docs.google.com
[4]http://github.com
[5]http://docs.google.com

Making a small change to this webpage only takes a few minutes, but in our experience, making any kind of change to a video takes an hour or more. In addition, most people are much less comfortable recording themselves than contributing written material.

The fourth factor, and the most important one in practice, is *how willing people are to collaborate.* The tools needed to build a "Wikipedia for lessons" or a "GitHub for lessons" have been around for almost twenty years, but neither model has caught on. When asked why not, teachers raise many objects[6], none of which hold up to close inspection:

- *The most important thing about a lesson isn't having it, but* writing *it, because that gives you a chance to figure out what you think about the topic.* This objection rhymes with my personal experience, but the same is true of software, and somehow we get up-and-coming programmers to use and improve libraries rather than building their own stuff from scratch.

- *It's just more trouble than it's worth, because it's always easier in the short term to write something from scratch than to learn your way around someone else's material.* And yet most teachers use textbooks, and most actors perform other people's plays, and...

- *It doesn't pay off for most teachers because they only teach any particular lesson once a year (or once a quarter).* Infrequent teaching ought to push people *toward* re-use, not away from it.

- *Working at scale results in a more neutral point of view (the average of the contributors' personal views), but in many fields, lessons are valuable precisely because they're one person's opinion.* This is true for literature, but for basic algebra? And if the difference is one of teaching method rather than content, then yeah, there should be half a dozen different shared lessons on polynomials, each approaching the topic in a different way.

- *There's no onboarding process to teach people the mechanics of distributed ad hoc large-scale collaboration.* This is undoubtedly a contributing factor, but (a) teachers get more training in how to develop lessons than most programmers get in how to take part in an open source project and (b) lack of a formal onboarding process hasn't slowed down Wikipedia.

---

[6]http://blog.mrmeyer.com/2016/why-secondary-teachers-dont-want-a-github-for-lesson-plans/

- *Collaboration on lesson development gets squeezed out by more important things (where "important" means "to the principal or chair").* Again, this should push people *toward* collaboration (possibly under official radar), since every minute they don't spend writing a lesson is a minute they can use to satisfy the principal or chair.

- *The Firewall of Doom at many schools prevents people from working on shared materials.* Probably true for some people, but this is not true for all and most teachers in industrialized countries have access to a computer at home these days.

- *The stakes are too high for teachers who are going to be evaluated on their teaching.* This may be true for some teachers, but isn't a universal.

- *No measurable outcome will show improvement, so there's no incentive to do it.* The same is true of open source software, but while only a small minority of programmers contribute, that's still enough people for it to thrive.

- *It's a generational thing: as digital natives, tomorrow's teachers will just naturally do it.* Millenials don't actually act that differently from their elders, and "not yet" arguments are as unfalsifiable as the claims by members of millenarian movements that the apocalypse is definitely coming–yup, any day now.

- *You can't run regression tests on a lesson, so there's no easy way to* tell if my changes have broken something that you wrote.* But Wikipedia...

One interesting observation is that while teachers don't collaborate at scale, they *do* remix by finding other people's materials online or in textbooks and reworking them. That suggests that the root problem may be a flawed analogy: rather than lesson development being like writing Wikipedia articles or open source software, perhaps it's more like postmodern music.

If this is true, then lessons may be the wrong granularity for sharing, and collaboration might be more likely to take hold if the thing being collaborated on was smaller. This fits well with Caulfield's theory of choral explanations[7]. He argues that sites like Stack Overflow[8] succeed because they provide a chorus of answers for every question, each of which is most suitable for a slightly different

---

[7]https://hapgood.us/2016/05/13/choral-explanations/
[8]https://stackoverflow.com/

questioner. If Caulfield is right, the future of learning–particularly online learning–may lie in guided tours of community-curated Q&A repositories rather than in things we would recognize as "lessons" today.

6.4 A Reminder

When designing a lesson, you must always remember that *you are not your learners*. You may be older (or younger, if you're teaching seniors) or wealthier (and therefore able to afford to download videos without foregoing a meal to pay for the bandwidth), but you are almost certainly more knowledgeable about technology. Don't assume that you know what they need or will understand: ask them, and actually pay attention to their answer. After all, it's only fair that learning should go both ways.

6.5 Challenges

Learner Personas (30 minutes)

Working in pairs or small groups, create a five-point persona that describes one of your typical learners.

Write Learning Objectives (20 minutes)

Write one more learning objectives for something you currently teach or plan to teach using Bloom's Taxonomy. Working with a partner, critique and improve the objectives.

Write More Learning Objectives (20 minutes)

Write one more learning objectives for something you currently teach or plan to teach using Fink's Taxonomy. Working with a partner, critique and improve the objectives.

Inessential Weirdness (15 minutes)

Betsy Leondar-Wright coined the phrase inessential weirdness[9] to describe things groups do that aren't really necessary, but which alienate people who aren't already members of that group. Sumana HariHareswara later used this notion as the basis for a talk on

---

[9]http://www.classmatters.org/2006_07/its-not-them.php

inessential weirdnesses in open source software[10]. Take a few minutes to read both articles, and then make a list of inessential weirdnesses you think your learners might encounter when you first teach them. How many of these can you avoid with a little effort?

---

[10]https://www.harihareswara.net/sumana/2016/05/21/0

7   Motivation and Demotivation

Objectives

- *Learners can name and describe the three principal ways in which they can demotivate their own learners.*
- *Learners can define impostor syndrome and stereotype threat, and describe ways in which to combat each.*
- *Learners can describe the difference between fixed and growth mindset and explain the importance of encouraging the latter.*
- *Learners can describe and enact at least three things they can do to make their programming workshops more accessible.*
- *Learners can describe and enact at least three things they can do to make their programming workshops more inclusive.*

Learners need encouragement to step out into unfamiliar terrain, so this chapter discusses ways instructors can motivate them. More importantly, it discusses ways that we can accidentally *demotivate* them, and how we can avoid doing that.

People learn best when they care about the topic and believe they can master it [Wlodkowski2017]. This presents us with a problem because most people don't actually want to program: they want to make music or compare changes to zoning laws with family incomes, and rightly regard programming as a tax they have to pay in order to do so. In addition, their early experiences with programming are often demoralizing, and believing that something will be hard to learn is a self-fulfilling prophecy.

Imagine a grid whose axes are "mean time to master" and "usefulness once mastered":

Everything that's quick to master and immediately useful should be taught first; things in the opposite corner that are hard to learn and have little near-term application should be dropped, and topics along the diagonal should be weighed against each other.

Figure 7.1: What to Teach

Actual Time

Any useful estimate of how long something takes to master must take into account how frequent failures are and how much time is lost to them. For example, editing text files seems like a simple task, but what about finding those files? Most graphical editors save things to the user's desktop or home directory; if people need to run shell commands on the files they've edited, a substantial fraction won't be able to navigate to the right directory without help. If this seems like a small problem to you, please revisit the discussion of expert blind spot (§4).

Many of the foundational concepts of computer science, such as computability, inhabit the "useful but hard to learn" corner of the grid described above. This doesn't mean that they aren't worth learning, but if our aim is to convince people that they *can* learn this

stuff, and that doing so will help them do more science faster, they are less compelling than things like automating repetitive tasks.

We therefore recommend a "teach most immediately useful first" approach. Have learners do something that *they* think is useful in their daily work within a few minutes of starting each lesson. This not only motivates them, it also helps build their confidence in us, so that if it takes longer to get to the payoff of a later topic, they'll stick with us.

The best-studied use of this idea is the media computation approach developed by Guzdial and Ericson at Georgia Tech [Guzdial2013]. Instead of printing "hello world" or summing the first ten integers, their students' first program opens an image, resizes it to create a thumbnail, and saves the result. This is an *authentic task*, i.e., something that learners believe they would actually do in real life. It is also has a *tangible artifact*: if the image comes out the wrong size, learners have a concrete starting point for debugging.

Strategies for Motivating Learners

[Ambrose2010] contains a list of evidence-based methods to motivate learners. None of them are surprising–it's hard to imagine someone saying that we shouldn't identify and reward what we value–but it's useful to check lessons against these points to make sure they're doing at least a few of these things.

What's missing from this list is strategies to motivate the instructor. Learners respond to an instructor's enthusiasm, and instructors need to care about a topic in order to keep teaching it, particularly when they are volunteers.

7.1   Demotivation

Women aren't leaving computing because they don't know what it's like; they're leaving because they **do** know.
– variously attributed

If you are teaching free-range learners, they are probably already motivated–if they weren't, they wouldn't be in your classroom. The challenge is therefore not to demotivate them. Unfortunately, we can do this by accident much more easily than you might think.

The three most powerful demotivators are *unpredictability*, *indifference*, and *unfairness*. Unpredictability demotivates people because if there's no reliable connection between what they do and what outcome they achieve, there's no reason for them to try to

do anything. If learners believe that the instructor or the educational system doesn't care about them or the lesson, they won't care either. And if people believe the class is unfair, they will also be demotivated, even if it is unfair in their favor (because consciously or unconsciously they will worry that they will some day find themselves in the group on the losing end [Wilkinson2011]). In extreme situations, learners may develop learned helplessness: when repeatedly subjected to negative feedback that they have no way to escape, they may learn not to even try to escape when they could.

Here are some quick ways to demotivate your learners:

- A "holier-than-thou" or contemptuous attitude from an instructor.

- Tell learners they are rubbish because they use Excel and/or Word, don't modularize their code, etc.

- Repeatedly make digs about Windows and praise Linux, e.g., say that the former is for amateurs.

- Criticize GUI applications (and by implication their users) and describe command-line tools as the One True Way.

- Dive into complex or detailed technical discussion with the one or two people in the audience who clearly don't actually need to be there.

- Pretend to know more than you do. People will actually trust you more if you are frank about the limitations of your knowledge, and will be more likely to ask questions and seek help.

- Use the J word ("just"). As discussed earlier (§4), this signals to the learner that the instructor thinks their problem is trivial and by extension that they therefore must be stupid for not being able to figure it out.

- Feign surprise. Saying things like "I can't believe you don't know X" or "you've never heard of Y?" signals to the learner that they do not have some required pre-knowledge of the material you are teaching, that they are in the wrong place, and it may prevent them from asking questions in the future.

Code of Conduct Revisited

As noted at the start (§1), we believe very strongly that classes should have a Code of Conduct (§14). Its details are important, but the most important thing about it is that it exists: knowing

that we have rules tells people a great deal about our values and about what kind of learning experience they can expect.

Never Learn Alone

One way to support learners who have been subject to systematic exclusion or discrimination (overt or otherwise) is to have people sign up for workshops in small teams rather than as individuals. If an entire lab group comes, or if attendees are drawn from the same (or closely-related) disciplines, everyone in the room will know in advance that they will be with at least a few people they trust, which increases the chances of them actually coming. It also helps after the workshop: if people come with their friends or colleagues, they can work together to implement what they've learned.

7.2 Impostor Syndrome

Impostor syndrome is the belief that one is not good enough for a job or position, that one's achievements are lucky flukes, and an accompanying fear of being "found out". Impostor syndrome seems to be particularly common among high achievers who undertake publicly visible work[1].

Academic work is frequently undertaken alone or in small groups but the results are shared and criticized publicly. In addition, we rarely see the struggles of others, only their finished work, which can feed the belief that everyone else finds it easy. Women and minority groups who already feel additional pressure to prove themselves in some settings may be particularly affected.

Two ways of dealing with your own impostor syndrome are:

1. Ask for feedback from someone you respect and trust. Ask them for their honest thoughts on your strengths and achievements, and commit to believing them.

2. Look for role models. Who do you know who presents as confident and capable? Think about how they conduct themselves. What lessons can you learn from them? What habits can you borrow? (Remember, they quite possibly also feel as if they are making it up as they go.)

As an instructor, you can help people with their impostor syndrome by sharing stories of mistakes that you have made or things

---

[1]https://www.usenix.org/blog/impostor-syndrome-proof-yourself-and-your-community

you struggled to learn. This reassures the class that it's OK to find topics hard. Being open with the group makes it easier to build trust and make students confident to ask questions. (Live coding is great for this: typos let the class see you're not superhuman.)

You can also emphasize that you want questions: you are not succeeding as a teacher if no one can follow your class, so you're asking students for their help to help you learn and improve. (Remember, it's much more important to *be* smart than to *look* smart.) If you would like to learn more, the Ada Initiative has some excellent resources for teaching about and dealing with imposter syndrome [Ada2017].

7.3 Stereotype Threat

Reminding people of negative stereotypes, even in subtle ways, can make them anxious about the risk of confirming those stereotypes, which in turn reduces their performance. This is called *stereotype threat*, and the clearest examples in computing are gender-related. Depending on whose numbers you trust, only 12-18% of programmers are women, and those figures have actually been getting worse over the last 20 years. There are many reasons for this (see [Margolis2003] and [Margolis2010]), and [Steele2011] summarizes what we know about stereotype threat in general and presents some strategies for mitigating it in the classroom.

However, while there's lots of evidence that unwelcoming climates demotivate members of under-represented groups, it's not clear that stereotype threat is the underlying mechanism. Part of the problem is that the term has been used in many ways[2]; another is questions[3] about the replicability of key studies. What *is* clear is that we need to avoid thinking in terms of a deficit model (i.e., we need to change the members of under-represented groups because they have some deficit, such as lack of prior experience) and instead use a systems approach (i.e., we need to change the system because it produces these disparities).

A great example of how stereotypes work was presented in [Patitsas2016]. This thought-provoking paper showed that people see evidence for a "geek gene" where none exists:

> *Although it has never been rigorously demonstrated, there is a common belief that CS grades are bimodal. We statistically*

---

[2]http://www.europhd.net/html/_onda02/07/PDF/20th_lab_materials/jane/shapiro_neuberg_2007.pdf

[3]https://www.psychologytoday.com/blog/rabble-rouser/201512/is-stereotype-threat-overcooked-overstated-and-oversold

analyzed 778 distributions of final course grades from a large research university, and found only 5.8% of the distributions passed tests of multimodality. We then devised a psychology experiment to understand why CS educators believe their grades to be bimodal. We showed 53 CS professors a series of histograms displaying ambiguous distributions and asked them to categorize the distributions. A random half of participants were primed to think about the fact that CS grades are commonly thought to be bimodal; these participants were more likely to label ambiguous distributions as "bimodal". Participants were also more likely to label distributions as bimodal if they believed that some students are innately predisposed to do better at CS. These results suggest that bimodal grades are instructional folklore in CS, caused by confirmation bias and instructor beliefs about their students.

It's easy to use language that suggests that some people are natural programmers and others aren't, but Mark Guzdial has called this belief the biggest myth about teaching computer science[4].

7.4 Mindset

Learners can be demotivated in subtler ways as well. For example, Dweck and others have studied the differences of fixed mindset and growth mindset. If people believe that competence in some area is intrinsic (i.e., that you either "have the gene" for it or you don't), *everyone* does worse, including the supposedly advantaged. The reason is that if they don't get it at first, they figure they just don't have that aptitude, which biases future performance. On the other hand, if people believe that a skill is learned and can be improved, they do better on average.

A person's mindset can be shaped by subtle cues. For example, if a child is told, "You did a good job, you must be very smart," they are likely to develop a fixed mindset. If on the other hand they are told, "You did a good job, you must have worked very hard," they are likely to develop a growth mindset, and subsequently achieve more. Studies have also shown that the simple action of telling learners about the different mindsets before a course can improve learning outcomes for the whole group.

As with stereotype threat, there are concerns[5] that growth mindset has been oversold, or that research is much more difficult to

---

[4]http://cacm.acm.org/blogs/blog-cacm/189498-top-10-myths-about-teaching-computer-science/fulltext

[5]http://www.learningspy.co.uk/psychology/growth-mindset-bollocks/

put into practice than its more enthusiastic advocates have implied. While some people interpret this back and forth of claim and counter-claim as evidence than education research isn't reliable, what it really shows is that anything involving human subjects is both subtle and difficult.

7.5   Accessibility

Not providing equal access to lessons and exercises is about as demotivating as it gets. The older Software Carpentry lessons, for example, the text beside the slides includes all of the narration–but none of the Python source code. Someone using a screen reader[6] would therefore be able to hear what was being said about the program, but wouldn't know what the program actually was.

While it may not be possible to accommodate everyone's needs, it *is* possible to get a good working structure in place without any specific knowledge of what specific disabilities people might have. Having at least some accommodations prepared in advance also makes it clear that hosts and instructors care enough to have thought about problems in advance, and that any additional concerns are likely to be addressed.

It Helps Everyone

Curb cuts[7] (the small sloped ramps joining a sidewalk to the street) were originally created to make it easier for the physically disabled to move around, but proved to be equally helpful to people with strollers and grocery carts. Similarly, steps taken to make lessons more accessible to people with various disabilities also help everyone else. Proper captioning of images, for example, doesn't just give screen readers something to say: it also makes the images more findable by exposing their content to search engines.

The first and most important step in making lessons accessible is to *involve people with disabilities in decision-making*: the slogan *nihil de nobis, sine nobis*[8] (literally, "nothing about us, without us") predates accessibility rights, but is always the right place to start. A few other recommendations are:

- *Find out what you need to do.* The W3C Accessibility Initiative's checklist for presentations [W3C2017] is a good starting point

---

[6]https://en.wikipedia.org/wiki/Screen_reader
[7]https://en.wikipedia.org/wiki/Curb_cut
[8]https://en.wikipedia.org/wiki/Nothing_About_Us_Without_Us

focused primarily on assisting the visually impaired, while Liz Henry's blog post about accessibility at conferences [Henry2014] has a good checklist for people with mobility issues, and this interview with Chad Taylor is a good introduction to issues faced by the hearing impaired [Taylor2014].

- *Know how well you're doing.* For example, sites like WebAIM[9] allow you to check how accessible your online materials are to visually impaired users.

- *Don't do everything at once.* We don't ask learners in our workshops to adopt all our best practices or tools in one go, but instead to work things in gradually at whatever rate they can manage. Similarly, try to build in accessibility habits when preparing for workshops by adding something new each time.

- *Do the easy things first.* There are plenty of ways to make workshops more accessible that are both easy and don't create extra cognitive load for anyone: font choices, general text size, checking in advance that your room is accessible via an elevator or ramp, etc.

7.6 Inclusivity

Inclusivity is a policy of including people who might otherwise be excluded or marginalized. In computing, it means making a positive effort to be more welcoming to women, people of color, people with various sexual orientations, the elderly, the physically challenged, the formerly incarcerated, the economically disadvantaged, and everyone else who doesn't fit Silicon Valley's white/Asian male demographic. Lee's paper "What can I do today to create a more inclusive community in CS?" [Lee2017] is a brief, practical guide to doing that with references to the research literature. These help learners who belong to one or more marginalized or excluded groups, but help motivate everyone else as well; while they are phrased in terms of term-long courses, many can be applied in our workshops:

- Ask learners to email you before the workshop to explain how they believe the training could help them achieve their goals.

- Review notes to make sure they are free from gendered pronouns, that they include culturally diverse names, etc.

---

[9]http://webaim.org/

- Emphasize that what matters is the rate at which they are learning, not the advantages or disadvantages they had when they started.

- Encourage pair programming.

- Actively mitigate behavior that some learners may find intimidating, e.g., use of jargon or "questions" that are actually asked to display knowledge.

7.7 Challenges

Authentic Tasks (15 minutes)

Think about something you did this week that uses one or more of the skills you teach, (e.g., wrote a function, bulk downloaded data, did some stats in R, forked a repo) and explain how you would use it (or a simplified version of it) as an exercise or example in class.

1. Pair up with your neighbor and decide where this exercise fits on a 2x2 grid of "short/long time to master" and "low/high usefulness".
2. Write the task and where it fits on the grid.
3. As a group, discuss how these relate back to the "teach most immediately useful first" approach.

Implement One Strategy for Inclusivity (5 minutes)

Pick one activity or change in practice from Lee's paper [Lee2017] that you would like to work on. Put a reminder in your calendar three months in the future to self-check whether you have done something about it.

Brainstorming Motivational Strategies (20 minutes)

1. Think back to a programming course (or any other) that you took in the past, and identify one thing the instructor did that demotivated you, and describe what could have been done afterward to correct the situation.

2. Pair up with your neighbor and discuss your stories, then add your comments to the shared notes.

3. Review the comments in the shared notes as a group. Rather than read them all out loud, highlight and discuss a few of the things that could have been done differently. This will give

everyone some confidence in how to handle these situations in the future.

Demotivational Experiences (15 minutes)

Think back to a time when you demotivated a student (or when you were demotivated as a student). Pair up with your neighbor and discuss what you could have done differently in the situation, and then share the story and what could have been done in the group notes.

Walk the Route (15 minutes)

Find the nearest public transportation drop-off point to your building and walk from there to your office and then to the nearest washroom, making notes about things you think would be difficult for someone with mobility issues. Now borrow a wheelchair and repeat the journey. How complete was your list of challenges? And did you notice that the first sentence in this challenge assumed you could actually walk?

Who Decides? (15 minutes)

In [Littky2004], Kenneth Wesson wrote, "If poor inner-city children consistently outscored children from wealthy suburban homes on standardized tests, is anyone naive enough to believe that we would still insist on using these tests as indicators of success?" Read [Cottrill2016], and then describe an example from your own experience of "objective" assessments that reinforced the status quo.

Credibility (15 minutes)

[Fink2013] describes three things that make teachers credible in their learners' eyes:

- **Competence**, or knowledge of the subject, as shown by the ability to explain complex ideas or reference the work of others.

- **Trustworthiness**, or having the student's best interests in mind. This can be shown by giving individualized feedback, offering a rational explanation for grading decisions, and treating all students the same.

- **Dynamism**, or excitement about the subject. This was discussed in detail in the chapter on teaching as a performance art (§3).

Describe one thing you do when teaching that fits into each category, and then describe one thing you *don't* do but should for each category as well.

Common Stereotypes (10 minutes)

You will (still) sometimes hear people say, "It's so simple that even your grandmother could use it." In pairs, list two or three other phrases that reinforce stereotypes about computing.

8   Live Coding

Objectives

- *Learners can describe live coding and explain its advantages as a teaching practice for programming workshops.*
- *Learners can enact and critique live coding.*

Teaching is theater not cinema.
– Neal Davis

Teaching is a performance art, just like drama, music, and athletics. And as in those fields, we have a collection of small tips and tricks to make teaching work better.

The first of our recommended teaching practices is so central that it deserves a chapter of its own: *live coding.* When they are live coding, instructors don't use slides. Instead, they go through the lesson material, typing in the code or instructions, with their learners following along. Its advantages are:

1. Watching a program being written is more compelling than watching someone page through slides that present bits and pieces of the same code.

2. It enables instructors to be more responsive to "what if?" questions. Where a slide deck is like a railway track, live coding allows instructors to go off road and follow their learners' interests.

3. It facilitates lateral knowledge transfer: people learn more than we realized we were teaching by watching *how* instructors do things.

4. It slows the instructor down: if she has to type in the program as she goes along, she can only go twice as fast as her learners, rather than ten-fold faster as she could with slides.

5. Learners get to see instructors' mistakes *and how to diagnose and correct them.* Novices are going to spend most of their time doing this, but it's left out of most textbooks.

6. Watching instructors make mistakes shows learners that it's all right to make mistakes of their own. Most people model the behavior of their teachers: if the instructor isn't embarrassed about making and talking about mistakes, learners will be more comfortable doing so too.

Live coding does have some drawbacks, but with practice, these can be avoided or worked around:

1. Instructors can go too slowly, either because they are not good typists or by spending too much time looking at notes to try to remember what they meant to type.

2. Typing in boilerplate code that is needed by the lesson, but not directly relevant to it (such as library import statements) increases the extraneous cognitive load (§5) on your learners. Willingham says "Memory is the residue of thought" [Willingham2010], so if you spend your time typing boilerplate, that may be what learners will take away.

Live coding is an example of the "I/We/You" approach to teaching discussed in Performance (§3). It takes a bit of practice for instructors to get used to thinking aloud while coding in front of an audience, but most report that it is then no more difficult to do than talking off a deck of slides.

8.1   Be Seen and Heard

If you are physically able to stand up for a couple of hours, do it while you are teaching. When you sit down, you are hiding yourself behind others for those sitting in the back rows. Make sure to notify the workshop organizers of your wish to stand up and ask them to arrange a high table, standing desk, or lectern.

Regardless of whether you are standing or sitting, make sure to move around as much as reasonable. You can for example go to the screen to point something out, or draw something on the white/blackboard (see below). Moving around makes the teaching more lively, less monotonous. It draws the learners' attention away from their screens, to you, which helps get the point you are making across.

Even though you may have a good voice and know how to use it well, it may be a good idea to use a microphone, especially if the workshop room is equipped with one. Your voice will be less tired, and you increase the chance of people with hearing difficulties being able to follow the workshop.

8.2   Take It Slow

For every command you type, every word of code you write, every menu item or website button you click, say out loud what you are doing while you do it, then point to the command and its output on the screen and go through it a second time. This not only slows you down, it allows learners who are following along to copy what you do, or to catch up, even when they are looking at their screen while doing it. Whatever you do, *don't* copy and paste code: doing this practically guarantees that you'll race ahead of your learners.

If the output of your command or code makes what you just typed disappear from view, scroll back up so learners can see it again - this is especially needed for the Unix shell lesson. Other options are to execute the same command a second time, or to copy and paste the last command(s) into the workshop's shared notes.

8.3   Mirror Your Learner's Environment

You may have set up your environment to your liking, with a very simple or rather fancy Unix prompt, colour schemes for your development environment, keyboard shortcuts etc. Your learners usually won't have all of this. Try to create an environment that mirrors what your learners have, and avoid using keyboard shortcuts. Some instructors create a separate bare-bones user (login) account on their laptop, or a separate teaching-only account on the service being taught (e.g., Github).

8.4   Use the Screen Wisely

You will need to enlarge your font considerably in order for people to read it from the back of the room, which means you can put much less on the screen than you're used to. (You will often be reduced to 60-70 columns and 20-30 rows, which basically means that you're using a 21st Century supercomputer to emulate an early-1980s VT100 terminal.)

To cope with this, maximize your window, and then ask everyone to give you a thumbs-up or thumbs-down on its readability. Use a

black font on a lightly-tinted background rather than a light font on a dark background–the light tint will glare less than a pure white background.

When the bottom of the projector screen is at the same height, or below, the heads of the learners, people in the back won't be able to see the lower parts. Draw up the bottom of your window(s) to compensate.

Pay attention to the room lighting as well: it should not be fully dark, and there should be no lights directly on or above the presenter's screen. If needed, reposition the tables so all learners can see the screen.

If you can get a second screen, use it: the extra screen real estate will allow you to display your code on one side and its output or behavior on the other. The second screen may require its own PC or laptop, so you may need to ask a helper to control it.

Multiple Personalities

If you teach using a console window, such as a Unix shell, it's important to tell people when you run an in-console text editor and when you return to the console prompt. Most novices have never seen a window take on multiple personalities in this way, and can quickly become confused (particularly if the window is hosting an interactive interpreter prompt for Python or some other language as well as running shell commands and hosting an editor).

8.5   Double Devices

Many instructors now use two devices when teaching: a laptop plugged into the projector for learners to see, and a tablet beside it on which they can view their notes and the shared notes that the learners are taking. This seems to be more reliable than displaying one virtual desktop while flipping back and forth to another. Of course, printouts of the lesson material are still the most reliable backup technology. . .

8.6   Use Illustrations

Most lesson material comes with illustrations, and these may help learners to understand the stages of the lesson and to organize the material. What can work really well is when you as instructor generate the illustrations on the white/blackboard as you progress through the material. This allows you to build up diagrams, making

them increasingly complex in parallel with the material you are teaching. It helps learners understand the material, makes for a more lively workshop (you'll have to move between your laptop and the blackboard) and gathers the learners' attention to you as well.

8.7   Avoid Distractions

Turn off any notifications you may use on your laptop, such as those from social media, email, etc. Seeing notifications flash by on the screen distracts you as well as the learners, and may even result in awkward situations when a message pops up you'd rather not have others see.

8.8   Improvise *After* You Know the Material

The first time you teach a new lesson, you should stick fairly closely to the topics it lays out and the order they're in. It may be tempting to deviate from the material because you would like to show a neat trick, or demonstrate some alternative way of doing something. Don't do this, since there is a fair chance you'll run into something unexpected that you then have to explain.

Once you are more familiar with the material, though, you can and should start improvising based on the backgrounds of your learners, their questions in class, and what you find most interesting about the lesson. This is like a musician playing a new song: the first few times, you stick to the sheet music, but after you're comfortable with it, you can start to put your own stamp on it.

If you really want to use something outside of the material, try it out thoroughly before the workshop: run through the lesson as you would during the actual teaching and test the effect of your modification.

8.9   Embrace Mistakes

No matter how well prepared you are, you will be making mistakes. Typos are hard to avoid, you may overlook something from the lesson instructions, etc. This is OK! It allows learners to see instructors' mistakes and how to diagnose and correct them. Some mistakes are actually an opportunity to point something out, or reflect back on something covered earlier. Novices are going to spend most of their time making the same and other mistakes, but how to deal with them is left out of most textbooks.

The typos are the pedagogy.
– *Emily Jane McTavish*

Note: if you've given a lesson several times, you're unlikely to make anything other than basic typing mistakes (which usually aren't informative). It's worth remembering "real" mistakes and making them deliberately, but that often feels forced. A better approach is to get learners to tell you what to do next in the hope that this will get you into the weeds.

8.10   Face the Screen–Occasionally

It's OK to *face the screen occasionally*, particularly when you are walking through a section of code statement by statement or drawing a diagram, but you shouldn't do this for more than a few seconds at a time. Looking at the screen for a few seconds can help lower your anxiety levels, since it gives you a brief break from being looked at.

A good rule of thumb is to treat the screen as one of your learners: if it would be uncomfortable to stare at someone for as long as you are spending looking at the screen, it's time to turn around and face your audience.

8.11   Have Fun

Teaching is performance art and can be rather serious business. On the one hand, don't let this scare you - it is much easier than performing Hamlet. You have an excellent script at your disposal, after all! On the other hand, it is OK to add an element of play, i.e. use humor and improvisation to liven up the workshop. How much you are able and willing to do this is really a matter of personality and taste - as well as experience. It becomes easier when you are more familiar with the material, allowing you to relax more. Choose your words and actions wisely, though. Remember that we want the learners to have a welcoming experience and a positive learning environment - a misplaced joke can ruin this in an instant. Start small, even just saying 'that was fun' after something worked well is a good start. Ask your co-instructors and helpers for feedback when you are unsure of the effect your behavior has on the workshop.

8.12 Challenges

The Bad and the Good (20 minutes)

Watch the video of live coding done poorly [Nederbragt2016a] and then the video of live coding done well [Nederbragt2016b], summarize your feedback on both using the usual 2×2 grid, and then compare it with this list (§17.6). These videos assume learners know what a shell variable is, know how to use the head command, and are familiar with the contents of the data files being filtered.

See Then Do (30 minutes)

Teach 3-4 minutes of a lesson using live coding to a fellow trainee, then swap and watch while that person live codes for you. Don't bother trying to record the live coding sessions–we have found that it's difficult to capture both the person and the screen with a hand-held device–but give feedback the same way you have previously (positive and negative, content and presentation). Explain in advance to your fellow trainee what you will be teaching and what the learners you teach it to are expected to be familiar with.

- What felt different about live coding (versus standing up and lecturing)? What was harder/easier?

- Did you make any mistakes? If so, how did you handle them?

- Did you talk and type at the same time, or alternate?

- How often did you point at the screen? How often did you highlight with the mouse?

- What will you try to do differently next time?

9 Teaching Practices

Objectives

- *Learners can name, describe, and enact four teaching practices that are appropriate to use in programming workshops for adults, and give a pedagogical justification for each.*
- *Learners can explain why instructors should* not *introduce new pedagogical practices in a short workshop.*

Just as domain expertise is often a matter of pattern recognition (§5.1), teaching expertise often comes down to using good practices consistently. None of the practices described below are essential (except having a code of conduct), but each will improve lesson delivery.

9.1 Have a Code of Conduct

An important part of making a class productive is to treat everyone with respect. We therefore strongly recommend that every group offering classes based on this material adopt a Code of Conduct like this one (§14), and require people taking part in the class to abide by it.

We believe equally strongly that your actual programming classes should also have and enforce a Code of Conduct. Programming is a scary topic for many novices, and workshops are meant to be a judgment free space to learn and experiment. The behavior of the instructor and other participants may make more of an impression on a novice learner than any "technical" topic you teach.

If you do this, hosts should point people at it during registration, and instructors should remind attendees of it at the start of the workshop. The Code of Conduct doesn't just tell everyone what the rules are: it tells them that there *are* rules, and that they can therefore expect a safe and welcoming learning experience.

If you are an instructor, and believe that someone in a workshop has violated the Code of Conduct, you may warn them, ask them to apologize, and/or expel them, depending on the severity of the violation and whether or not you believe it was intentional. Whatever you do:

- Do it in front of witnesses. Most people will tone down their language and hostility in front of an audience, and having someone else present ensures that later discussion doesn't degenerate into conflicting claims about who said what.

- Contact the organizer or host of your class as soon as you can and describe what happened. Remember, a Code of Conduct is meaningless without a procedure for enforcing it.

A Code of Conduct cannot stop people from being offensive, any more than laws against theft stop people from stealing. What it *can* do is make expectations and consequences clear. In our experience, people rarely violate the Code of Conduct in person, though some are more likely to online, where they feel less inhibited. And remember, a Code of Conduct is *not* an infringement on free speech. People have a right to say what they think, but that doesn't mean they have a right to speak wherever and whenever they want. If someone wishes to say something disparaging about someone else, they can go and find a space of their own in which to say it.

9.2   Starting Out

To begin your class, the instructors should give a brief introduction that will convey their capacity to teach the material, accessibility and approachability, desire for student success, and enthusiasm. Tailor your introduction to the students' skill level so that you convey competence (without seeming too advanced) and demonstrate that you can relate to the students. Throughout the workshop, continually demonstrate that you are interested in student progress and that you are enthusiastic about the topics.

Students should also introduce themselves (preferably verbally). At the very least, everyone should add their name to the Etherpad, but its also good for everyone at a given site to know who all is in the group. Note: this can be done while setting up before the start of the class.

9.3   Overnight Homework

In a two-day class, have learners read the operations checklists as overnight homework and do their demotivational story just before lunch on day 2: it means day 2 starts with *their* questions (which wakes them up), and the demotivational story is a good lead-in to lunchtime discussion.

9.4   Never a Blank Page

Programming workshops (and other kinds of classes) can be built around a set of independent exercises, develop a single extended example in stages, or use a mixed strategy. The main advantages of independent exercises are that people who fall behind can easily re-synchronize, and that lesson developers can add, remove, and rearrange material at will. A single extended example, on the other hand, will show learners how the bits and pieces they're learning fit together: in educational parlance, it provides more opportunity for them to integrate their knowledge.

Whichever approach you take, learners should never start with a blank page (or screen), since they often find this intimidating or bewildering. Modifying existing code instead of writing new code from scratch doesn't just give them structure: it is also more realistic. Keep in mind, however, that starter code may increase cognitive load, since learners can be distracted by trying to understand it all before they start their own work.

9.5   Take Notes Together

Many studies have shown that taking notes while learning improves retention [Aiken1975], [Bohay2011]. As discussed in Memory (§4), this happens because taking notes forces you to organize and reflect on material as it's coming in, which in turn increases the likelihood that you will transfer it to long-term memory in a usable way.

Our experience, and some recent research findings, lead us to believe that taking notes *collaboratively* helps learning even more [Orndorff2015], even though taking notes on a computer is generally less effective than taking notes using pen and paper [Mueller2014]. Taking notes collaboratively:

- It allows people to compare what they think they're hearing with what other people are hearing, which helps them fill in gaps and correct misconceptions right away.

- It gives the more advanced learners in the class something useful to do. Rather than getting bored and checking Twitter during class, they often take the lead in recording what's being said, which keeps them engaged, and allows less advanced learners to focus more of their attention on new material. Keeping the more advanced learners busy also helps the whole class stay engaged because boredom is infectious: if a handful of people start updating their Facebook profiles, the people around them will start checking out too.

- The notes the learners take are usually more helpful *to them* than those the instructor would prepare in advance, since the learners are more likely to write down what they actually found new, rather than what the instructor predicted would be new.

- Glancing at the notes as they're being taken helps the instructor discover that the class didn't hear something important, or misunderstood it.

We usually use Etherpad[1] or Google Docs[2] for collaborative note-taking. The former makes it easy to see who's written what, while the latter scales better and allows people to add images to the notes. Whichever is chosen, classes also use it to share snippets of code and small datasets, and as a way for learners to show instructors their work (by copying and pasting it in).

Shared note-taking is almost always mentioned positively in post-workshop feedback. However, it's also common for participants to report that they find it distracting, as it's one more thing they have to keep an eye on. We believe the positives outweigh the negatives, but think that some careful controlled studies would tell us whether we're right, and how to use it better.

If you are going to have a group take notes together, make a list of everyone's name and paste it into the document each time you want every person to respond individually (e.g., to answer a question or contribute an exercise solution). This prevents the situation in which everyone is trying to edit the same couple of lines at the same time.

9.6 Assess Motivation and Prior Knowledge

Most formal educational systems train people to treat all assessment as summative, i.e., to think of every interaction with a teacher as

---

[1] http://etherpad.org

[2] https://docs.google.com

an evaluation, rather than as a chance to shape instruction. For example, we use a short pre-assessment questionnaire to profile learners before workshops to help instructors tune the pace and level of material. We send people this questionnaire out after they have registered rather than making it part of the sign-up process because when we did the latter, many people concluded that since they couldn't answer all the questions, they shouldn't enrol. We were therefore scaring off many of the people we most wanted to help.

Instead of asking people how easily they could complete specific tasks, we could just ask them to rate their knowledge of various subjects on a scale from 1 to 5. However, self-assessments of this kind are usually inaccurate because of the Dunning-Kruger effect[3]: the less people know about a subject, the less accurate their estimate of their knowledge is.

That said, there *are* things we can do:

- Before running a workshop, communicate its level clearly to everyone who's thinking of signing up by listing the topics that will be covered and showing a few examples of exercises that people will be asked to complete.

- Provide multiple exercises for each teaching episode so that more advanced learners don't finish early and get bored.

- Ask more advanced learners to help people next to them. They'll learn from answering their peers' questions (since it will force them to think about things in new ways).

- The helpers and the instructor who aren't teaching the particular episode should keep an eye out for learners who are falling behind and intervene early so that they don't become frustrated and give up.

The most important thing is to accept that no class can possibly meet everyone's individual needs. If the instructor slows down to accommodate two people who are struggling, the other 38 are not being well served. Equally, if she spends a few minutes talking about an advanced topic because two learners are bored, the 38 who don't understand it will feel left out. All we can do is tell our learners what we're doing and why, and hope that they'll understand.

It's important to design lessons with a particular audience in mind. It's equally important to find out who's in each specific audience, since this will influence how you introduce yourself, motivate

---

[3]https://en.wikipedia.org/wiki/Dunning%E2%80%93Kruger_effect

topics, and pace the lessons. Before the start of a Software Carpentry instructor training class, we ask people to fill in a short questionnaire like the one below. It doesn't tell us everything we might want to know, but it does give trainers a pretty clear idea of who they're speaking to.

1. Have you ever participated in a Software Carpentry workshop?

   - Yes, as a learner.
   - Yes, as a helper.
   - Yes, as an organizer.
   - Yes, as an instructor.
   - No, but I am familiar with what is taught at a workshop.
   - No, and I am not familiar with what is taught.

2. Which of these describes your teaching experience?

   - I have none.
   - I have taught a workshop or other informal course.
   - I have been a teaching assistant for a college-level course.
   - I have been the instructor for a college-level course.
   - I have taught at the K-12 level.

3. Which of these describes your previous formal training in teaching?

   - None
   - A few hours
   - A workshop
   - A certification or short course
   - A full degree

4. How frequently do you work with the tools that Software Carpentry teaches, such as R, Python, MATLAB, Perl, SQL, Git, and the Unix Shell?

   - Every day
   - A few times a week
   - A few times a month
   - A few times a year
   - Never or almost never

5. How often would you expect to teach a Software Carpentry workshop after this training?

   - Not at all
   - Once a year
   - Several times a year

6. Why do you want to take this training course?

9.7   Sticky Notes as Status Flags

Give each learner two sticky notes of different colours, e.g., red and green. These can be held up for voting, but their real use is as status flags. If someone has completed an exercise and wants it checked, they put the green sticky note on their laptop; if they run into a problem and need help, the put up the red one. This is better than having people raise their hands because:

- it's more discreet (which means they're more likely to actually do it),

- they can keep working while their flag is raised, and

- the instructor can quickly see from the front of the room what state the class is in.

   Sometimes a red sticky involves a technical problem that takes a bit more time to solve. To prevent this issue from slowing down the whole class too much, you could use the occasion to take the small break you had planned to take a bit later, giving the helper(s) time to fix the problem.

9.8   Sticky Notes to Distribute Attention

Sticky notes can also be used to ensure that the instructor's attention is fairly distributed. Have each learner write their name on a sticky note and put it on their laptop. Each time the instructor calls on them or answers one of their questions, their sticky note comes down. Once all the sticky notes are down, everyone puts theirs up again.

   This technique makes it easy for the instructor to see who they haven't spoken with recently, which in turn helps them avoid the un-conscious trap of only interacting with the most extroverted of their learners. It also shows learners that attention is being distributed fairly, so that when they *are* called on, they won't feel like they're being picked on.

9.9   Never Touch the Learner's Keyboard

It's often tempting to fix things for learners, but when you do, it can easily seem like magic (even if you narrate every step). Instead, talk your learners through whatever they need to do. It will take longer, but it's more likely to stick.

9.10 Repeat the Question

Whenever someone asks a question in class, repeat it back to them before answering it to check that you've understood it, and to give people who might not have heard it a chance to do so. This is particularly important when presentations are being recorded or broadcast, since your microphone will usually not pick up what other people are saying. Repeating questions back also gives you a chance to redirect the question to something you're more comfortable answering if need be...

9.11 Minute Cards

We frequently use sticky notes as *minute cards*: before each break, learners take a minute to write one positive thing on the green sticky note (e.g., one thing they've learned that they think will be useful), and one thing they found too fast, too slow, confusing, or irrelevant on the red one. They can use the red sticky note for questions that haven't yet been answered. While they are enjoying their coffee or lunch, the instructors review and cluster these to find patterns. It only takes a few minutes to see what learners are enjoying, what they still find confusing, what problems they're having, and what questions are still unanswered.

9.12 One Up, One Down

We frequently ask for summary feedback at the end of each day. The instructors ask the learners to alternately give one positive and one negative point about the day, without repeating anything that has already been said. This requirement forces people to say things they otherwise might not: once all the "safe" feedback has been given, participants will start saying what they really think.

Minute cards are anonymous; the alternating up-and-down feedback is not. Each mode has its strengths and weaknesses, and by providing both, we hope to get the best of both worlds.

9.13 Pair Programming

Pair programming is a software development practice in which two programmers share one computer. One person (called the driver) does the typing, while the other (called the navigator) offers comments and suggestions. The two switch roles several times per hour.

Pair programming is a good practice in real life, and also a good way to teach [Hannay2009], [Porter2013]. Partners can not only help each other out during the practical, but can also clarify each other's misconceptions when the solution is presented, and discuss common research interests during breaks. To facilitate this, we strongly prefer flat (dinner-style) seating to banked (theater-style) seating; this also makes it easier for helpers to reach learners who need assistance.

When pair programming is used it's important to put *everyone* in pairs, not just the learners who are struggling, so that no one feels singled out. It's also useful to have people sit in new places (and hence pair with different partners) after each coffee or meal break. It's also important to have people switch roles within each pair three or four times per hour, so that the stronger personality in each pair doesn't dominate the session.

Confidence is Not Knowledge

The Dunning-Kruger Effect[4] can easily come into play in pair programming: whoever thinks *they know the most can dominate the session regardless of how much they* actually *know.*

Switching Partners

Instructors have mixed opinions on whether people should be required to change partners at regular intervals. On the one hand, it gives everyone a chance to gain new insights and make new friends. On the other, it is uncomfortable for introverts, and moving computers and power adapters to new desks several times a day is disruptive.

9.14 Have Learners Make Predictions

Research has shown that people learn more from demonstrations if they are asked to predict what's going to happen [Miller2013]. Doing this fits naturally into live coding: after adding or changing a few lines of a program, ask someone what is going to happen when it's run.

9.15 Collaborative Debugging

If you are live coding and your program doesn't work, explain the symptoms to your learners. The underlying cause often then becomes clear; if it doesn't, have them take turns suggesting things

---

[4]https://en.wikipedia.org/wiki/Dunning%E2%80%93Kruger_effect

to try next. Be careful not to let one or two people dominate the discussion.

9.16   Peer Instruction

No matter how good a teacher is, she can only say one thing at a time. How then can she clear up many different misconceptions in a reasonable time?

The best solution developed so far is a technique called *peer instruction*. Originally created by Eric Mazur at Harvard, it has been studied extensively in a wide variety of contexts, including programming [Porter2013]. Peer instruction combines formative assessment with student discussion and looks something like this:

1. Give a brief introduction to the topic.

2. Give students a multiple choice question (MCQ) that probes for misconceptions (rather than simple factual knowledge).

3. Have all the students vote on their answers to the MCQ.

   - If the students all have the right answer, move on.
   - If they all have the same wrong answer, address that specific misconception.
   - If they have a mix of right and wrong answers, give them several minutes to discuss those answers with one another in small groups (typically 2-4 students) and then reconvene and vote again.

As [Avanti2013] shows, group discussion significantly improves students' understanding because it forces them to clarify their thinking, which can be enough to call out gaps in reasoning. Re-polling the class then lets the instructor know if they can move on, or if further explanation is necessary. A final round of additional explanation and discussion after the correct answer is presented gives students one more chance to solidify their understanding.

Peer instruction is essentially a way to provide one-to-one mentorship in a scalable way. Despite this, we usually do not use it in either programming workshops or instructor training workshops because it takes people time to adapt to a new way of learning–time that we typically don't have in our compressed two-day format.

Note that it is important to have learners record their votes so that they can't change their minds afterward and rationalize it by making excuses to themselves like "I just misread the question". Much of the value of peer instruction comes from the jarring experience of having their answer be wrong and having to think through the reasons why.

9.17   Setting Up Your Learners

Adult learners tell us that it is important to them to leave programming workshops with their own machine set up to do real work. We therefore strongly recommend that instructors be prepared to teach on all three major platforms (Linux, Mac OS, and Windows), even though it would be simpler to require learners to use just one.

To aid in this, put detailed setup instructions for all three platforms on the workshop's website, and email learners a couple of days before the workshop starts to remind them to do the setup. Even with this, a few people will always show up without the right software, either because their other commitments didn't allow them to go through the setup or because they ran into problems. To detect this, have everyone run some simple command as soon as they arrive and show the instructors the result, and then have helpers and other learners assist people who have run into trouble.

Common Denominators

If you have participants using several different operating systems, avoid using features which are OS-specific, and point out any that you do *use. For example, some shell commands take different options on Mac OS than on Linux, while the "minimize window" controls and behavior on Windows are different from those on other platforms.*

Virtual Machines

We have experimented with virtual machines (VMs) on learners' computers to reduce installation problems, but those introduce problems of their own: older or smaller machines simply aren't fast enough, and learners often struggle to switch back and forth between two different sets of keyboard shortcuts for things like copying and pasting.

Some instructors use VPS over SSH or web browser pages instead. This solve the installation issues, but makes us dependent on

*host institutions' WiFi (which can be of highly variable quality),
and has the issues mentioned above with things like keyboard
shortcuts.*

9.18   Setting Up Tables

You may not have any control over the layout of the desks or tables
in the room in which your programming workshop takes place, but
if you do, we find it's best to have:

- all tables on the same level (rather than banked seating),

- four people per table (so that each table can have two pairs if
  you choose to use pair programming), and

- in-floor power outlets (so that you don't have to run power cords
  across the floor that people can trip over).

Whatever layout you have, try to make sure the seats have good
back support, since people are going to be in them for an extended
period, and check that every seat has an unobstructed view of the
screen.

9.19   Setting Up Your Own Environment

Setting up your environment is just as important as setting up your
learners', but more involved. As well as having all the software that
they need, and network access to the tool they're using to take notes,
you should also have a glass of water, or a cup of tea or coffee. This
helps keep your throat lubricated (as discussed in the next section),
but its real purpose is to give you an excuse to pause for a couple of
seconds and think when someone asks a hard question or you lose
track of what you were going to say next.

You will probably also want some whiteboard pens and a few of
the other things described in the travel kit checklist (§21.5).

9.20   Cough Drops

If you talk all day to a room full of people, your throat gets raw
because you are irritating the epithelial cells in your larynx and
pharynx. This doesn't just make you hoarse–it also makes you more
vulnerable to infection (which is part of the reason people often come
down with colds after teaching).

The best way to protect yourself against this is to keep your throat lined, and the best way to do that is to use cough drops early and often. The right ones will also help delay the onset of coffee breath, for which your learners will probably be grateful.

9.21   Think-Pair-Share

Think-pair-share is a lightweight technique that helps refine their ideas and compare them with others'. Each person starts by thinking individually about a question or problem and jotting down a few notes. Participants are then paired to explain their ideas to each another, and possibly to merge them or select the more interesting ones. Finally, a few pairs present their ideas to the whole group.

Think-pair-share works because, to paraphrase Oscar Wilde's Lady Windermere, people often can't know what they're thinking until they've heard themselves say it. Pairing gives people new insight into their own thinking, and forces them to think through and resolve any gaps or contradictions *before* exposing their ideas to a larger group. We used think-pair-share in several of the challenges in our discussion of motivation (§1.5).

9.22   Humor

Humor should be used sparingly when teaching: most jokes are less funny when written down, and become even less funny with each re-reading. Being spontaneously funny while teaching usually works better, but can easily go wrong: what's a joke to your circle of friends may turn out to be a serious political issue to your audience. If you do make jokes when teaching, don't make them at the expense of any group, or of anyone except possibly yourself.

9.23   Challenges

Create a Questionnaire (20 minutes)

Using the questionnaire shown earlier as a template, create a short questionnaire you could give learners before teaching a class of your own. What do you most want to know about their background?

10 Teaching Online

Objectives

- *Learner can explain why expectations for massive online courses were unrealistic, and ways in which ad hoc use online learning is changing education.*
- *Learners can explain several key features of successful online courses, and judge whether a particular course meets those criteria.*
- *Learners can implement several logically different kinds of exercises using multiple-choice questions and/or write-and-run coding.*

If you use robots to teach, you teach people to be robots.
– variously attributed

Five years ago, you couldn't cross the street on a major university campus without hearing some talking about how teaching online was going to revolutionize education–or destroy it, or possibly both. Now that the hype has worn off, it's clear that while almost all learning now has an online component, MOOCs (massive open online courses) aren't as effective as their more enthusiastic proponents claimed they would be [Ubell2017]. As one specific example, [Koedinger2015] estimated "...the learning benefit from extra doing (1 SD increase) to be more than six times that of extra watching or reading." "Doing", in this case, refers to completing an interactive task with feedback, while "benefit" refers to both completion rates and overall performance.

One reason is that recorded content is ineffective for most novices learners because it cannot intervene to clear up specific learners' misconceptions. Some people happen to already have the right conceptual categories for a subject, or happen to form them correctly early on; these are the ones who stick with most massive online

courses, but many discussions of the effectiveness of such courses ignore this survivor bias.

Many enthusiasts also didn't realize that MOOCs and other forms of online education were just the latest in a long series of proposals to use machines to teach. As Audrey Watters has chronicled[1], every innovation in communications–from the printing press through radio and television to desktop computers and now mobile devices–has spawned a wave of aggressive optimists who believe that education is broken (without actually knowing much about education) and that the latest technology is the solution (without knowing much about what's been tried before, why it failed, or what success would actually look like) [Watters2014].

That said, technology *has* changed teaching and learning. Before blackboards were introduced into schools in the early 1800s, there was no way for a teacher to share an improvised example, diagram, or exercise with an entire class at once. Combining low cost, low maintenance, reliability, ease of use, and flexibility, blackboards enabled teachers to do things quickly and at scale that they had only been able to do slowly and piecemeal before. Similarly, the hand-held video camera revolutionized athletics training, just as the tape recorder revolutionized music instruction a decade earlier.

Today's revolutionary technology is the Internet, whose key characteristics are:

1. Students can access far more information, far more quickly, than ever before–provided, of course, that a search engine considers it worth indexing, that their internet service provider and government don't block it, and that the truth isn't drowned in a sea of attention-sapping disinformation.

2. Students can access far more people than ever before as well–again, provided that they aren't driven offline by harassment or marginalized because they don't conform to the social norms of whichever group is talking loudest.

3. Courses can reach far more people than before–but only if those students actually have access to the required technology, can afford to use it, and aren't being used as a way to redistribute wealth from the have-nots to the haves [Cottom2017].

4. Teachers can get far more detailed insight into how students work–so long as students are doing things that are amenable to large-scale automated analysis and aren't in a position to object to the use of surveillance into the classroom.

---

[1]http://teachingmachin.es/timeline.html

The caveats in this list are a big part of why this book exists. Right now, most of the discussion about using technology in teaching is shaped by what tech can do rather than what students need, and by Silicon Valley's craving for quick profits rather than by society's need for competent, well-informed citizens. If people in tech know more about teaching, they'll know more about what they should build to teach more effectively. And if they do that, they'll be able to reach more people from more diverse backgrounds, so that there will be more voices in the room when key decisions about the future of education are being made.

10.1   General Guidance

MOOCs *are* useful for people who already a mental model and wish to remind themselves of things they already know or fill in gaps in their knowledge. Most of today's online classes use some mix of recorded video presentations, automated exercise submission and grading, and web-based discussion, which can either be synchronous using chat tools and video calls or asynchronous using mailing lists and bulletin boards.

The two greatest strengths of this model are that learners can work when it's convenient for them, and that they have access to a wider range of courses, both because the Internet brings them all next door and because online courses typically have lower direct and indirect costs than in-person courses. The disadvantages are that they have to shoulder much more of the burden of staying focused, and that the impersonality of working online can demotivate people and encourage uncivil behavior.

Hybrid Vigor

The "pure MOOC" model described above is actually only a small subset of what could be done: [Brookfield2016] describes many other ways groups can discuss things, only a handful of which have ever been implemented online.) A hybrid approach that has proven quite successful is real-time remote instruction, in which the learners are co-located at one (or a few) sites, with helpers present, while the instructor(s) teaching via online video.

As an instructor teaching online, you should take advantages of the pros and do what you can to minimize or avoid the cons [Nilson2017,Miller2016]:

1. Deadlines should be frequent, well-publicized, and enforced, so that learners will get into a work rhythm.

2. Keep other all-class synchronous activities like live lectures to a minimum so that people don't miss things because of scheduling conflicts.

3. Encourage or require students to do some of their work in small groups (2-6 people) that *do* have synchronous activities such as a weekly online discussion. This will help students stay engaged and motivated without creating too many scheduling headaches.

4. Create, publicize, and enforce a code of conduct (§9.1) so that everyone can actually (as opposed to theoretically) take part in online discussions.

5. Remember that people learn best in small chunks (§4), so use lots of small lesson episodes rather than a handful of lecture-length chunks. Also remember that, disabilities aside, they can read faster than you can talk, so use video to engage rather than instruct. The one exception to this is that video is actually the best way to teach people verbs (actions), so use short screencasts to show people how to use an editor, step through code in a debugger, and so on.

6. Remember that the goal when teaching novices (§2) is to identify and clear up misconceptions. If early data shows that learners are struggling with some parts of a lesson, create extra alternative explanations of those points and extra exercises for them to practice on.

7. Remember that you are not the first person to record educational videos: everything from short guides[2] to entire books are there to help you.

Two-Way Video

Just as video lets an instructor show learners how she's doing things (rather than just what she has done), learners can use it to show instructors how they are working. If you are teaching programming using desktop or laptop computers, have your learners record a 5-minute screencast to show how they solved a problem. You can then watch that video at 4X or faster to see how proficient they are with the tools they're supposed to be using.

---

[2]https://www.adelaide.edu.au/learning/teaching/communities-of-practice/elearning/Guide_to_Creating_Educational_Videos.pdf

Freedom To and Freedom From

Isaiah Berlin's 1958 essay "Two Concepts of Liberty[3]" made a distinction between positive liberty, which is the ability to actually do something, and negative liberty, which is the absence of rules saying that you can't do it. Unchecked, online discussions usually offer negative liberty (nobody's stopping you from saying what you think) but not positive liberty (many people can't actually be heard). One way to address this is to introduce some kind of throttling, *such as a rule that says each learner is only allowed to contribute one message per discussion thread per day. Doing this allows those who have something to say to say it, while clearing space for others to say things as well.*

When it comes to teaching platforms, you can either use an all-in-one learning management system (LMS) such as Moodle[4] or assemble something ad hoc: Slack[5] or Zulip[6] for chat, Google Hangouts[7] for video conversations, and WordPress[8], Google Docs[9], or any number of wikis for collaborative authoring. If you are just starting out, then use whatever requires the least installation and administration on your side, and the least extra learning effort on your learners' side. (I once taught a half-day class using group text messages because that was the only tool everyone was already familiar with.)

The most important thing when choosing technology is to *ask your learners what they are already using.* Most people don't use IRC[10], and find its arcane conventions and interface offputting. Similarly, while this book lives in a GitHub[11] repository, requiring non-experts to submit exercises as pull requests proved to be an unmitigated disaster, even with its supposedly easy-to-use web editing tools. As an instructor, you're asking people to learn a lot; the least you can do in return is learn how to use their preferred tools.

Points for Improvement

One way to demonstrate to learners that they are learning with *you, not just* from *you, is to allow them to edit your course notes. In live courses, we recommend that you enable them to do this as*

---

[3]https://en.wikipedia.org/wiki/Two_Concepts_of_Liberty
[4]http://moodle.org
[5]http://slack.com
[6]https://zulipchat.com/
[7]http://hangouts.google.com
[8]https://wordpress.org/
[9]http://docs.google.com
[10]https://en.wikipedia.org/wiki/Internet_Relay_Chat
[11]http://github.com

you lecture (§9.5); in online courses, you can put your nodes into a wiki, a GitHub[12] repository, or anything else so long as it allows you to review and comment on their proposed changes, and the learner to make revisions based on your feedback, before those changes go live. Giving credit (or at least thanks) to people for fixing mistakes, clarifying explanations, adding new examples, and writing new exercises does increase the short-term load on the instructor, but reduces the lesson's long-term maintenance costs (§6.3).

A major concern with any online community, learning or otherwise, is how to actually make it a community. Hundreds of books and presentations discuss this, but almost all are based on their authors' personal experiences. [Kraut2012] is a welcome exception: while it predates the accelerating decline of Twitter and Facebook into weaponized abuse and misinformation, most of what was true then is true now. [Fogel2017] is also full of useful tips for the community of practice that students may hope to join.

One other concern people often have about teaching online is cheating. Day-to-day dishonesty is no more common in online classes than in face-to-face settings, but the temptation to have someone else write the final exam, and the difficulty of checking whether this happened, is one of the reasons educational institutions have been reluctant to offer credit for pure online classes. Remote exam proctoring is possible, usually by watching the student take the exam using a webcam. Before investing in this, read [Lang2013], which explores why and how students cheat, and how courses often give them incentives to do so.

10.2   Different Types of Exercises

Every mechanic has her favorite screwdrivers, and every good teacher has different kinds of exercises to check that her students are actually learning, let them practice their new skills, and keep them engaged. Some types of exercise are well known, but others aren't as widely used as they should be.

The two key requirements are that an exercise has to be quick for learners to do, and it has to be possible to check the answer automatically. These requirements rule out some useful kinds of assessment, but many remain.

The first type of exercise that works well online is a *multiple choice question* that presents a question and asks the student to

---

[12]http://github.com

pick the correct answer from a list. Doing this might (in fact, should) require them to do more than just read and remember, and as a previous post discussed, multiple-choice questions are most effective when their wrong answers probe for specific misconceptions on the student's part.

Example: You are in `/home/marg`. Use `ls` with an appropriate argument to get a listing of the files in the directory `/home/marg/seasonal`. Which of the following files is not *in that directory?*

- *`autumn.csv`*
- *`fall.csv`*
- *`spring.csv`*
- *`winter.csv`*

The second type of exercise is *write and run*, in which the student has to write code that produces a specified output. When the code is submitted, we check its structure and/or output and give feedback. Write and run exercises can be as simple or as complex as the instructor wants. For example, it's often enough with novices to simply ask them to call a function or method: experienced instructors often forget how hard it can be to figure out which parameters go where.

Example: the matrix M contains data read from a file. Using one function or method call, create a matrix Z that has the same shape as M but contains only zeroes.

Write and run exercises help students practice the skills they most want to learn, but writing good automated checks is hard: students can find very creative ways to get the right answer, and it's demoralizing to give them a false negative. One way to reduce how often this occurs is to give them a small test suite they can run their code against before they submit it (at which point it is run against a more comprehensive set of tests). Doing this helps to catch cases in which students have completely misunderstood the written spec of the exercise.

To help students realize just how hard it is to write good tests instructors can get them to do it themselves. Instead of writing code that satisfies some specification, they can be asked to write tests to determine whether a piece of code conforms to a spec.

Example: the function `monotonic_sum` calculates the sum of each section of a list of numbers in which the values are monotonically increasing. For example, given the input `[1, 3, 3, 4, 5, 1]`,

the output should be [4, 3, 9, 1]. Write and run unit tests to determine which of the following bugs the function contains:

- *Considers every negative number the start of a new sub-sequence.*
- *Does not include the first value of each sub-sequence in the sub-sum.*
- *Does not include the last value of each sub-sequence in the sub-sum.*
- *Only re-starts the sum when values decrease rather than fail to increase.*

Fill in the blanks is a refinement of write and run in which the student is given some starter code and asked to complete it. (In practice, many write and run exercises are actually fill in the blanks because the instructor will provide comments to remind the students of what steps they should take.) Novices often find fill in the blanks less intimidating than writing all the code from scratch, and since the instructor has provided most of the answer's structure, submissions are much easier to check.

Example: fill in the blanks so that the code below prints the string 'hat'.

```
text = 'all that it is'
slice = text[____:____]
print(slice)
```

As described earlier (§5), a *Parsons Problem* is another kind of exercise that avoids the "blank screen of terror" problem: the student is given the lines of code needed to solve a problem, but has to put them in the right order. Research over the past few years has shown that Parsons Problems are effective because they allow students to concentrate on control flow separately from vocabulary [Ericons2017]. The same research shows that giving the student more lines than she needs, or asking her to rearrange some lines and add a few more, makes this kind of problem significantly harder [Harms2017]. Parsons Problems can be emulated (albeit somewhat clumsily) by asking students to rearrange code in an editor.

Example: rearrange and indent these lines to calculate the sums of the positive and negative values in a list.

```
positive = 0
return negative, positive
if v > 0
else
positive += v
negative = 0
for v in values
negative += v
```

Tracing execution is the inverse of a Parsons Problem: given a few lines of code, the student has to trace the order in which those lines are executed. This is an essential debugging skill, and is a good way to solidify students' understanding of loops, conditionals, and the evaluation order of function and method calls. Again, we don't yet support this directly, but it can be emulated by having students type in a list of line labels.

Example: in what order are the labelled lines in this block of code executed?

```
A)      vals = [-1, 0, 1]
B)      inverse_sum = 0
        try:
            for v in vals:
C)              inverse_sum += 1/v
        except:
D)          pass
```

Tracing values is similar to tracing execution, but instead of spelling out the order in which code is executed, the student is asked to list the values that one or more variables take on as the program runs. Again, it can be implemented by having students type in their answers, but this quickly becomes impractical. In practice, the best approach is to give the student a table whose columns are labelled with variable names and whose rows are labelled with line numbers.

Example: what lines of text pass through the pipes and the final redirect when this file:

```
2017-11-01,Akeratu,9
2017-11-01,Monona,3
2017-11-02,Monona,1
2017-11-03,Monona,1
2017-11-03,Akeratu,7
```

is run through this Unix shell command:

```
cut -d , -f 2 filename | sort | uniq > result.txt
```

Returning to debugging skills, another exercise that helps student develop them is *minimal fixes*. Given a few lines of code that contain a bug, the student must either make or identify the smallest change that will produce the correct output. Making the change can be done as using write and run, while identifying it can be done as a multiple choice question.

Example: this function is supposed to test whether a point (x, y) lies strictly within a rectangle defined by (x_min, y_min, x_max, y_max). Change one line to make it do so correctly.

```
def inside(point, rect):
    if (point.x <= rect.x_min): return false
    if (point.y <= rect.y_min): return false
    if (point.x >= rect.y_max): return false
    if (point.y >= rect.y_max): return false
    return true
```

Theme and variation exercises are similar, but instead of making a change to fix a bug, the student is asked to make a small alteration that changes the output in some specific way. These alterations can include:

- replacing one function call with another
- changing one variable's initial value
- swapping an inner and outer loop
- changing the order of tests in a chain of conditionals
- changing the nesting of function calls or the order in which methods are chained

Again, this gives students a chance to practice a useful real-world skill: the fastest way to produce a working program is often to tweak one that already does something useful.

Example: change the inner loop control in the function below so that it sets the upper left triangle of the matrix to zero.

```
def zeroTriangle(matrix):
    for c in range(matrix.cols):
        for r in range(matrix.rows):
            matrix[r, c] = 0
```

Matching problems are another entire family of exercises. *One-to-one matching* gives the student two lists of equal length and asks her to pair corresponding items, e.g., "match each piece of code with the output it produces".

Example: match each function's name with the operation it implements.

Function	Operation
SGEMV	triangular banded matrix-vector multiply
STBMV	solve triangular matrix with multiple right-hand sides
STRSM	matrix-vector multiply

Many-to-many matching is similar, but the lists aren't the same length, so some items may be matched to several others. Both kinds require students to use higher-order thinking skills, but many-to-many are more difficult because students can't do easy matches first to reduce their search space. (In technical terms, there is a higher cognitive load (§5).)

Matching problems can be emulated by having students submit lists of pairs as text (such as "A3, B1, C2"), but that's clumsy and error-prone. If available, the machinery built for Parsons Problems can be used let students drag and drop blocks of text to form matches.

Drag-and-drop opens many other doors: for example, tracing execution is easy to implement this way. So is *labelling diagrams*: rather than students typing in the labels, it is faster and more reliable for them to drag labels around to attach to the correct elements. The picture can be a complex data structure ("after this code is executed, which variables point to which parts of this structure?"), the graph that a program produces ("match each of these pieces of code with the part of the graph it generated"), the code itself ("match each term to an example of that program element"), or many other things.

Example: label the following diagram to show which structures the variables x, y, and z refer to after these three lines of code are executed.

```
x = 3
y = [x, x]
z = [x, y]
```

Figure 10.1: Labelling

Drawing diagrams of things like data structures is also straight-forward to do on paper but very difficult to grade automatically. One way to make solutions gradable may be to constrain the drawing in the same way that Parsons Problems constrain code construction, i.e., give students the pieces of the diagram and ask them to arrange them correctly, but this is a long way off.

We mentioned earlier that matching problems require students to use higher-order thinking skills. *Summarization* also does this, and gives them a chance to practice a skill that is very useful when *reporting* bugs rather than fixing them. For example, students can be asked, "Which sentence best describes how the output of f changes as x varies from 0 to 10?" and then given several options as a multiple choice question. Similarly, ranking problems present the student with several choices and ask them to order them from fastest to slowest, most robust to most brittle, and so on. (Ranking is more manageable when implemented with drag and drop than as a multiple choice question.)

One other kind of exercise that can be implemented as a multiple choice question is *fault mapping*: given a piece of buggy code and an error message, the student has to identify the line on which the error occurred. In simple cases this will be the line mentioned in the error message, but in more subtle cases, the student will have to trace execution forward and backward to figure out where things first went wrong.

Other kinds of exercises are hard for any automated platform to provide. *Refactoring exercises* are the complement of theme and variation exercises: given a working piece of code, the student has to modify it in some way *without* changing its output. For example, the student could be asked to replace loops with vectorized expressions, to simplify the condition in a while loop, etc. The challenge here is that there are often so many ways to refactor a piece of code that

grading requires human intervention.

Example: write a single list comprehension that has the same effect as this loop.

```
result = []
for v in values:
    if len(v) > threshold:
        result.append(v)
```

Code review is hard to grade automatically in the general case, but can be tackled if the student is given a rubric (i.e., a list of faults to look for) and asked to match particular comments against particular lines of code. For example, the student can be told that there are two indentation errors and one bad variable name, and asked to point them out; if she is more advanced, she could be given half a dozen kinds of remarks she could make about the code without guidance as to how many of each she should find. As with tracing values, this is easiest for students to do when presented as a table, which we currently don't support.

Example: using the rubric provided, mark each line of the code below.

```
01)  def addem(f):
02)      x1 = open(f).readlines()
03)      x2 = [x for x in x1 if x.strip()]
04)      changes = 0
05)      for v in x2:
06)          print('total', total)
07)          tot = tot + int(v)
08)      print('total')
```

1. *poor variable name*
2. *unused variable*
3. *use of undefined variable*
4. *missing values*
5. *fossil code*

All of this discussion has assumed that grading must be fully automatic in order to scale to large classes, but that is not necessarily true. [Paré2008] and [Kulkarni2013] report experiments in which learners grade each other's work, and the grades they assign are then compared with grades given by graduate-level teaching assistants

or other experts. Both found that student-assigned grades agreed with expert-assigned grades as often as the experts' grades agreed with each other, and that a few simple steps (such as filtering out obviously unconsidered responses or structuring rubrics) decreased disagreement even further. Much more research needs to be done, but given that critical reading is an effective way to learn, this result may point to a future in which learners use technology to make judgments, rather than being judged by technology.

10.3   Challenges

Give Feedback (20 minutes)

1. Watch [Wilson2017] as a group and give feedback on it. Organize feedback along two axes: positive vs. negative and content vs. presentation.

2. Have each person in the class add one point to a 2x2 grid on a whiteboard (or in the shared notes) without duplicating any points that are already up there.

What did other people see that you missed? What did they think that you strongly agree or disagree with? (You can compare your answers with these lists (§17.5).)

Adapting Multiple Choices Questions (30 minutes)

Pick one of the examples given in this chapter of using multiple choice questions to implement some other kind of online programming exercise, create an example, and swap with one of your fellow learners.

Adapting Write and Run Exercises (30 minutes)

Pick one of the examples given in this chapter of using write and runs exericses to implement some other kind of online programming exercise, create an example, and swap with one of your fellow learners.

11 Building Community

Objectives

- *Learners can make an informed decision about whether to create something new, or join an existing effort.*
- *Learners can judge whether a meeting is well organized and well run or not.*
- *Learners can outline a three-step plan for recruiting, retaining, and retiring volunteers and other organization participants.*
- *Learners can explain the difference between a service board and a governance board, and judge which kind an organization has.*

Many well-intentioned people want the world to be a better place, but don't actually want anything important to change. A lot of grassroots efforts to teach programming fall into this category: they want to teach children and adults how to program so that they can get good jobs, rather than empower them to change the system that has shut them (and people like them) out of those jobs in the past.

If you are going to build a community, the first and most important thing you have to decide is what *you* want: to help people succeed in the world we have, or to give them a way to make a better one. If you choose the latter, you have to accept that one person can only do so much. Just as we learn best together, we teach best when we are teaching with other people, and the best way to achieve that is to build a community.

11.1 Learn, Then Do

The first step in building a community is to decide if you really need to, or whether you would be more effective joining an existing organization. Thousands of groups are already teaching people tech

skills, from the 4-H Club[1] and literacy programs[2] to get-into-coding non-profits like Black Girls Code[3]. Joining an existing group will give you a head start on teaching, an immediate set of colleagues, and a chance to learn more about how to run things. The only thing it *won't* give you is the ego gratification and control that comes from being a founder.

Whether you join an existing group or set up one of your own, you owe it to yourself and everyone who's going to work with you to find out what's been done before. People have been writing about grassroots organizing for decades; [Alinsky1989] is probably the best-known work on the subject, while [Brown2007] and [Midwest2010] are practical manuals rooted in decades of practice. If you want to read more deeply, [Adams1975] is a history of the Highlander Folk School, whose approach has been emulated by many successful groups, while [Spalding2014] is a guide to teaching adults written by someone with deep personal roots in organizing.

Meetings, Meetings, Meetings

Knowing how to run a meeting efficiently is *the* core skills of community organizers. (Knowing how to take part in someone else's meeting is just as important, but gets far less attention–as a colleague once said, everyone offers leadership training, nobody offers followership training.) The most important rules for making meetings efficient are not secret, but are rarely followed:

1. *Decide if there actually needs to be a meeting.* If the only purpose is to share information, have everyone send a brief email instead.

2. *Choose a chair.* A meeting without a chair is going to work just about as well as an orchestra without a conductor.

3. *Write down the agenda, with timings.* List the items that are to be discussed and estimate the time allowed for each. Your first estimates with any new group will be wildly optimistic, so revise them upward for subsequent meetings.

4. *No technology.* Insist that everyone put their phones, tablets, and laptops into politeness mode (i.e., closes them).

5. *Record minutes.* Someone other than the chair should take point-form notes about the most important pieces of information that

---

[1] http://www.4-h-canada.ca/
[2] https://www.frontiercollege.ca/
[3] http://www.blackgirlscode.com/

were shared, and about every decision that was made or every task that was assigned to someone.

6. *Make sure the chair doesn't speak more than anyone else.* The chair is there to make sure that the meeting keeps moving, *not* to do most of the talking.

7. *Require politeness.* No one gets to interrupt anyone, no one gets to ramble, and if someone goes off topic, it's the chair's job to say, "Let's discuss that elsewhere."

8. *End early.* If your meeting is scheduled for 10:00-11:00, you should aim to end at 10:55 to give people time to get where they need to go next.

[Brown2007] and [Brookfield2016] have lots of good advice on running meetings, and if you want to "learn, then do", an hour of training on chairing meetings is the most effective place to start.

Sticky Notes and Interruption Bingo

Some people are so used to the sound of their own voice that they will insist on talking half the time no matter how many other people are in the room. One way to combat this is to give everyone three sticky notes at the start of the meeting. Every time they speak, they have to take down one sticky note. When they're out of notes, they aren't allowed to speak until everyone has used at least one, at which point everyone gets all of their sticky notes back. This ensures that nobody talks more than three times as often as the quietest person in the meeting, and completely changes the dynamics of most groups: people who have given up trying to be heard because they always get trampled suddenly have space to contribute, and the overly-frequent speakers quickly realize just how unfair they have been.

Another useful technique is called interruption bingo. Draw a grid, and label the rows and columns with the participants' names. Each time someone interrupts someone else, add a tally mark to the appropriate cell. Halfway through the meeting, take a moment to look at the results. In most cases, you will see that one or two people are doing all of the interrupting, often without being aware of it. After that, saying, "All right, I'm adding another tally to the bingo card," is often enough to get them to throttle back.

Dealing With Disruptors

The National Oceanic and Atmospheric Administraiton (NOAA) has published a short, amusing, and above all useful guide to

dealing with disruptive behaviors[4]. It categorizes those behaviors under labels like "talkative", "indecisive", and "shy", and outlines strategies for handling each.

11.2   Three Steps

- *Me in 2012: I'm not going to worry about retaining volunteers until I have a few.*
- *My Dad: If you don't think about how you're going to keep them, you probably won't get any.*

Everyone who gets involved with your organization, including you, goes through three phases: recruitment, retention, and retirement (from the organization). You don't need to worry about this cycle when you're just getting started, but it *is* worth thinking about as soon as you have more than a couple of non-founders involved.

The first step is recruiting volunteers. Your marketing (§12) should help you with this by making your organization findable, and by making its mission and its value to volunteers clear to people who might want to get involved. Share stories that exemplify the kind of help you want as well as stories about the people you're helping, and make it clear that there are many ways to get involved. (We discuss this in more detail in the next section.)

Your best source of new recruits is your own classes: "see one, do one, teach one" has worked well for volunteer organizations for as long as there have *been* volunteer organizations. Make sure that every class or other encounter ends with two sentences explaining how people can help, and that help is welcome. People who come to you this way will know what you do, and will have recent experience of being on the receiving end of what you offer that they can draw on, which helps your organization avoid collective expert blind spot.

Start Small

As Ben Franklin[5] observed, a person who has performed a favor for someone is more likely to do another favor for that person than they would be if they had received a favor from that person. Asking people to do something small for you is therefore a good step toward getting them to do something larger. One natural way to do this when teaching is to ask people to submit fixes for your lesson materials for typos or unclear wording, or to suggest new exercises or examples. If your materials are written in a

---

[4]https://coast.noaa.gov/ddb/story_html5.html
[5]https://en.wikipedia.org/wiki/Ben_Franklin_effect

maintainable way (§6.3), this gives them a chance to practice some useful skills, and gives you an opportunity to start a conversation that might lead to a new recruit.

Recruiting doesn't end when someone first shows up: if you don't follow through, people will come out once or twice, then decide that what you're doing isn't for them and disappear. Two things you can do to get newcomers over this initial hump are:

1. Have them take part in group activities before they do anything on their own, both so that they get a sense of how your organization does things, and so that they build social ties that will keep them involved.

2. Give newcomers a mentor, and make sure the mentors actually do some proactive mentoring. The most important things a mentor can do are make introductions and explain the unwritten rules, so make it clear to mentors that these are their primary responsibilities, and they are to report back to you every few weeks to tell you what they've done.

The second part of the volunteer lifecycle is retention, which is a large enough topic to deserve its own section. The third and final part is retirement. Sooner or later, everyone moves on (including you). When this happens:

1. Ask people to be explicit about their departure.

2. Make sure they don't feel embarrassed or ashamed about leaving.

3. Give them an opportunity to pass on their knowledge. For example, you can ask them to mentor someone for a few weeks as their last contribution, or to be interviewed by someone who's staying with the organization to collect any stories that are worth re-telling.

4. Make sure they hand over the keys. It's awkward to discover six months after someone has left that they're the only person who knows how to book a playing field for the annual softball game.

5. Follow up 2-3 months after they leave to see if they have any further thoughts about what worked and what didn't while they were with you, or any advice to offer that they either didn't think to give or were uncomfortable giving on their way out the door.

6. Thank them, both when they leave and the next time your group gets together.

11.3 Retention

The community organizer Saul Alinksy said, "If your people aren't having a ball doing it, there is something very wrong." Community members shouldn't expect to enjoy every moment of their work with your organization, but if they don't enjoy any of it, they won't stay.

Enjoyment doesn't necessarily mean having an annual party: people may enjoy cooking, coaching, or just working quietly beside others. There are several things every organization should do to ensure that people are getting something they value out of their work:

1. *Ask people what they want rather than guessing.* Just as you are not your learners (§6.4), you are probably different from other members of your organization. Ask people what they want to do, what they're comfortable doing (which may not be the same thing), what constraints there are on their time, and so on.

2. *Provide many ways to contribute.* The more ways there are for people to help, the more people will be able to help. Someone who doesn't like standing in front of an audience may be able to maintain your organization's website or handle its accounts; someone who doesn't know how to do anything else may be able to proof-read lessons, and so on. The more kinds of tasks you do yourself, the fewer opportunities there are for others to get involved.

3. *Recognize contributions.* Everyone likes to be appreciated, so communities should acknowledge their members' contributions both publicly and privately.

4. *Make space.* Micromanaging or trying to control everything centrally means people won't feel they have the autonomy to act, which will probably cause them to drift away. In particular, if you're too engaged or too quick on the reply button, people have less opportunity to grow as members and to create horizontal collaborations. The community can continue to be "hub and spoke", focused around one or two individuals, rather than a highly-connected network in which others feel comfortable participating.

Another way to make participation rewarding is to provide training. Organizations require committees, meetings, budgets, grant proposals, and dispute resolution; most people are never taught how to do any of this, any more than they are taught how to teach,

but training people to do these things helps your organization run more smoothly, and the opportunity to gain transferable skills is a powerful reason for people to get and stay involved. If you are going to do this, don't try to provide the training yourself (unless it's what you specialize in). Many civic and community groups have programs of this kind, and you can probably make a deal with one of them.

Other groups may be useful in other ways as well, and you may be useful to them–if not immediately, then tomorrow or next year. You should therefore set aside an hour or two every month to find allies and maintain your relationships with them. One way to do this is to ask them for advice: how do they think you ought to raise awareness of what you're doing? Where have they found space to run classes? What needs do they think aren't being met, and would you be able to meet them (either on your own, or in partnership with them)? Any group that has been around for a few years will have useful advice; they will also be flattered to be asked, and will know who you are the next time you call.

Government Matters

Have you ever spoken to someone from the public relations office at your local college or school board, or in your city councilor's office? If not, what are you waiting for?

Soup, Then Hymns

Manifestos are fun to write, but most people join a volunteer community to help and be helped rather than to argue over the wording of a grand vision statement. (Most people who prefer the latter are only *interested in arguing. . .) To be effective you should therefore focus on things that are immediately useful, e.g., on what people can create that will be used by other community members right away. Once your organization shows that it can actually achieve things–even small things–people will be more confident that it's worth thinking about bigger issues.*

One important special case of making things rewarding is to pay people. Volunteers can do a lot, but eventually tasks like system administration and accounting need full-time paid staff. This is often a difficult transition for grassroots organizations, but is outside the scope of this book.

11.4 Governance

As Jo Freeman pointed out in her influential essay "The Tyranny of Structurelessness[6]", every organization has a power structure: the only question is whether it's formal and accountable, or informal and unaccountable. Make yours one of the first kind: write and publish the rules governing everything from who's allowed to use the name and logo to who gets to decide whether people are allowed to charge money to teach with whatever materials your group has worked up.

Organizations can govern themselves in many different ways, and a full discussion of the options is outside the scope of this book. The most important thing to keep in mind is that countries and corporations are only two of many governance models, and that a commons is often a better model for volunteer teaching organizations. A commons is "something managed jointly by a community according to rules they themselves have evolved and adopted"; as [Bollier2014] emphasizes, all three parts of that definition are essential: a commons isn't just a shared pasture, but also includes the community that shares it and the rules they use to do so.

Most resources, throughout most of human history, have been commons: it is only in the last few hundred years that impersonal markets have pushed them to the margins. In order to do so, free-market advocates have had to convince us we're something we're not (dispassionate calculators of individual advantage) and erase or devalue local knowledge and custom. Both have had tragic consequences for us individually and communally, and now for our whole planet.

Since society has difficulty recognizing commons organizations, and since most of the people you will want to recruit don't have experience with them, you will probably wind up having some sort of board, a director, and other staff. Broadly speaking, your organization can have either a *service board*, whose members also take on other roles in the organization, or a *governance board* whose primary responsibility is to hire, monitor, and if need be fire the director. Board members can be elected by the community or appointed; in either case, it's important to prioritize competence over passion (the latter being more important for the rank and file), and to try to recruit for particular skills such as accounting, marketing, and so on.

Don't worry about drafting a constitution when you first get started: it will only result in endless wrangling about what we're

---

[6]http://www.jofreeman.com/joreen/tyranny.htm

going to do rather than formalization of what you're already doing. When the time does come to formalize your rules, though, make your organization a democracy: sooner or later (usually sooner), every appointed board turns into a mutual agreement society and loses sight of what the community it's meant to serve actually needs. Giving the community power is messy, but is the only way invented so far to ensure that an organization continues to meet people's actual needs.

11.5   Final Thoughts

As [Pigni2016] discusses, burnout is a chronic risk in any community activity. If you don't take care of yourself, you won't be able to take care of your community.

Every organization eventually needs fresh ideas and fresh leadership. When that time comes, train your successors and then move on. They will undoubtedly do things you wouldn't have, but the same is true of every generation. Few things in life are as satisfying as watching something you helped build take on a life of its own. Celebrate that–you won't have any trouble finding something else to keep you busy.

11.6   Challenges

Several of these exercises are taken from [Brown2007], which is an exceptionally useful book on building community organizations.

Who Are You? (20 minutes)

Read NOAA's descriptions of disruptive behaviors[7].

1. Which of them best describes you?
2. Do you think the strategies described for handling people like you are effective?
3. Are other strategies equally or more effective?

People You May Meet (30 minutes)

As an organizer, part of your job is sometimes to help people find a way to contribute despite themselves. In small groups, pick three of the people below and discuss how you would help them become a better contributor to your organization.

---

[7]https://coast.noaa.gov/ddb/story_html5.html

- *Anna* knows more about every subject than everyone else put together–at least, she thinks she does. No matter what you say, she'll correct you; no matter what you know, she knows better.

- *Catherine* has so little confidence in her own ability that she won't make any decision, no matter how small, until she has checked with someone else.

- *Frank* believes that knowledge is power, and enjoys knowing things that other people don't. He can make things work, but when asked how he did it, he'll grin and say, "Oh, I'm sure you can figure it out."

- *Hediyeh* is quiet. She never speaks up in meetings, even when she knows that what other people are saying is wrong. She might contribute to the mailing list, but she's very sensitive to criticism, and will always back down rather than defending her point of view.

- *Kenny* has discovered that most people would rather shoulder his share of the work than complain about him, and he takes advantage of it at every turn. The frustrating thing is that he's so damn *plausible* when someone finally does confront him. "There have been mistakes on all sides," he says, or, "Well, I think you're nit-picking."

- *Melissa* means well, but somehow something always comes up, and her tasks are never finished until the last possible moment. Of course, that means that everyone who is depending on her can't do their work until *after* the last possible moment. . .

- *Raj* is rude. "It's just the way I talk," he says, "If you can't hack it, maybe you should find another team." His favorite phrase is, "That's stupid," and he uses obscenity as casually as minor characters in Tarantino films.

Values (45 minutes)

Answer the following questions on your own, and then compare your answers to those given by other members of your group.

1. What are the values your organization expresses?
2. Are these the values you want the organization to express?
3. If not, what values would you like it to express?
4. What are the specific behaviors that demonstrate those values?

5. What are some key behaviors that would demonstrate the values you would like for your group?
6. What are the behaviors that would demonstrate the opposite of those values?
7. What are some key behaviors that would demonstrate the opposite of the values you want to have?

Meeting Procedures (30 minutes)

Answer the following questions on your own, and then compare your answers to those given by other members of your group.

1. How are your meetings run?
2. Is this how you want your meetings to be run?
3. Are the rules for running meetings explicit or just assumed?
4. Are these the rules you want?
5. Who is eligible to vote/make decisions?
6. Is this who you want to be vested with decision-making authority?
7. Do you use majority rule, make decisions by consensus, or use some other method?
8. Is this the way you want to make decisions?
9. How do people in a meeting know when a decision has been made?
10. How do people who weren't at a meeting know what decisions were made?
11. Is this working for your group?

Size (20 minutes)

Answer the following questions on your own, and then compare your answers to those given by other members of your group.

1. How big is your group?
2. Is this the size you want for your organization?
3. If not, what size would you like it to be?
4. Do you have any limits on the size of membership?
5. Would you benefit from setting such a limit?

Staffing (30 minutes)

Answer the following questions on your own, and then compare your answers to those given by other members of your group.

1. Do you have paid staff in your organization?

2. Or is it all-volunteer?
3. Should you have paid staff?
4. Do you want/need more or less staff?
5. What do you call the staff (e.g., organizer, director, coordinator, etc.)?
6. What do the staff members do?
7. Are these the primary roles and functions that you want the staff to be filling?
8. Who supervises your staff?
9. Is this the supervision process and responsibility chain that you want for your group?
10. What is your staff paid?
11. Is this the right salary to get the needed work done and to fit within your resource constraints?
12. What benefits does your group provide to its staff (health, dental, pension, short and long-term disability, vacation, comp time, etc.)?
13. Are these the benefits that you want to give?

Collaborations (30 minutes)

Answer the following questions on your own, and then compare your answers to those given by other members of your group.

1. Do you have any agreements or relationships with other groups?
2. Do you want to have relationships with any other groups?
3. How would having (or not having) collaborations help you to achieve your goals?
4. What are your key collaborative relationships?
5. Are these the right collaborators for achieving your goals?
6. With what groups or entities would you like your organization to have agreements or relationships?

Money (30 minutes)

Answer the following questions on your own, and then compare your answers to those given by other members of your group.

1. Who pays for what?
2. Is this who you want to be paying?
3. Where do you get your money?
4. Is this how you want to get your money?
5. If not, do you have any plans to get it another way?
6. If so, what are they?

7. Who is following up to make sure that happens?
8. How much money do you have?
9. How much do you need?
10. What do you spend most of your money on?
11. Is this how you want to spend your money?

Becoming a Member (45 minutes)

Answer the following questions on your own, and then compare your answers to those given by other members of your group.

1. How does someone join?
2. Does this process work for your organization?
3. What are the membership criteria?
4. Are these the membership criteria you want?
5. Are people required to agree to any rules of behavior upon joining?
6. Are these the rules for behavior you want?
7. Are there membership dues?

12 Marketing

Objectives

- *Learners can explain what marketing actually is.*
- *Learners can clearly explain the value of what they are offering to different potential stakeholders.*
- *Learners can explain what a brand is and determine whether they or their organization have one.*

It's hard to get people with technical backgrounds to think about marketing, not least because it's perceived as being about spin and misdirection. In reality, *marketing* is the craft of seeing things from other people's perspective, understanding their wants and needs, and finding ways to meet them. This should sound familiar: many of the techniques introduced earlier in this book are intended to do exactly this for lessons. This chapter will look at how to apply similar ideas to the larger problem of getting people to support the work you're trying to do.

12.1 What Are You Offering to Whom?

The first step is to figure out what you are offering to whom, i.e., what actually brings in the funding and other support you need to keep operating. As [Kuchner2011] points out, the answer is often counter-intuitive. For example, most scientists think their product is papers, but their actual product is their grant proposals, because those are what brings in money. Their papers are the advertising that persuades people to buy (fund) those proposals, just as albums are now the advertising that persuades people to buy musicians' concert tickets and t-shirts.

You may or may not be a scientist, so suppose instead that your group is offering weekend programming workshops to people who are re-entering the workforce after taking several years out to look

after young children. If your learners are paying enough for your workshops to cover your costs, then the learners are your customers and the workshops are the product. If, on the other hand, the workshops are free, or the learners are only paying a token amount (to cut the no-show rate), then your actual product may be some mix of:

- your grant proposals,
- the alumni of your workshops that the companies sponsoring you would like to hire,
- the half page summary of your work in the mayor's annual report to city council that shows how she's supporting the local tech sector, or
- the personal satisfaction that teaching gives your volunteer instructors.

As with our recommended lesson design process (§20), you should try to identify specific people who might be interested in what you're doing and figure out which of *their* needs *your* program will meet. Personas (§6.1) are one way to do this. Another is to write a set of *elevator pitches*, each aimed at a different stakeholder. A widely-used template for these pitches looks like this:

1. For *target audience*
2. who *dissatisfaction with what's currently available*
3. our *category*
4. provide *key benefit.*
5. Unlike *alternatives*
6. our program *key distinguishing feature.*

Continuing with our weekend workshop example, we might use this for potential attendees:

For people re-entering the workforce after taking time out to raise children *who* still have regular childcare responsibilities*, our* introductory programming workshops *provide* weekend classes with on-site childcare. *Unlike* online classes, *our program* gives participants a chance to meet people who are at the same stage of life.

but use this to characterize the companies that we would like to donate staff time for teaching:

For a company that wants to recruit entry-level software developers *that* is struggling to find mature, diverse candidates

our introductory programming workshops *provide* a pool of potential recruits in their thirties that includes large numbers of people from underrepresented groups. *Unlike* college recruiting fairs, *our program* connects companies directly with a diverse audience.

If you don't know why different potential stakeholders might be interested in what you're doing, ask them. If you do know, ask them anyway: answers can change over time, and it's a good way to discover things that you might have missed. Once you have written these pitches, you should use them to drive what you put on your organization's web site and in other publicity material, since it will help people figure out as quickly as possible whether you and they have something to talk about. However, you probably *shouldn't* copy them verbatim, since many people in tech have seen this template so often that their eyes will glaze over if they encounter it again.

As you are writing these pitches, remember that people are not just economic animals. A sense of accomplishment, control over their own lives, and being part of a community motivates them just as much as money. People may volunteer to teach with you because it's what their friends are doing; similarly, a company may say that they're sponsoring classes for economically disadvantaged high school students because they want a larger pool of potential employees further down the road, but in reality, the CEO might actually be doing it simply because it's the right thing to do.

12.2   Branding and Positioning

A *brand* is someone's first reaction to a mention of a product; if their reaction is "what's that?", you don't have a brand yet. Branding is important because people aren't going to help with something they don't know about or don't care about.

Most discussion of branding today focuses on ways to build awareness online. Mailing lists, blogs, and Twitter all give you ways to reach people, but as the volume of (mis)information steadily increases, the attention paid to any particular interruption decreases. As this happens, *positioning* becomes more important. Positioning (sometimes also called "differentiation") is what sets your offering apart from others: it's the "unlike" section of your elevator pitches. When you are reaching out to people who are already generally familiar with your field, this is what you should emphasize, since it's what will catch their attention.

There are other things you can do to help build your brand as well. One is to use props: a robot car that one of your students made from

scraps she found around the house, the website another student made for his parents' retirement home, or anything else that makes what you're doing seem real. Another is to make a short video–no more than a few minutes long–showcasing the backgrounds and accomplishments of your students. The aim of both is to tell a story: while people always ask for data, stories are what they believe.

Foundational Myths

One of the most compelling stories a person or organization can tell is about why and how they got started. Are you teaching what you wish someone had taught you but didn't? Was there one particular person you wanted to help, and that opened the floodgates? Are you picking up where someone else left off, and if so, why?

Free samples are also compelling. Put some lesson materials online so that people can see what you teach; post a few (short) videos from actual workshops, or go to where your hoped-for learners or sponsors are and run a lunchtime drop-in session.

Whatever else you do, make your organization findable by doing what you can to make you and your organization rank highly in Google searches. There's a lot of folklore about how to do this under the label "SEO" (for "search engine optimization"); given Google's near-monopoly powers and lack of transparency, most of it boils down to trying to stay one step ahead of algorithms designed to prevent people from gaming rankings. Search for yourself and for your organization on a regular basis and see what comes up, then read these guidelines from Moz[1] and do what you can to improve your site. Keep this cartoon[2] in mind: people don't (initially) want to know about your org chart or get a virtual tour of your site; they want your address, parking information, and above all, some idea of what you teach, when you teach it, how to get in touch, and how it's going to change their life.

Offline findability is equally important for new organizations. Many of the people you hope to reach might not be online, or might not be online as often as you; notice boards in schools, local libraries, drop-in centers, and grocery stores are still an effective way to reach them.

Build Alliances

As discussed in the previous chapter (§11), building alliances with other groups that are doing things related to what you're doing

---

[1]https://moz.com/learn/seo/on-page-factors
[2]https://xkcd.com/773/

pays off in many ways. One of those is referrals: if someone approaches you for help, but would be better served by some other organization, take a moment to make an introduction. If you've done this several times, add something to your website to help the next person find what they need. The organizations you are helping will soon start to help you in return.

12.3   The Art of the Cold Call

Building a web site and hoping that people find it is one thing; calling people up or knocking on their door without any sort of prior introduction is another. As with standing up and teaching, though, it's a craft that can be learned like any other, and there are a few simple rules you can follow:

1. Start by establishing a point of connection: "I was speaking to X" or "You attended bootcamp Y". This must be specific: spammers and headhunters have trained us all to ignore anything that starts, "I recently read your website".

2. Explain how you are going to help make their lives better (e.g., "Your students will be able to do their math homework much faster if you let us help them").

3. Be specific about what you are offering (e.g., "Our usual two-day curriculum includes...") so that they can figure out right away whether this is worth pursuing, but keep it to one or two sentences.

4. Mention your backers, your size, how long you've been around, or your instructors's backgrounds to make yourself credible.

5. Create a slight sense of urgency ("we're booking workshops right now").

6. Tell them what your terms are: do you charge money, do they need to cover instructors' travel costs, can they reserve seats for their own staff, etc.

7. Above all, *keep it short.* The message below takes 30 seconds or less to scan; by the end, either they're interested enough to reply or they're not.

This template works pretty well, but "pretty well" is relative. Most organizations expect a 2-3% response rate to cold calls; for Software Carpentry, we found that about half of emails were answered, about

half of those answers were, "Sure, let's talk more," and about half of those led to workshops, which means that 10-15% of targeted emails to people we had some sort of connection with turned into workshops.

Mail Out of the Blue

Hi [name],

I hope you don't mind mail out of the blue, but I wanted to follow up on our conversation at the tech showcase last week to see if you would be interested having us run an instructor training workshop - we're scheduling the next batch over the next couple of weeks.

As background, hundreds of grassroots groups are now teaching people how to code, but ironically, most of their founders are teaching themselves how to teach. Just as a few lessons about variables, loops, and functions can save people hours of frustration, a few lessons about formative assessment, live coding, and other practices can help busy volunteers do more good in less time and with less pain.

This one-day class will introduce participants to a handful of key ideas that have a firm foundation in education research and proven useful in practice. The class is based on the Software Carpentry instructor training course, which has been delivered dozens of times on four continents. Instructions will be hands-on: short lessons will alternate with individual and group practical exercises, including practice teaching sessions.

We've run dozens of workshops like this since 2010, and if it sounds interesting, please give me a shout.

12.4   Making Change

Everyone is afraid of the unknown and of embarrassing themselves. As a result, most people would rather fail than change. Marketing is therefore not just about communicating clearly: it is also about figuring out why people are resisting your offer of help and then finding a way past that resistance.

For example, Lauren Herckis looked at why university faculty don't adopt better teaching methods[3]. She found that the main reason is a fear of looking stupid in front of their students, and

---

[3]https://www.insidehighered.com/news/2017/07/06/anthropologist-studies-why-professors-dont-adopt-innovative-teaching-methods

that secondary reasons were concern that the inevitable bumps in switching how they taught would affect course evaluations, and a desire to continue emulating the lecturers who had inspired them. It's pointless to argue about whether these issues are "real" in some objective sense: faculty believe they are, so any marketing aimed at faculty needs to address them.

Medical researchers realized several decades ago that there's no point coming up with a better way to do things if practitioners won't adopt it. The growing field of *implementation science* explores evidence-based ways to improve transference, and [Borrego2014] categories some related ideas for effecting change in higher education. The bulk of that paper expands upon this table (which is included as an image because rotating text in a simple cross-browser fashion is apparently still beyond present-day technology):

Figure 12.1: Change in Higher Education

Each of the major categories is defined by whether the change is individual or to the system as a whole, and whether it is pre-

scribed (top-down) or emergent (bottom-up). The person trying to make the changes–and make them stick–has a different role in each situation, and should pursue different strategies accordingly. In our experience, the most important things are:

1. *Ask, don't tell.* Teachers know their students and their needs much better than you do, so start by asking what they think the most pressing needs are.

2. *Find allies.* Many colleges and universities have teaching and learning centers whose staff are keen to improve teaching practices, and who also know how to navigate the local bureaucracy. Similarly, there are often tech meetup groups or other local organizations whose members are likely helpers.

3. *Start small.* [Lang2016] describes evidence-based teaching practices that can be put in place with minimal effort and at low cost. These may not have the most impact, but scoring a few early wins helps build support for larger and riskier efforts.

12.5   A Final Thought

As [Kuchner2011] says, if you can't be first in a category, create a new category that you can be first in; if you can't do that, think about doing something else entirely. This isn't as defeatist as it sounds: if someone else is already doing what you're doing better than you, there are probably lots of other equally useful things you could be doing instead.

12.6   Challenges

Write an Elevator Pitch for a City Councilor (15 minutes)

This chapter described an organization that offers weekend programming workshops for people re-entering the workforce after taking a break to raise children. Write an elevator pitch for that organization aimed at a city councilor whose support the organization needs.

Write Elevator Pitches for Your Organization (30 minutes)

Identify two groups of people your organization needs support from, and write an elevator pitch aimed at each one.

Identify Causes of Passive Resistance (30 minutes)

People who don't want change will sometimes say so out loud, but will also often use various forms of passive resistance, such as just not getting around to it over and over again, or raising one possible problem after another to make the change seem riskier and more expensive than it's actually likely to be. Working in small groups, list three or four reasons why people might not want your teaching initiative to go ahead, and explain what you can do with the time and resources you have to counteract each.

Why Teach Programming? (20 minutes)

As Mark Guzdial has pointed out[4], just learning to program doesn't automatically improve people's general problem-solving skills (i.e., there is no evidence for transference). However, there are lots of other reasons to learn to program:

1. To understand our world.
2. To study and understand processes.
3. To be able to ask questions about the influences on their lives.
4. To use an important new form of literacy.
5. To have a new way to learn science and mathematics.
6. As a job skill.
7. To use computers better.
8. As a medium in which to learn problem-solving.

Read the full explanations of these points[5], then draw a 3x3 grid whose X and Y axes are labelled "low", "medium", and "high" and place each point in one sector according to how important it is to you (the X axis) and to your learners (the Y axis).

1. Which points are closely aligned in importance (i.e., on the diagonal in your grid)?
2. Which points are mis-aligned (i.e., in the off-diagonal corners)?
3. How does this affect your marketing strategy?

---

[4]https://computinged.wordpress.com/2017/10/18/why-should-we-teach-programming-hint-its-not-to-learn-problem-solving/
[5]https://computinged.wordpress.com/2017/10/18/why-should-we-teach-programming-hint-its-not-to-learn-problem-solving/

13 License

This work is licensed under the Creative Commons Attribution 3.0 Unported license (CC-BY-3.0). You are free:

- to Share–to copy, distribute and transmit the work
- to Remix–to adapt the work

under the following conditions:

- Attribution–you must attribute the work in the manner specified by the author or licensor (but not in any way that suggests that they endorse you or your use of the work).

with the understanding that:

- Waiver–Any of the above conditions can be waived if you get permission from the copyright holder.

- Public Domain–Where the work or any of its elements is in the public domain under applicable law, that status is in no way affected by the license.

- Other Rights–In no way are any of the following rights affected by the license:

  - Your fair dealing or fair use rights, or other applicable copyright exceptions and limitations;
  - The author's moral rights;
  - Rights other persons may have either in the work itself or in how the work is used, such as publicity or privacy rights.

- Notice–For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to http://creativecommons.org/licenses/by/3.0/.

14 Code of Conduct

To ensure a welcoming environment for all, we require everyone participating in our classes to conform to the Code of Conduct given below. This code applies to all spaces managed by our group including, but not limited to, workshops, mailing lists, and online forums (including source code repositories). We strongly recommend that anyone running workshops or classes of any kind choose and publish a similar code so that everyone will know what is expected of them and what to do when those expectations are not met.

We are dedicated to providing a welcoming and supportive environment for all people, regardless of background or identity. However, we recognize that some groups in our community are subject to historical and ongoing discrimination, and may be vulnerable or disadvantaged. Membership in such a specific group can be on the basis of characteristics such as such as gender, sexual orientation, disability, physical appearance, body size, race, nationality, sex, color, ethnic or social origin, pregnancy, citizenship, familial status, veteran status, genetic information, religion or belief, political or any other opinion, membership of a national minority, property, birth, age, or choice of text editor. We do not tolerate harassment of participants on the basis of these categories, or for any other reason.

Harassment is any form of behavior intended to exclude, intimidate, or cause discomfort. Because we are a diverse community, we may have different ways of communicating and of understanding the intent behind actions. Therefore we have chosen to prohibit certain forms of behavior in our community, regardless of intent. Prohibited harassing behavior includes but is not limited to:

- written or verbal comments which have the effect of excluding people on the basis of membership of a specific group listed above;

- causing someone to fear for their safety, such as through stalking, following, or intimidation;

- the display of sexual or violent images;

- unwelcome sexual attention;

- non-consensual or unwelcome physical contact;

- sustained disruption of talks, events or communications;

- incitement to violence, suicide, or self-harm;

- continuing to initiate interaction (including photography or recording) with someone after being asked to stop; and

- publication of private communication without consent.

Behavior not explicitly mentioned above may still constitute harassment. The list above should not be taken as exhaustive but rather as a guide to make it easier to enrich all of us and the communities in which we participate. All interactions should be professional regardless of location: harassment is prohibited whether it occurs on or offline, and the same standards apply to both.

Enforcement of the Code of Conduct will be respectful and not include any harassing behaviors.

Thank you for helping make this a welcoming, friendly community for all.

This code of conduct is a modified version of that used by PyCon, which in turn is forked from a template written by the Ada Initiative and hosted on the Geek Feminism Wiki.

15 Citation

Please cite this work as:

Greg Wilson (ed.): *How to Teach Programming (And Other Things).* Second edition, Lulu.com, 2017, 978-1-365-98428-0, http://third-bit.com/teaching.

16 How to Contribute

We welcome contributions of all kinds, from errata to suggestions for improvements to new material. The source for this book is stored in a GitHub repository at http://github.com/gvwilson/thirdbit, so if you know how to use Git, and would like to add or fix something, please send us a pull request.

If you don't know how to use Git, you can file an issue in that repository or email the editor at `gvwilson@third-bit.com`. Please note that by doing so, you are agreeing that we may incorporate your changes in either original or edited form and release them under the same license as the rest of this material. Please also note that we require everyone involved in this project to abide by our Code of Conduct (§14).

Finally, we are always grateful to hear how you have used this material and how we could make it better: please email us if you have a story you would like to share.

17 Extra Material

This section includes material that doesn't naturally fit anywhere else.

17.1 How to Use This Material

This material has been taught as a multi-week online class, as a two-day in-person class, and as a two-day class in which the learners are in co-located groups and the instructor participates remotely.

Terminology

When we talk about workshops, we will try to be clear about whether we're discussing ones whose subject is programming, which are aimed at general learners, and those whose subject is how to teach, which are using this material.

In-Person

In our experience, this is the most effective way to deliver an instructor training workshop.

- Participants are physically together for one or two days. When they need to work in small groups (e.g., for practice teaching), some or all of them go to nearby breakout spaces. Participants bring their own tablets or laptops to view and edit online material during the class, and use pen and paper and/or whiteboards for some exercises.

- Participants use Etherpad or Google Doc for in-person training, both for shared note-taking (§9.5) and for posting exercise solutions and feedback on recorded lessons. Questions and discussion are done aloud.

- Several times during the training, participants are put in groups of three to teach for 2-3 minutes. The mechanics are described later (§3.2), and while participants are initially intimidated at first, they routinely rank it as the most useful part of the class.

Two-Day Online With Groups

In this format, learners are in groups of 4-12, but those groups are geographically distributed.

- Each class uses an Etherpad or Google Doc for shared note-taking (§9.5), and more importantly for asking and answering questions: having several dozen people try to talk on a call works poorly, so in most sessions, the instructor does the talking and learners respond through the note-taking tool's chat.

- Each group of learners is together in a room using one camera and microphone, rather than each being on the call separately. We have found that having good audio matters more than having good video, and that the better the audio, the more learners can communicate with the instructor and other rooms by voice rather than by using the Etherpad chat.

- We do the video lecture exercise as in the two-day in-person training.

Multi-Week Online

This was the first format we used, and we no longer recommend it.

- We met every week or every second week for an hour via web conferencing. Each meeting was held twice (or even three times) to accommodate learners' time zones and because video conferencing systems can't handle 60+ people at once.

- We used web conferencing and shared note-taking as described above for online group classes.

- Learners posted homework online between classes, and commented on each other's work. (In practice, comments were relatively rare: people seemed to prefer to discuss material in the web chats.)

- We used a WordPress blog for the first ten rounds of training, then a GitHub-backed blog, and finally Piazza. WordPress worked best: setting up accounts was tedious, but everything

after that ran smoothly. Using a GitHub blog worked so poorly that we didn't try it again: a third of the participants found it extremely frustrating, and post-publication commentary was awkward. Piazza was better than GitHub, but still not as easy for participants to pick up as WordPress. In particular, it was hard to find things once there were more than a dozen homework categories.

17.2   Key Terms

Educational psychology is the study of how people learn. It touches on everything from the neuropsychology of perception and the mechanisms of memory to the sociology of school systems and the philosophical question of what we actually mean by "learning" (which turns out to be pretty complicated once you start looking beyond the standardized Western classroom). Within the broad scope of educational psychology, two specific perspectives have primarily influenced our teaching practices (and by extension, this instructor training).

The first perspective is *cognitivism*, which treats learning as a problem in neuropsychology. Cognitivists focus their attention on things like pattern recognition, memory formation, and recall. It is good at answering low-level questions, but generally ignores larger issues like, "What do we mean by 'learning'?" and, "Who gets to decide?"

The second perspective is *situated learning*, which we discussed in Building Community (§11). For example, Software Carpentry aims to serve researchers who are exploring data management and programming on their own (legitimate peripheral practice) and make them aware of other people doing that work (simply by attending the workshop) and the best practices and ideas of that community of practice, thereby giving them a way to become members of that community. Situated learning thus describes why we teach, and recognizes that teaching and learning are always rooted in a social context. We then depend on the cognitivist perspective to drive *how* we teach the specific content associated with the community of practice.

Other Perspectives

There are many other perspectives outside cognitivist theory— see the Learning Theories site [Learning2017] for summaries. Besides cognitivism, those encountered most frequently include behaviorism *(which treats education as stimulus/response con-*

ditioning), constructivism *(which considers learning an active process during which learners construct knowledge for themselves), and* connectivism *(which emphasizes the social aspects of learning, particularly those made possible by the Internet). And yes, it would help if their names were less similar...*

Educational psychology does not tell us how to teach on its own because it under-constrains the problem: in real life, several different teaching methods might be consistent with what we currently know about how learning works. We therefore have to try those methods in the class, with actual learners, in order to find out how well they balance the different forces in play.

Doing this is called *instructional design.* If educational psychology is the science, instructional design is the engineering. For example, there are good reasons to believe that children will learn how to read best by starting with the sounds of letters and working up to words. However, there are equally good reasons to believe that children will learn best if they are taught to recognize entire simple words like "open" and "stop", so that they can start using their knowledge sooner.

The first approach is called "phonics", and the second, "whole language". The whole language approach may seem upside down, but more than a billion people have learned to read and write Chinese and similar ideogrammatic languages in exactly this way. The only way to tell which approach works best for most children, most of the time, is to try them both out. These studies have to be done carefully, because so many other variables can have an impact on rules. For example, the teacher's enthusiasm for the teaching method may matter more than the method itself, since children will model their teacher's excitement for a subject.

Unsurprisingly, effective teaching depends on what the teacher knows, which can be divided into:

- *content knowledge*, such as the "what" of programming;

- *general pedagogical knowledge*, i.e., an understanding of the psychology of learning; and

- the *pedagogical content knowledge* (PCK) that connects the two. PCK is things like what examples to use when teaching how parameters are passed to a function, or what misconceptions about wildcard expansion are most common. For example, an instructor could write variable names and values on paper plates and then stack and unstack them to show how the call stack works.

A great example of PCK is [Gelman2002], which is full of PCK for teaching introductory statistics. The CS Teaching Tips site [Tips2017] is gathering similar ideas for computing.

17.3   Situated Learning

A framework in which to think about grassroots education is *situated learning*, which focuses on how *legitimate peripheral participation* leads to people becoming members of a *community of practice*. Unpacking those terms, a community of practice is a group of people bound together by interest in some activity, such as knitting or particle physics. Legitimate peripheral participation means doing simple, low-risk tasks that community nevertheless recognizes as valid contributions: making your first scarf, stuffing envelopes during an election campaign, or proof-reading documentation for open source software.

Situated learning focuses on the transition from being a newcomer to being accepted as a peer by those who are already community members. This typically means starting with simplified tasks and tools, then doing similar tasks with more complex tools, and finally tackling the challenges of advanced practitioners. For example, children learning music may start by playing nursery rhymes on a recorder or ukulele, then play other simple songs on a trumpet or saxophone in a band, and finally start exploring their own musical tastes. Healthy communities understand and support these progressions, and recognize that each step is meant to give people a ramp rather than a cliff.

Whatever the domain, situated learning emphasizes that learning is a social activity. In order to be effective and sustainable, teaching therefore needs to be rooted in a community; if one doesn't exist, you need to build one.

17.4   Myths

One well-known scheme[1] characterizes learners as visual, auditory, or kinesthetic according to whether they like to see things, hear things, or do things. This scheme is easy to understand, but as de Bruyckere and colleagues point out in *Urban Myths About Learning and Education* [DeBruyckere2015], it is almost certainly false. Unfortunately, that hasn't stopped a large number of companies from marketing products based on it to parents and school boards.

---

[1]https://en.wikipedia.org/wiki/Learning_styles#Learning_modalities

This is not the only myth to plague education. The learning pyramid that shows we remember 10% of what we read, 20% of what we hear, and so on? Myth. The idea that "brain games" can improve our intelligence, or at least slow its decline in old age? Also a myth, as are the claims that the Internet is making us dumber or that young people read less than they used to.

Computing education has its own myths. Mark Guzdial's "Top 10 Myths About Teaching Computer Science" [Guzdial2015a] are:

1. The lack of women in Computer Science is just like all the other STEM fields.

2. To get more women in CS, we need more female CS faculty.

3. A good CS teacher is a good lecturer.

4. Clickers and the like are an add-on for a good teacher.

5. Student evaluations are the best way to evaluate teaching.

6. Good teachers personalize education for students' learning styles.

7. High schools just can't teach CS well, so they shouldn't do it at all.

8. The real problem is to get more CS curriculum into the hands of teachers.

9. All I need to do to be a good CS teacher is model good software development practice, because my job is to produce excellent software engineers.

10. Some people are just born to program.

The last of these is the most pervasive and most damaging. As discussed in Motivation (§7), Elizabeth Patitsas and others have shown that grades in computing classes are *not* bimodal [Patitsas2016], i.e., there isn't one group that gets it and another that doesn't. Many of the participants in our workshops have advanced degrees in intellectually demanding subjects, but have convinced themselves that they just don't have what it takes to be programmers. If all we do is dispel that belief, we will have done them a service.

17.5 Feedback on Bad Teaching Demo Videos

The two lists below summarize key feedback on the videos of bad live teaching and bad recorded teaching used in the chapters on performance (§3) and online teaching (§10).

Part 1: Bad Teaching in Person

- Starts by being rude to audience ("Could you sit down? Yeah, now please.")
- Text on the screen is far too small.
- Speaker frequently interrupts himself.
- Uses nonsensical names like "foo" instead of authentic tasks.
- First example returns the type of its argument: again, not a common task.
- "This is really simple stuff - even Excel users can understand it." Don't denigrate people's existing knowledge.
- "This is instantiating a new function object." Too much unnecessary jargon.
- "Yes, and it's lexical binding."
- "Which of course is polymorphic."
- "And as you'd expect from something which is doing lazy binding on types."
- "It works like you'd expect." No it doesn't–it took some of the brightest minds of the 20th Century several decades to come up with our current model of computing. To suggest otherwise is to imply that people who don't already understand it are stupid.
- "Trust me." Never a good thing to hear a teacher say...
- 01:29 check phone. The audience will never care more about the material than the presenter, and it's pretty clear at this point that the presenter doesn't care a lot.
- 01:40 defining higher-level functions.
- Any audience that needed basic function definition explained won't be ready for this.
- Any audience that's ready for this didn't need basic functions defined.
- Conclusion: the presenter has no idea who his audience actually is.
- 02:01 Corrected the mistake without explaining it.
- Failed to turn the mistake into a teachable moment.
- Will leave audience confused: what was wrong and why?

Part 2: Bad Teaching in Recorded Video

- Clipped at the start: normally start talking about a second in to give people a chance to collect themselves.
- Flubbed the first sentence: don't do a hundred takes, but don't just do one either.
- Screen is very cluttered: several irrelevant background windows open.

- Terminal is very cluttered: what does all that text have to do with the lesson?
- Font is much too small.
- Flipped between windows at 00:15 for no apparent reason.
- Speaking far too quickly.
- Starts by talking about some other platform that isn't on the screen, rather than what is.
- Don't draw attention at 00:30 to the Python version unless it's important.
- Background noise at 00:40.
- Explanation of what functions are seems to trying to motivate the lesson, but:
- uses jargon like "parameterize",
- nothing is happening on the screen while the presenter is talking, and
- reference to previous courses the learner might have done are obviously improvised.
- Audio is garbed starting at 01:16.
- Closing the background windows at 1:33 just draws attention to them at this point.
- Typing in code starting at 01:52 introduces a lot of background noise (keyboard clicking).
- "And it takes some sort of parameter" introduces important jargon (parameter), but then the presenter talks about the colon instead.
- Presenter goes from "body" to "scoping" starting at 2:07 without any clear explanation.
- Then dives down a rabbit hole briefly discussing scoping rules: this will be unintelligible to anyone new to functions.
- 02:28 "to get back to a top-level prompt" - what's a "top-level prompt"?
- At least the presenter is using a meaningful name for the function. . .
- 02:43 "this is all pretty simple stuff" - no! It's only simple once you know it.
- 02:50 "even if all you know is something like R": don't denigrate people's existing knowledge, or other communities.
- Pasting in a block of code around 03:05
- The code isn't explained. . .
- . . . and nobody is going to be able to keep up with the speed.
- 03:11 more background noise
- 03:12 Since there is a built-in function called 'sum', use some other name for the example function to avoid confusion.
- 03:18 "Because we've got to write something": the audience will

never care more about the material than the presenter, and it's pretty clear at this point that the presenter doesn't care a lot.
- 03:22 "Let me just rewrite this because we haven't introduced variable arguments."
- What bug is the presenter fixing?
- What are "variable arguments"? (Remember, this is the viewer's first introduction to functions.)
- The replacement code is identical to the original code except for one '*' character: beginners aren't going to notice that or know why it's important.
- 03:34 more keyboard clicking.
- 03:40 background voices.
- 03:53 "This is so simple, even Excel users could follow along." Again, never be derogatory in a lesson.
- 04:02 "This is actually polymorphic on types." Presenter clearly has no audience level in mind.
- Anyone who is new to functions isn't going to understand what "polymorphic on types" means.
- 04:09 Summing 'a', 'b', and 'c' fails.
- Don't include failing examples in videos unless the failure is purposeful.
- If a failure is included, explain what the fault was and how it was fixed (more than "I can't initialize..." and then self-interruption).
- 04:18 Abandoning example and going back to 'double' is going to confuse viewers.
- 04:40 Explanation of why doubling 'x' produces 'xx' is garbled and has nothing to do with functions.
- 04:48 "As you'd probably expect, you can nest the function calls": there are many things that are more important to explain before nested function calls.
- 04:58 "Yeah, I guess that's right." Does not inspire confidence in the presenter.
- 05:13 Audio is garbled again.
- 05:34 List comprehensions are another distraction.
- Last few seconds should have been edited out.

17.6 Feedback on Live Coding Demo Videos

The two lists below summarize key feedback on the two videos used in the discussion of live coding (§8).

Part 1: How Not to Do It

• Instructor ignores a red sticky clearly visible on a learner's laptop.

• Instructor is sitting, mostly looking at the laptop screen.

• Instructor is typing commands without saying them out loud.

• Instructor uses fancy shell prompt in the console window.

• Instructor uses small font in not full-screen console window with black background.

• The console window bottom is partially blocked by the learner's heads for those sitting in the back.

• Instructor receives a a pop-up notification in the middle of the session.

• Instructor makes a mistake (a typo) but simply fixes it without pointing it out, and redoes the command.

Part 2: How to Do It Right

• Instructor checks if the learner with the red sticky on her laptop still needs attention.

• Instructor is standing while instructing, making eye-contact with participants.

• Instructor is saying the commands out loud while typing them.

• Instructor moves to the screen to point out details of commands or results.

• Instructor simply uses $ as shell prompt in the console window.

• Instructor uses big font in wide-screen console window with white background.

• The console window bottom is above the learner's heads for those sitting in the back.

• Instructor makes mistake (a typo) and uses the occasion to illustrate how to interpret error-messages.

17.7   Evaluating Impact

A key part of effecting change (§12.4) is to convince people that what you're doing is having a positive impact. That turns out to be surprisingly hard for free-range programming workshops:

1. *Ask learners if the workshop was useful.* Study after study has shown that there is no correlation between how highly learners rate a course and how much they actually learn [Uttl2016], and most people working in education are now aware of that.

2. *Give them an exam at the end of the workshop.* Doing that dramatically changes the feel of the workshop, and how much they know at the end of the day is a poor predictor of how much they will remember two or three months later.

3. *Give them an exam two or three months later.* That's hard enough to do in a traditional battery-farmed learning environment; doing it with free-range learners is even harder. In addition:

   - The people who didn't get anything out of the workshop are probably less likely to take part in follow-up, so feedback gathered this way will be subject to self-selection bias.

   - The fact that learners *remember* something doesn't necessarily mean it was useful (although they are more likely to remember things that are useful than things that aren't).

4. *See if they keep using what they learned.* This is a good way to evaluate employment-oriented skills, but equally useful for things people have learned for fun. The problem is how to do it: you probably shouldn't put spyware on their computers, and follow-up surveys suffer from the same low return rate and self-selection bias as exams.

5. *See if they recommend the workshop to friends.* This method often strikes the best balance between informative and doable: if people are recommending your workshop to other people, that's a pretty good sign.

There are many other options; the most important thing is to figure out early on how you're going to know whether you're teaching the right things the right way, and how you're going to convince potential backers that you're doing so.

17.8   Three Kinds of Thinking

[Fink2013] suggests designing exercises to prompt three kinds of thinking, and provides these examples (among others):

Field	Critical Thinking	Creative Thinking	Practical Thinking
Biology	Evaluate the validity of the bacterial theory of ulcers.	Design a experiment to test the bacterial theory of ulcers.	How would the bacterial theory of ulcers change conventional treatment regimes?
Art	Compare and contrast how Rembrandt and Van Gogh used light in these two paintings.	Draw a beam of light.	How could we reproduce the lighting in this painting in an actual room?

To ensure that key concepts are truly understood, instructors should give learners exercises of all three types for each concept.

18 Bibliography

18.1 What to Read Next

- [Ambrose2010] Susan A. Ambrose, Michael W. Bridges, Michele DiPietro, Marsha C. Lovett, and Marie K. Norman: *How Learning Works: Seven Research-Based Principles for Smart Teaching*[1]. Jossey-Bass, 2010, 978-0470484104. *An excellent overview of what we know about education and why we believe it's true, covering everything from cognitive psychology to social factors.*

- [Brookfield2016] Stephen D. Brookfield and Stephen Preskill: *The Discussion Book: 50 Great Ways to Get People Talking*[2]. Jossey-Bass, 2016, 978-1119049715. *Describes fifty different ways to get groups talking productively.*

- [Brown2007] Michael Jacoby Brown: *Building Powerful Community Organizations*[3]. Long Haul Press, 2007, 978-0977151806. *An excellent practical introduction to creating effective organizations in and for communities written by someone with decades of experience doing exactly that.*

- [Didau2016] David Didau and Nick Rose: *What Every Teacher Needs to Know About Psychology*[4]. John Catt Educational, 2016, 978-1909717855. *An informative, opinionated survey of what modern psychology has to say about teaching.*

---

[1] https://www.amazon.com/How-Learning-Works-Research-Based-Principles/dp/0470484101/

[2] https://www.amazon.com/Discussion-Book-Great-People-Talking/dp/1119049717/

[3] https://www.amazon.com/Building-Powerful-Community-Organizations-Personal/dp/0977151808/

[4] https://www.amazon.com/Every-Teacher-Needs-About-Psychology/dp/1909717851/

- [Green2014] E. Green: *Building a Better Teacher: How Teaching Works (and How to Teach It to Everyone)*[5]. W. W. Norton, 2014, 978-0393244151. *A well-written look at why educational reforms in the past 50 years have mostly missed the mark, and what we should be doing instead.*

- [Guzdial2015b] Mark Guzdial: _Learner-Centered Design of Computing Education: Research on Computing for Everyone[6] Morgan & Claypool, 2015, 978-1627053518. *An evidence-based argument that we must design computing education for everyone, not just people who think they are going to become professional programmers.*

- [Huston2009] Therese Huston: *Teaching What You Don't Know*[7]. Harvard University Press, 2009, 978-0674035805. *A pointed, funny, and very useful book that explores exactly what the title suggests.*

- [Lang2016] James M. Lang: *Small Teaching: Everyday Lessons from the Science of Learning*[8]. Jossey-Bass, 2016, 978-1118944493. *Presents a selection of accessible evidence-based practices that teachers can adopt when they little time and few resources.*

- [Lemov2014] Doug Lemov: *Teach Like a Champion 2.0: 62 Techniques that Put Students on the Path to College*[9] (2nd edition). Jossey-Bass, 2014, 978-1118901854. *Presents 62 classroom techniques drawn from intensive study of thousands of hours of video of good teachers in action.*

- [Margolis2003] J. Margolis and A. Fisher: *Unlocking the Clubhouse: Women in Computing*[10]. MIT Press, 2003, 978-0262632690. *A groundbreaking report on the gender imbalance in computing, and the steps Carnegie-Mellon took to address the problem.*

---

[5] https://www.amazon.com/Building-Better-Teacher-Teaching-Everyone/dp/0393351084/

[6] https://www.amazon.com/Learner-Centered-Design-Computing-Education-Human-Centered/dp/1627053514/

[7] https://www.amazon.com/Teaching-What-You-Dont-Know/dp/0674035801/

[8] https://www.amazon.com/Small-Teaching-Everyday-Lessons-Learning/dp/1118944496/

[9] https://www.amazon.com/Teach-Like-Champion-2-0-Techniques/dp/1118901851/

[10] https://www.amazon.com/Unlocking-Clubhouse-Women-Computing-Press/dp/0262632691/

18.2 Other References

- [Abela2009] Andrew Abela: "Chart Suggestions–A Thought Starter". http://extremepresentation.typepad.com/files/choosing-a-good-chart-09.pdf, viewed April 2017.

- [Ada2017] Ada Initiative: "Imposter Syndrome training". https://adainitiative.org/continue-our-work/impostor-syndrome-training/, viewed May 2017.

- [Adams1975] Frank Adams: *Unearthing Seeds of Fire: The Idea of Highlander*[11]. John F. Blair, 1975, 978-0895870193. *A history of the Highlander Folk School and its founder, Myles Horton, who inspired many other social change organizations.*

- [Aiken1975] Edwin G. Aiken, Gary S. Thomas, and William A. Shennum: "Memory for a Lecture: Effects of Notes, Lecture Rate, and Informational Density." *Journal of Educational Psychology*, 67(3), June 1975, 10.1037/h0076613. *A landmark study showing that taking notes improves retention when learning.*

- [Alinsky1989] Saul Alinsky: *Rules for Radicals: A Practical Primer for Realistic Radicals*[12]. Vintage, 1989, 978-0679721130. *A widely-read guide to community organization written by one of the 20th Century's great organizers.*

- [Avanti2013]: Avanti Learning Centre: "ConcepTests at Avanti's Learning Centre in Kanpur." https://www.youtube.com/watch?v=2LbuoxAy56o, viewed April 2017.

- [Aveling2013] Emma-Louise Aveling, Peter McCulloch, and Mary Dixon-Woods: "A Qualitative Study Comparing Experiences of the Surgical Safety Checklist in Hospitals in High-Income and Low-Income Countries." *BMJ Open*, 3(8), 2013, 10.1136/bmjopen-2013-003039. *Reports on surgical checklist implementations and effects in the UK and Africa.*

- [Barker2015] Lecia Barker, Christopher Lynnly Hovey, and Jane Gruning: "What Influences CS Faculty to Adopt Teaching Practices?", *Proc. 46th ACM Technical Symposium on Computer Science Education*, 2015, 10.1145/2676723.2677282. *Describes findings from a two-part study of how computer science educators adopt new teaching practices.*

---

[11]https://www.amazon.com/Unearthing-Seeds-Fire-Idea-Highlander/dp/0895870193/

[12]https://www.amazon.com/Rules-Radicals-Practical-Primer-Realistic/dp/0679721134/

- [Benner2000] Patricia Benner: *From Novice to Expert: Excellence and Power in Clinical Nursing Practice*[13]. Pearson, 2000, 978-0130325228. *A classic study of clinical judgment and how expertise develops.*

- [Biggs2011] John Biggs and Catherine Tang: *Teaching for Quality Learning at University*[14]. Open University Press, 2011, 978-0335242757. *A step-by-step guide to lesson development, delivery, and evaluation for people working in higher education.*

- [Bohay2011] Mark Bohay, Daniel P. Blakely, Andrea K. Tamplin, and Gabriel A. Radvansky: "Note Taking, Review, Memory, and Comprehension." *American Journal of Psychology*, 124(1), 2011, 10.5406/amerjpsyc.124.1.0063. *Presents a study showing that note-taking improves retention most at deeper levels of understanding.*

- [Bollier2014] David Bollier: *Think Like a Commoner: A Short Introduction to the Life of the Commons*[15]. New Society Publishers, 2014, 978-0865717688. *A short introduction to one of the most widely used kinds of governance in human societies throughout history.*

- [Borrego2014] Maura Borrego and Charles Henderson: "Increasing the Use of Evidence-Based Teaching in STEM Higher Education: A Comparison of Eight Change Strategies." *Journal of Engineering Education*, 103(2), April 2014, 10.1002/jee.20040. *Categories different approaches to effecting change in higher education.*

- [Brown2014] Neil C. C. Brown and Amjad Altadmri: "Investigating Novice Programming Mistakes: Educator Beliefs vs. Student Data." *Proc Tenth Annual Conference on International Computing Education Research*, 2014, 10.1145/2632320.2632343. *Uses data from over 100,000 students to show that educators know less than they think about what mistakes novice programmers actually make.*

---

[13]https://www.amazon.com/Novice-Expert-Excellence-Clinical-Commemorative/dp/0130325228/

[14]https://www.amazon.com/Teaching-Learning-University-Research-Education/dp/0335242758/

[15]https://www.amazon.com/Think-Like-Commoner-Introduction-Commons/dp/0865717680/

- [Cottom] Tressie McMillan Cottom: *Lower Ed: The Troubling Rise of For-Profit Colleges in the New Economy*[16]. The New Press, 2017, 978-1620970607. *Lays bare the dynamics of this growing "educational" industry to show how it leads to greater inequality rather than less.*

- [Cottrill2016] Cameron Cottrill: "Why Talented Black and Hispanic Students Can Go Undiscovered." https://mobile.nytimes.com/2016/04/10/upshot/why-talented-black-and-hispanic-students-can-go-undiscovered.html, viewed April 2017.

- [Deathbulge Feedback Feelings] Deathbulge: "Feedback Feelings". http://www.deathbulge.com/comics/155, viewed May 2017. *How many of us react to feedback.*

- [DeBruyckere2015] Pedro De Bruyckere, Paul A. Kirschner, and Casper D. Hulshof: *Urban Myths about Learning and Education*[17]. Academic Press, 2015, 978-0128015377. *Describes and debunks some widely-held myths about how people learn.*

- [Epstein2002] L.C. Epstein: *Thinking Physics is Gedanken Physics*[18]. Insight Press, 2002: 978-0935218084. *An entertaining problem-based introduction to thinking like a physicist.*

- [Ericson2017] Barbara J. Ericson, Lauren E. Margulieux, and Jochen Rick: "Solving Parsons Problems Versus Fixing and Writing Code." Proc. Koli Calling 2017, November 2017, 10.1145/3141880.3141895. *Another in a series of studies showing that Parsons Problems are an effective way to teach programming.*

- [Fehily2008] Chris Fehily: *SQL: Visual QuickStart Guide*[19] (3rd edition). Peachpit Press, 2008 978-0321553577. *An introduction to SQL that is both a good tutorial and a good reference guide.*

- [Fincher2012] Sally Fincher, Brad Richards, Janet Finlay, Helen Sharp, and Isobel Falconer: "Stories of Change: How Educators

---

[16]https://www.amazon.com/Lower-Ed-Troubling-Profit-Colleges/dp/1620970600/

[17]https://www.amazon.com/Urban-Myths-about-Learning-Education/dp/0128015373/

[18]https://www.amazon.com/Thinking-Physics-Understandable-Practical-Reality/dp/0935218084/

[19]https://www.amazon.com/SQL-Visual-QuickStart-Guide-3rd/dp/0321553578/

Change Their Practice." *Proc. Frontiers in Education Conference*, 2012, 10.1109/FIE.2012.6462317. *A detailed look at how educators actually adopt new teaching practices.*

- [Fincher2007] Sally Fincher and Josh Tenenberg: "Warren's Question." *Proc. Third International Workshop on Computing Education Research*, 2007, 10.1145/1288580.1288588. *A detailed look at a particular instance of transferring a teaching practice.*

- [Fink2003] L. Dee Fink: "A Self-Directed Guide to Designing Courses for Significant Learning." https://www.deefinkandassociates.com/GuidetoCourseDesignAug05.pdf, viewed April 2017.

- [Fink2013] L. Dee Fink: *Creating Significant Learning Experiences: An Integrated Approach to Designing College Courses*[20] (2nd edition). Jossey-Bass, 2013, 978-1118124253. *A step-by-step guide to a systematic lesson design process.*

- [Fogel2017] Karl Fogel: *Producing Open Source Software*[21], viewed September 2017. *The second edition of the definitive description of how to set up and run an open software development project.*

- [Gawande2011] Atul Gawande: "Personal Best." *The New Yorker*, October 3, 2011. *Describes how having a coach can improve practice in a wide variety of fields.*

- [Gelman2002] Andrew Gelman and Deborah Nolan: *Teaching Statistics: A Bag of Tricks*[22]. Oxford University Press, 2002, 978-0198572244. *A collection of useful motivating examples for teaching statistics.*

- [Gormally2014] Cara Gormally, Mara Evans, and Peggy Brickman: "Feedback about Teaching in Higher Ed: Neglected Opportunities to Promote Change." *CBE Life Sciences Education*, 13(2), 2014, 10.1187/cbe.13-12-0235. *Summarizes the best practices for providing instructional feedback, and recommends specific strategies for providing feedback.*

- [Guzdial2017] Mark Guzdial: "Computing Education Blog". https://computinged.wordpress.com/, viewed April 2017. *An informative, frequently-updated blog about computing education.*

---

[20]https://www.amazon.com/Creating-Significant-Learning-Experiences-Integrated/dp/1118124251/

[21]http://producingoss.com/

[22]https://www.amazon.com/Teaching-Statistics-Tricks-Andrew-Gelman/dp/0198572247/

- [Guzdial2013] Mark Guzdial: "Exploring Hypotheses About Media Computation." *Proc. Ninth Annual International ACM Conference on International Computing Education Research*, 2013, 10.1145/2493394.2493397. *A look back on 10 years of media computation research.*

- [Guzdial2016]: "Five Principles for Programming Languages for Learners". *Communications of the ACM*, 2016, https://cacm.acm.org/blogs/blog-cacm/203554-five-principles-for-programming-languages-for-learners/fulltext. *Five rules for choosing a programming environment for novices.*

- [Guzdial2015a] Mark Guzdial: "Top 10 Myths About Teaching Computer Science". *Communications of the ACM*, 2015, https://cacm.acm.org/blogs/blog-cacm/189498-top-10-myths-about-teaching-computer-science/fulltext. *Ten things many people believe that aren't true.*

- [Hannay2009] Jo E. Hannay, Tore Dybå, Erik Arisholm, and Dag I.K. Sjøberg: "The Effectiveness of Pair Programming: A Meta-Analysis." *Information and Software Technology*, 51(7), 2009, 10.1016/j.infsof.2009.02.001. *A summary of research on the effectiveness of pair programming.*

- [Harms2016] Kyle James Harms, Jason Chen, and Caitlin L. Kelleher: "Distractors in Parsons Problems Decrease Learning Efficiency for Young Novice Programmers." *Proc. ICER 2016*, September 2016, 10.1145/2960310.2960314. *Shows that adding distractors to Parsons Problems does not improve learning while increasing the time spent solving them.*

- [Henderson2011] Charles Henderson, Andrea Beach, and Noah Finkelstein: "Facilitating Change in Undergraduate STEM Instructional Practices: An Analytic Review of the Literature." *Journal of Research in Science Teaching*, 48(8), 2011, 10.1002/tea.20439. *Describes eight approaches to effecting change in STEM education that form a useful framework for thinking about how free-range workshops can go mainstream.*

- [Henry2014] Liz Henry: "Unlocking the Invisible Elevator: Accessibility at Tech Conferences." https://modelviewculture.com/pieces/unlocking-the-invisible-elevator-accessibility-at-tech-conferences, viewed April 2017.

- [ISW2017] Instructional Skills Workshop Network: "ISW and FDSW Handbooks", https://iswnetwork.ca/resources/pd-resources/isw-and-fdw-handbooks/, accessed May 2017.

- [Kernighan1982] Brian W. Kernighan and P.J. Plauger: *The Elements of Programming Style*[23] (2nd edition). McGraw-Hill, 1982, 978-0070342071. *An early and influential description of the Unix programming philosophy.*

- [Kernighan1984] Brian W. Kernighan and Rob Pike: *The UNIX Programming Environment*[24]. Prentice Hall, 1984, 978-0139376818. *An influential early description of Unix.*

- [Kernighan1988] Brian W. Kernighan and Dennis M. Ritchie: *The C Programming Language*[25] (2nd edition). Prentice Hall, 1988, 978-0131103709. *The book that made C a popular programming language.*

- [Kirschner2006] Paul Kirschner, John Sweller, and Richard Clark: "Why Minimal Guidance During Instruction Does Not Work: An Analysis of the Failure of Constructivist, Discovery, Problem-Based, Experiential, and Inquiry-Based Teaching." *Educational Psychologist*, 41(2), 2006. *Argues that inquiry-based learning is less effective for novices than guided instruction.*

- [Koedinger2015] Kenneth R. Koedinger, Jihee Kim, Julianna Zhuxin Jia, Elizabeth A. McLaughlin, and Norman L. Bier: "Learning is Not a Spectator Sport: Doing is Better Than Watching for Learning from a MOOC" *Proc. Second ACM Conference on Learning @ Scale*, 2015, 10.1145/2724660.2724681. *Measures the benefits of doing rather than watching.*

- [Kraut2012] Robert E. Kraut and Paul Resnick: *Building Successful Online Communities: Evidence-Based Social Design*[26]. MIT Press, 2012, 978-0262016575. *Sums up what we actually know about making thriving online communities and why we believe it's true.*

- [Kuchner2011] Marc Kuchner: *Marketing for Scientists: How to Shine in Tough Times*[27]. Island Press, 2011, 978-1597269940. *A short, readable guide to making people aware of, and care about, your work.*

---

[23] https://www.amazon.com/Elements-Programming-Style-2nd/dp/0070342075/

[24] https://www.amazon.com/Unix-Programming-Environment-Prentice-Hall-Software/dp/013937681X/

[25] https://www.amazon.com/Programming-Language-Brian-W-Kernighan/dp/0131103628/

[26] https://www.amazon.com/Building-Successful-Online-Communities-Evidence-Based/dp/0262016575/

[27] https://www.amazon.com/Marketing-Scientists-Shine-Tough-Times/dp/1597269948/

- [Kuittinen2004] Marja Kuittinen and Jorma Sajaniemi: "Teaching Roles of Variables in Elementary Programming Courses" *Proc. 9th Annual SIGCSE Conference on Innovation and Technology in Computer Science Education*, 2004, 10.1145/1007996.1008014. *Presents a few patterns used in novice programming and looks at the pedagogical value of teaching them.*

- [Kulkarni2013] Chinmay Kulkarni, Koh Pang Wei, Huy Le, Daniel Chia, Kathryn Papadopoulos, Justin Cheng, Daphne Koller, and Scott R. Klemmer: "Peer and Self Assessment in Massive Online Classes". *ACM Transactions on Computer-Human Interaction*, 20(6), 2013, 10.1145/2505057. *Replicates the finding of [Paré2008] that peer grading can be as effective at scale as expert grading.*

- [Lang2013] James M. Lang: *Cheating Lessons: Learning From Academic Dishonesty*[28]. Harvard University Press, 2013, 978-0674724631. *Explores why students cheat, and how courses often give them incentives to do so.*

- [Learning2017] "Learning Theories" website. http://www.learning-theories.com/, viewed April 2017.

- [Lee2017] Cynthia Lee: "What Can I Do Today to Create a More Inclusive Community in CS?" http://bit.ly/2oynmSH, viewed April 2017. *A practical checklist of things instructors can do to make their computing classes more inclusive.*

- [Littky2004] D. Littky and S. Grabelle: *The Big Picture: Education is Everyone's Business*[29]. Association for Supervision and Curriculum Development, 2004, 978-0871209719. *A personal exploration of the purpose of education and how to make schools better.*

- [Macnamara2014] Brooke N. Macnamara, David Z. Hambrick, and Frederick L. Oswald: "Deliberate Practice and Performance in Music, Games, Sports, Education, and Professions." *Psychological Science*, 25(8), 2014, 10.1177/0956797614535810. *A meta-study of the effectiveness of deliberate practice.*

- [Margolis2010] J. Margolis, R. Estrella, J. Goode, J.J. Holme, and K. Nao: *Stuck in the Shallow End: Education, Race, and*

---

[28]https://www.amazon.com/Cheating-Lessons-Learning-Academic-Dishonesty/dp/0674724631/
[29]https://www.amazon.com/Big-Picture-Education-Everyones-Business/dp/0871209713/

Computing[30]. MIT Press, 2010, 978-0262260961. *A hard-hitting look at racial inequities in computing education.*

- [Marsh2002] Herbert W. Marsh and John Hattie: "The Relation Between Research Productivity and Teaching Effectiveness." *Journal of Higher Education*, 73(5), 2002. *One study of many showing there is zero correlation between research ability and teaching effectiveness.*

- [Mayer2003] Richard E. Mayer and Roxana Moreno: "Nine Ways to Reduce Cognitive Load in Multimedia Learning." *Educational Psychologist*, 38, 2003. *Shows how research into how we absorb and process information can be applied to the design of instructional materials.*

- [Midwest2010]: Midwest Academy: *Organizing for Social Change*[31] (4th edition). Forum Press, 2010, 978-0984275212. *A practical step-by-step handbook for those who want to build effective social change organizations.*

- [Miller2013] Kelly Miller, Nathaniel Lasry, Kelvin Chu, and Eric Mazur: "Role of Physics Lecture Demonstrations in Conceptual Learning." *Physical Review Physics Education Research*, 9(2), 2013, 10.1103/PhysRevSTPER.9.020113. *Reports a detailed study of what students learn during demonstrations and why.*

- [Miller2016] Michelle D. Miller: *Minds Online: Teaching Effectively with Technology*[32]. Harvard University Press, 2016, 978-0674660021. *Desribes ways that insights from neuroscience can be used to improve online teaching.*

- [Mueller2014] Pam A. Mueller and Daniel M. Oppenheimer: "The Pen Is Mightier Than the Keyboard." *Psychological Science*, 25(6), 2014, 10.1177/0956797614524581. *Presents evidence that taking notes by hand is more effective than taking notes on a laptop.*

- [Muller2011] Derek Muller: "Khan Academy and the Effectiveness of Science Videos". https://fnoschese.wordpress.com/2011/03/17/khan-academy-and-the-effectiveness-of-

---

[30]https://www.amazon.com/Stuck-Shallow-End-Education-Computing/dp/0262514044/

[31]https://www.amazon.com/Organizing-Social-Change-Bobo-Kendall/dp/0984275215/

[32]https://www.amazon.com/Minds-Online-Teaching-Effectively-Technology/dp/0674660021/

science-videos/, viewed April 2017. *A hard look at how and whether educational video works.*

- [Nederbragt2016a] Lex Nederbragt: "A Video Introduction to Live Coding Part 1". https://youtu.be/bXxBeNkKmJE, viewed April 2017. *How not to use live coding when teaching.*

- [Nederbragt2016b] Lex Nederbragt: "A Video Introduction to Live Coding Part 2". https://youtu.be/SkPmwe_WjeY, viewed April 2017. *How to use live coding effectively as a teaching technique.*

- [Nilson2017] Linda B. Nilson and Ludwika A. Goodson: *Online Teaching at Its Best: Merging Instructional Design with Teaching and Learning Research*[33]. Jossey-Bass, 2017, 978-1119242291. *A sequel to the first author's widely-regarded guide for college-level instructors that focuses on online teaching.*

- [Orndorff2015] Harold N. Orndorff III: "Collaborative Note-Taking: The Impact of Cloud Computing on Classroom Performance." *International Journal of Teaching and Learning in Higher Education*, 27(3), 2015. *Presents a study showing that collaborative note-taking improves grades and learning outcomes.*

- [Paré2008] D.E. Paré and S. Joordens: "Peering Into Large Lectures: Examining Peer and Expert Mark Agreement Using peerScholar, an Online Peer Assessment Tool." *Journal of Computer Assisted Learning*, August 2008, 10.1111/j.1365-2729.2008.00290.x. *Shows that peer grading by small groups can be as effective as expert grading once accountability features are introduced.*

- [Patitsas2016] Elizabeth Patitsas, Jesse Berlin, Michelle Craig, and Steve Easterbrook: "Evidence That Computer Science Grades Are Not Bimodal." *Proc. 2016 ACM Conference on International Computing Education Research*, 2016, 10.1145/2960310.2960312. *Presents a statistical analysis and an experiment which jointly show that grades in computing classes are not bimodal–i.e., there is no geek gene.*

- [Petre2016] Marian Petre, André van der Hoek, and Yen Quach: *Software Design Decoded: 66 Ways Experts Think*[34]. MIT Press, 2016, 978-0262035187. *A short illustrated overview of how expert software developers think.*

---

[33]https://www.amazon.com/Online-Teaching-Its-Best-Instructional/dp/1119242290/
[34]https://www.amazon.com/Software-Design-Decoded-Experts-Think/dp/0262035189/

- [Pigni2016] Alessandra Pigni: *The Idealist's Survival Kit: 75 Simple Ways to Avoid Burnout*[35]. Parallax Press, 2016, 978-1941529348. *A guide to staying sane and healthy while doing good.*

- [Porter2013] Leo Porter, Mark Guzdial, Charlie McDowell, and Beth Simon: "Success in Introductory Programming: What Works?" *Communications of the ACM*, 56(8), August 2013, 10.1145/2492007.2492020. *Summarizes the evidence that peer instruction, media computation, and pair programming can significantly improve outcomes in introductory programming courses.*

- [Ray2014] Eric J. Ray and Deborah S. Ray: *Unix and Linux: Visual QuickStart Guide*[36] (5th edition). Peachpit Press, 2014, 978-0321997548. *An introduction to Unix that is both a good tutorial and a good reference guide.*

- [Schön1984] Donald A. Schön: *The Reflective Practitioner: How Professionals Think In Action*[37]. Basic Books, 1984, 978-0465068784. *A groundbreaking look at how professionals in different fields actually solve problems.*

- [Scott1999] J. C. Scott: *Seeing Like a State: How Certain Schemes to Improve the Human Condition Have Failed*[38]. Yale University Press, 1999, 978-0300128789. *Argues that large organizations consistently prefer uniformity over productivity.*

- [Spalding2014] Dan Spalding: *How to Teach Adults*[39]. Jossey-Bass, 2014, 978-1118841365. *A short guide to teaching adult free-range learners informed by the author's social activism.*

- [Spannaus2012] Tim Spannaus: *Creating Video for Teachers and Trainers: Producing Professional Video with Amateur Equipment*[40]. Pfeiffer, 2012, 978-1118088098. *A short, practical guide to doing exactly what the title says.*

---

[35] https://www.amazon.com/Idealists-Survival-Kit-Prevent-Burnout/dp/1941529348/

[36] https://www.amazon.com/Unix-Linux-Visual-QuickStart-Guide/dp/0321997549/

[37] https://www.amazon.com/Reflective-Practitioner-Professionals-Think-Action/dp/0465068782/

[38] https://www.amazon.com/Seeing-like-State-Certain-Condition/dp/0300078153/

[39] https://www.amazon.com/How-Teach-Adults-Jossey-Bass-Education/dp/1118841360/

[40] https://www.amazon.com/Creating-Video-Teachers-Trainers-Professional/dp/1118088093/

- [Steele2011] C. M. Steele: *Whistling Vivaldi: And Other Clues to How Stereotypes Affect Us*[41]. W. W. Norton, 2011, 978-0393341485. *Explains and explores stereotype threat and strategies for addressing it.*

- [Tips2017] CS Teaching Tips website, http://csteachingtips.org/, viewed April 2017.

- [Taylor2014] Chad Taylor: "Q&A: Making Tech Events Accessible to the Deaf Community." https://modelviewculture.com/pieces/qa-making-tech-events-accessible-to-the-deaf-community, viewed April 2017.

- [Ubell2017] Robert Ubell: "How the Pioneers of the MOOC Got It Wrong." http://spectrum.ieee.org/tech-talk/at-work/education/how-the-pioneers-of-the-mooc-got-it-wrong, January 16, 2017. *A brief exploration of why MOOCs haven't lived up to initial hype.*

- [Urbach2014] David R. Urbach, Anand Govindarajan, Refik Saskin, Andrew S. Wilton, and Nancy N. Baxter: "Introduction of Surgical Safety Checklists in Ontario, Canada." *New England Journal of Medicine*, 370(11), 2014, 10.1056/NEJMsa1308261. *Reports a study showing that the introduction of surgical checklists did not have a significant effect on operative outcomes.*

- [Uttl2016] Bob Uttl, Carmela A. White, and Daniela Wong Gonzalez: "Meta-Analysis of Faculty's Teaching Effectiveness: Student Evaluation of Teaching Ratings and Student Learning Are Not Related" *Studies in Educational Evaluation*, 2016, 10.1016/j.stueduc.2016.08.007 *A summary of studies shown that how students rate a course and how much they actually learn are not related.*

- [W3C2017] W3C Web Accessibility Initiative: "How to Make Presentations Accessible to All." https://www.w3.org/WAI/training/accessible, accessed May 2017.

- [Watters2014] Audrey Watters: *The Monsters of Education Technology*[42]. CreateSpace, 2014, 978-1505225051. *A collection of essays about the history of educational technology and the exaggerated claims repeatedly made for it. There's more criticism than prescription, but the former is well-informed and sharp-edged.*

---

[41]https://www.amazon.com/Whistling-Vivaldi-Stereotypes-Affect-Issues/dp/0393339726/

[42]https://www.amazon.com/Monsters-Education-Technology-Audrey-Watters/dp/1505225051/

- [Wiggins2005] G.P. Wiggins and J. McTighe: *Understanding by Design*[43]. Association for Supervision and Curriculum Development, 2005, 978-1416600350. *A lengthy presentation of reverse instructional design.*

- [Wilkinson2011] Richard Wilkinson and Kate Pickett: *The Spirit Level: Why Greater Equality Makes Societies Stronger*[44]. Bloomsbury Press, 2011, 978-1608193417. *Presents evidence that inequality harms everyone, both economically and otherwise.*

- [Willingham2010] Daniel T. Willingham: *Why Don't Students Like School?*[45]. Jossey-Bass, 2010, 978-0470591963. *A cognitive scientist looks at how the mind works and what it means in the classroom.*

- [Wilson2016] Greg Wilson: "How to Teach Badly". https://www.youtube.com/watch?v=-ApVt04rB4U, viewed May 2017.

- [Wilson2017] Greg Wilson: "How to Teach Badly (Part 2)". https://youtu.be/xcnoHaxXvdQ, viewed November 2017.

- [Wlodkowski2017] Raymond J. Wlodkowski and Margery B. Ginsberg: *Enhancing Adult Motivation to Learn: A Comprehensive Guide to Teaching All Adults*[46] (4th edition). Jossey-Bass, 2017, 978-1119077992. *The standard reference for understanding adult motivation, which more than lives up to the "comprehensive" in its name.*

---

[43] https://www.amazon.com/Understanding-Design-Grant-Wiggins/dp/1416600353/

[44] https://www.amazon.com/Spirit-Level-Equality-Societies-Stronger/dp/1608193411/

[45] https://www.amazon.com/Why-Dont-Students-Like-School/dp/047059196X/

[46] https://www.amazon.com/Enhancing-Adult-Motivation-Learn-Comprehensive/dp/1119077990/

19 Glossary

- Authentic Task: A task which contains important elements of things that learners would do in real (non-classroom situations). To be authentic, a task should require learners to construct their own answers rather than choose between provided answers, and to work with the same tools and data they would use in real life.

- Behaviorism: A theory of learning whose central principle is stimulus and response, and whose goal is to explain behavior without recourse to internal mental states or other unobservables. See also cognitivism.

- Bloom's Taxonomy: A six-part hierarchical classification of understand whose levels are *knowledge*, *comprehension*, *application*, *analysis*, *synthesis*, and *evaluation* that has been widely adopted[1]. See also Fink's Taxonomy.

- Chunking: The act of grouping related concepts together so that they can be stored and processed as a single unit.

- Cognitive Load Theory: Cognitive load[2] is the amount of mental effort required to solve a problem. Cognitive load theory divides this effort into *intrinsic*, *extraneous*, and *germane*, and holds that people learn faster and better when extraneous load is reduced.

- Cognitivism: A theory of learning that holds that mental states and processes can and must be included in models of learning. See also behaviorism.

- Community of Practice: A self-perpetuating group of people who share and develop a craft or occupation, such as knitters,

---

[1]https://en.wikipedia.org/wiki/Bloom's_taxonomy
[2]https://en.wikipedia.org/wiki/Cognitive_load

musicians, or programmers. See also legitimate peripheral participation.

- Competent Practitioner: Someone who can do normal tasks with normal effort under normal circumstances. See also novice and expert.

- Concept Map: A picture of a mental model in which concepts are nodes in a graph and relationships are (labelled) arcs.

- Connectivism: A theory of learning which emphasizes its social aspects, particularly as enabled by the Internet and other technologies.

- Constructivism: A theory of learning that views learners as actively constructing knowledge.

- Content Knowledge: A person's understanding of a subject. See also general pedagogical knowledge and pedagogical content knowledge.

- Deliberate Practice: The act of observing performance of a task while doing it in order to improve ability.

- Diagnostic Power: The degree to which a wrong answer to a question or exercise tells the instructor what misconceptions a particular learner has.

- Educational Psychology: The study of how people learn. See also instructional design.

- Expert: Someone who can diagnose and handle unusual situations, knows when the usual rules do not apply, and tends to recognize solutions rather than reasoning to them. See also competent practitioner and novice.

- Expert Blind Spot: The inability of experts to empathize with novices who are encountering concepts or practices for the first time.

- Externalized Cognition: The use of graphical, physical, or verbal aids to augment thinking.

- Faded Example: A series of examples in which a steadily increasing number of key steps are blanked out. See also scaffolding.

- Fink's Taxonomy: A six-part non-hierarchical classification of understanding first proposed in [Fink2013] whose categories are *foundational knowledge*, *application*, *integration*, *human dimension*, *caring*, and *learning how to learn*. See also: Bloom's Taxonomy.

- Fixed Mindset: The belief that an ability is innate, and that failure is due to a lack of some necessary attribute. See also growth mindset.

- Fluid Representation: The ability to move quickly between different models of a problem.

- Formative Assessment: Assessment that takes place during a lesson in order to give both the learner and the instructor feedback on actual understanding. See also summative assessment.

- General Pedagogical Knowledge: A person's understanding of the general principles of teaching. See also content knowledge and pedagogical content knowledge.

- Growth Mindset: The belief that ability comes with practice. See also fixed mindset.

- Implementation Science: the study of how to translate research findings to everyday clinical practice.

- Impostor Syndrome: A feeling of insecurity about one's accomplishments that manifests as a fear of being exposed as a fraud.

- Inclusivity: Working actively to include people with diverse backgrounds and needs.

- Inquiry-Based Learning: The practice of allowing learners to ask their own questions, set their own goals, and find their own path through a subject.

- Instructional Design: The craft of creating and evaluating specific lessons for specific audiences. See also educational psychology.

- Jugyokenkyu: Literally "lesson study", a set of practices that includes having teachers routinely observe one another and discuss lessons to share knowledge and improve skills.

- Lateral Knowledge Transfer: The "accidental" transfer of knowledge that occurs when an instructor is teaching one thing, and the learner picks up another.

- Learned Helplessness: A situation in which people who are repeatedly subjected to negative feedback that they have no way to escape learn not to even try to escape when they could.

- Learner Persona: A brief description of a typical target learner for a lesson that includes their general background, what they already know, what they want to do, how the lesson will help them, and any special needs they might have.

- Learning Objective: What a lesson is trying to achieve.

- Learning Outcome: What a lesson actually achieves.

- Legitimate Peripheral Participation: Newcomers' participation in simple, low-risk tasks that a community of practice recognizes as valid contributions.

- Live Coding: The act of teaching programming by writing software in front of learners as the lesson progresses.

- Long-Term Memory: The part of memory that stores information for long periods of time. Long-term memory is very large, but slow. See also short-term memory.

- Minute Cards: A feedback technique in which learners spend a minute writing one positive thing about a lesson (e.g., one thing they've learned) and one negative thing (e.g., a question that still hasn't been answered).

- Novice: Someone who has not yet built a usable mental model of a domain. See also competent practitioner and expert.

- Pair Programming: A software development practice in which two programmers share one computer. One programmer (the driver) does the typing, while the other (the navigator) offers comments and suggestions in real time. Pair programming is often used as a teaching practice in programming classes.

- Parsons Problem: An assessment technique developed by Dale Parsons and others in which learners rearrange given material to construct a correct answer to a question.

- Pedagogical Content Knowledge (PCK): The understanding of how to teach a particular subject, i.e., the best order in which to introduce topics and what examples to use. See also content knowledge and general pedagogical knowledge.

- Peer Instruction: A teaching method in which an instructor poses a question and then students commit to a first answer, discuss answers with their peers, and commit to a (revised) answer.

- Persistent Memory: see long-term memory.

- Plausible Distractor: A wrong answer to a multiple-choice question that looks like it could be right. See also diagnostic power.

- Reflective Practice: see deliberate practice.

- Reverse Instructional Design: An instructional design method that works backwards from a summative assessment to formative assessments and thence to lesson content.

- Scaffolding: Extra material provided to early-stage learners to help them solve problems.

- Short-Term Memory: The part of memory that briefly stores information that can be directly accessed by consciousness.

- Situated Learning: A model of learning that focuses on people's transition from being newcomers to be accepted members of a community of practice.

- Stereotype Threat: A situation in which people feel that they are at risk of being held to stereotypes of their social group.

- Summative Assessment: Assessment that takes place at the end of a lesson to tell whether the desired learning has taken place.

- Tangible Artifact: Something a learner can work on whose state gives feedback about the learner's progress and helps the learner diagnose mistakes.

- Test-Driven Development: A software development practice in which programmers write tests first in order to give themselves concrete goals and clarify their understanding of what "done" looks like.

- Understanding by Design: see reverse instructional design.

- Working Memory: see short-term memory.

20 Lesson Design Template

Designing a good course is as hard as designing good software. To help you, this appendix summarizes a process based on evidence-based teaching practices:

- It lays out a step-by-step progression to help you figure out what to think about in what order.
- It provides spaced check-in points so you can re-scope or redirect effort.
- The end product specifies deliverables clearly so you can finish development without major surprises.
- Everything from Step 2 onward goes into your final course, so there is no wasted effort.
- Writing sample exercises early lets you check that everything you want your students to do actually works.

This backward design process was developed independently by [Wiggins2005], [Biggs2011], and [Fink2013]. We have slimmed it down by removing steps related to meeting curriculum guidelines and other institutional requirements.

Note: the steps are described in order of increasing detail, but the process itself is always iterative. You will frequently go back to revise earlier work as you learn something from your answer to a later question or realize that your initial plan isn't going to play out the way you first thought.

20.1 Step 1: Brainstorming

The first step is to throw together some rough ideas so that you and your colleagues can make sure your thoughts about the course are aligned. To do this, write some point-form answers to three or four of the questions listed below. You aren't expected to answer all of

them, and you may pose and answer others if you think it's helpful, but you should always include a couple of answers to the first.

1. What problem(s) will student learn how to solve?
2. What concepts and techniques will students learn?
3. What technologies, packages, or functions will students use?
4. What terms or jargon will you define?
5. What analogies will you use to explain concepts?
6. What heuristics will help students understand things?
7. What mistakes or misconceptions do you expect?
8. What datasets will you use?

You may not need to answer every question for every course, and you will often have questions or issues we haven't suggested, but couple of hours of thinking at this stage can save days of rework later on.

Checkin: a rough scope for the course that you have agreed with your colleagues.

20.2   Step 2: Who Is This Course For?

"Beginner" and "expert" mean different things to different people, and many factors besides pre-existing knowledge influence who a course is suitable for. The second step in designing a course is therefore to figure out who your audience is. To do this, you should either create some learner personas[1], or (preferably) reference ones that you and your colleagues have drawn up together.

After you are done brainstorming, you should go through these personas and decide which of them your course is intended for, and how it will help them. While doing this, you should make some notes about what specific prerequisite skills or knowledge you expect students to have above and beyond what's in the persona.

Checkin: brief summaries of who your course will help and how.

20.3   Step 3: What Will Learners Do Along the Way?

The best way to make the goals in Step 1 firmer is to write full descriptions of a couple of exercises that students will be able to do toward the end of the course. Writing exercises early is directly analogous to test-driven development[2]: rather than working forward

---

[1]lessons.md
[2]https://en.wikipedia.org/wiki/Test-driven_development

from a (probably ambiguous) set of learning objectives, designers work backward from concrete examples of where their students are going. Doing this also helps uncover technical requirements that might otherwise not be found until uncomfortably late in the lesson development process.

To complement the full exercise descriptions, you should also write brief point-form descriptions of one or two exercises per lecture hour to show how quickly you expect learners to progress. (Again, these serve as a good reality check on how much you're assuming, and help uncover technical requirements.) One way to create these "extra" exercises is to make a point-form list of the skills needed to solve the major exercises and create an exercise that targets each.

Checkin: 1-2 fully explained exercises that use the skills the student is to learn, plus half a dozen point-form exercise outlines.

Note: be sure to include solutions with example code so that you can check that your software can do everything you need.

20.4   Step 4: How Are Concepts Connected?

In this stage, you put the exercises in a logical order then derive a point-form course outline for the entire course from them. This is also when you will consolidate the datasets your formative assessments have used.

Checkin: a course outline.

Note:

- The final outline should be at the lecture and formative assessment level, e.g., one major bullet point for each hour of work with 3-4 minor bullet points for the episodes in that hour.

- It's common to change assessments in this stage so that they can build on each other.

- You are likely to discover things you forgot to list earlier during this stage, so don't be surprised if you have to double back a few times.

20.5   Step 5: Course Overview

You can now write a course overview consisting of:

- a one-paragraph description (i.e., a sales pitch to students)
- half a dozen learning objectives
- a summary of prerequisites

Doing this earlier often wastes effort, since material is usually added, cut, or moved around in earlier steps.

Checkin: course description, learning objectives, and prerequisites.

Note: see the appendix for a discussion of how to write good learning objectives.

20.6   Reminder

As noted at the start, this process is described as a sequence, but in practice you will loop back repeatedly as each stage informs you of something you overlooked.

21 Checklists

Atul Gawande's 2007 article "The Checklist[1]" popularized the idea that using checklists can save lives (and make many other things better too). The results of recent studies have been more nuanced [Aveling2013], [Urbach2014], but we still find them useful, particularly when bringing new instructors onto a team.

The checklists below are used before, during, and after instructor training events, and can easily be adapted for end-learner workshops as well. We recommend that every group build and maintain its own checklists customized for its instructors' and learners' needs.

21.1 Scheduling the Event

1. Decide if it will be in person, online for one site, or online for several sites.
2. Talk through expectations with the host(s) and make sure that everyone agrees on who is covering travel costs.
3. Determine who is allowed to take part: is the event open to all comers, restricted to members of one organization, or something in between?
4. Arrange trainers.
5. Arrange space, including breakout rooms for video recording.
6. Choose dates. If it is in person, book travel.
7. Get names and email addresses of attendees from host(s).
8. Make sure they are added to the registration system.

21.2 Setting Up

1. Set up a web page with details on the workshop, including date, location, and a list of what participants need to bring.
2. Check whether any attendees have special needs.

---

[1]http://www.newyorker.com/magazine/2007/12/10/the-checklist

3. If the workshop is online, test the video conferencing link.
4. Make sure attendees will all have network access.
5. Create an Etherpad or Google Doc for shared notes.
6. Email attendees a welcome message that includes a link to the workshop home page, background readings, and a description of any pre-requisite tasks.

21.3 At the Start of the Event

1. Remind everyone of the code of conduct.
2. Collect attendance.
3. Distribute sticky notes.
4. Collect any relevant online account IDs.

21.4 At the End of the Event

1. Update attendance records. Be sure to also record who participated as an instructor or helper.
2. Administer a post-workshop survey.
3. Update the course notes and/or checklists.

21.5 Travel Kit

Here are a few things instructors take with them when they travel to teach:

- sticky notes
- cough drops
- comfortable shoes
- a small notepad
- a spare power adapter
- a spare shirt
- deodorant
- a variety of video adapters
- laptop stickers
- a toothbrush or some mouthwash
- a granola bar or some other emergency snack
- Eno or some other antacid (because road food)
- business cards
- a printed copy of the notes, or a tablet or other device
- an insulated cup for tea/coffee
- spare glasses/contacts
- a notebook and pen

- a portable WiFi hub (in case the room's network isn't working)
- extra whiteboard markers
- a laser pointer
- a packet of wet wipes (because spills happen)
- USB drives with installers for various operating systems
- running shoes, a bathing suit, a yoga mat, or whatever else you exercise in or with

22 The Rules

It's impossible to put everything that matters about teaching and learning on a single page, but these points are always worth remembering.

1. Be kind: all else is details.

2. You are not your learners.

3. Most people would rather fail than change.

4. Never teach alone.

5. No lesson survives first contact with learners.

6. Nobody will be more excited about the lesson than you are.

7. Every lesson is too short from the teacher's point of view and too long from the learner's.

8. Never hesitate to sacrifice truth for clarity.

9. Every mistake is a lesson.

10. "I learned this a long time ago" is not the same as "this is easy".

11. Ninety percent of magic consists of knowing one extra thing.

12. You can't help everyone, but you can always help someone.

23 Why I Teach

When I first started volunteering at the University of Toronto, students sometimes asked me why I did it. This was my answer:

When I was your age, I thought universities existed to teach people how to learn. Later, in grad school, I thought universities were about doing research and creating new knowledge. Now that I'm in my forties, though, I've realized that what we're really teaching you is how to take over the world, because you're going to have to one day whether you like it or not.

My parents are in their seventies. They don't run the world any more: it's people my age who pass laws, set interest rates, and make life-and-death decisions in hospitals. As scary as it is, we have become the grownups.

Twenty years from now, though, we'll be heading for retirement and you will be in charge. That may sound like a long time when you're nineteen, but take three breaths and it's gone. That's why we give you problems whose answers can't be cribbed from last year's notes. That's why we put you in situations where you have to figure out what needs to be done right now, what can be left for later, and what you can simply ignore. It's because if you don't learn how to do these things now, you won't be ready to do them when you have to.

It's all true, but isn't the whole story. I don't want people to make the world a better place so that I can retire in comfort. I want them to do it because it's the greatest adventure of our time. A hundred and fifty years ago, most societies still practiced slavery. A hundred years ago, when my grandmother was young, she wasn't legally a person[1] in Canada. Fifty years ago, most of the world's people suffered under totalitarian rule; in the year I was born, judges could–and did–order

---

[1]https://en.wikipedia.org/wiki/The_Famous_Five_(Canada)

electroshock therapy to "cure" homosexuals. Yes, there's still a lot wrong with the world, but look at how many more choices we have than our grandparents did. Look at how many more things we can know, and be, and enjoy.

This didn't happen by chance. It happened because millions of people made millions of little decisions, the sum of which was a better world. We don't think of these day-to-day decisions as political, but every time we buy one brand of running shoe instead of another or shout an anatomical insult instead of a racial one at a cab driver, we're choosing one vision of the world instead of another.

In his 1947 essay "Why I Write[2]", George Orwell said:

> *In a peaceful age I might have written ornate or merely descriptive books, and might have remained almost unaware of my political loyalties. As it is I have been forced into becoming a sort of pamphleteer... Every line of serious work that I have written since 1936 has been written, directly or indirectly, against totalitarianism... It seems to me nonsense, in a period like our own, to think that one can avoid writing of such subjects. Everyone writes of them in one guise or another. It is simply a question of which side one takes...*

Replace "writing" with "teaching" and you'll have the reason I do what I do. The world doesn't get better on its own. It gets better because people make it better: penny by penny, vote by vote, and one lesson at a time.

---

[2]http://www.resort.com/~prime8/Orwell/whywrite.html